

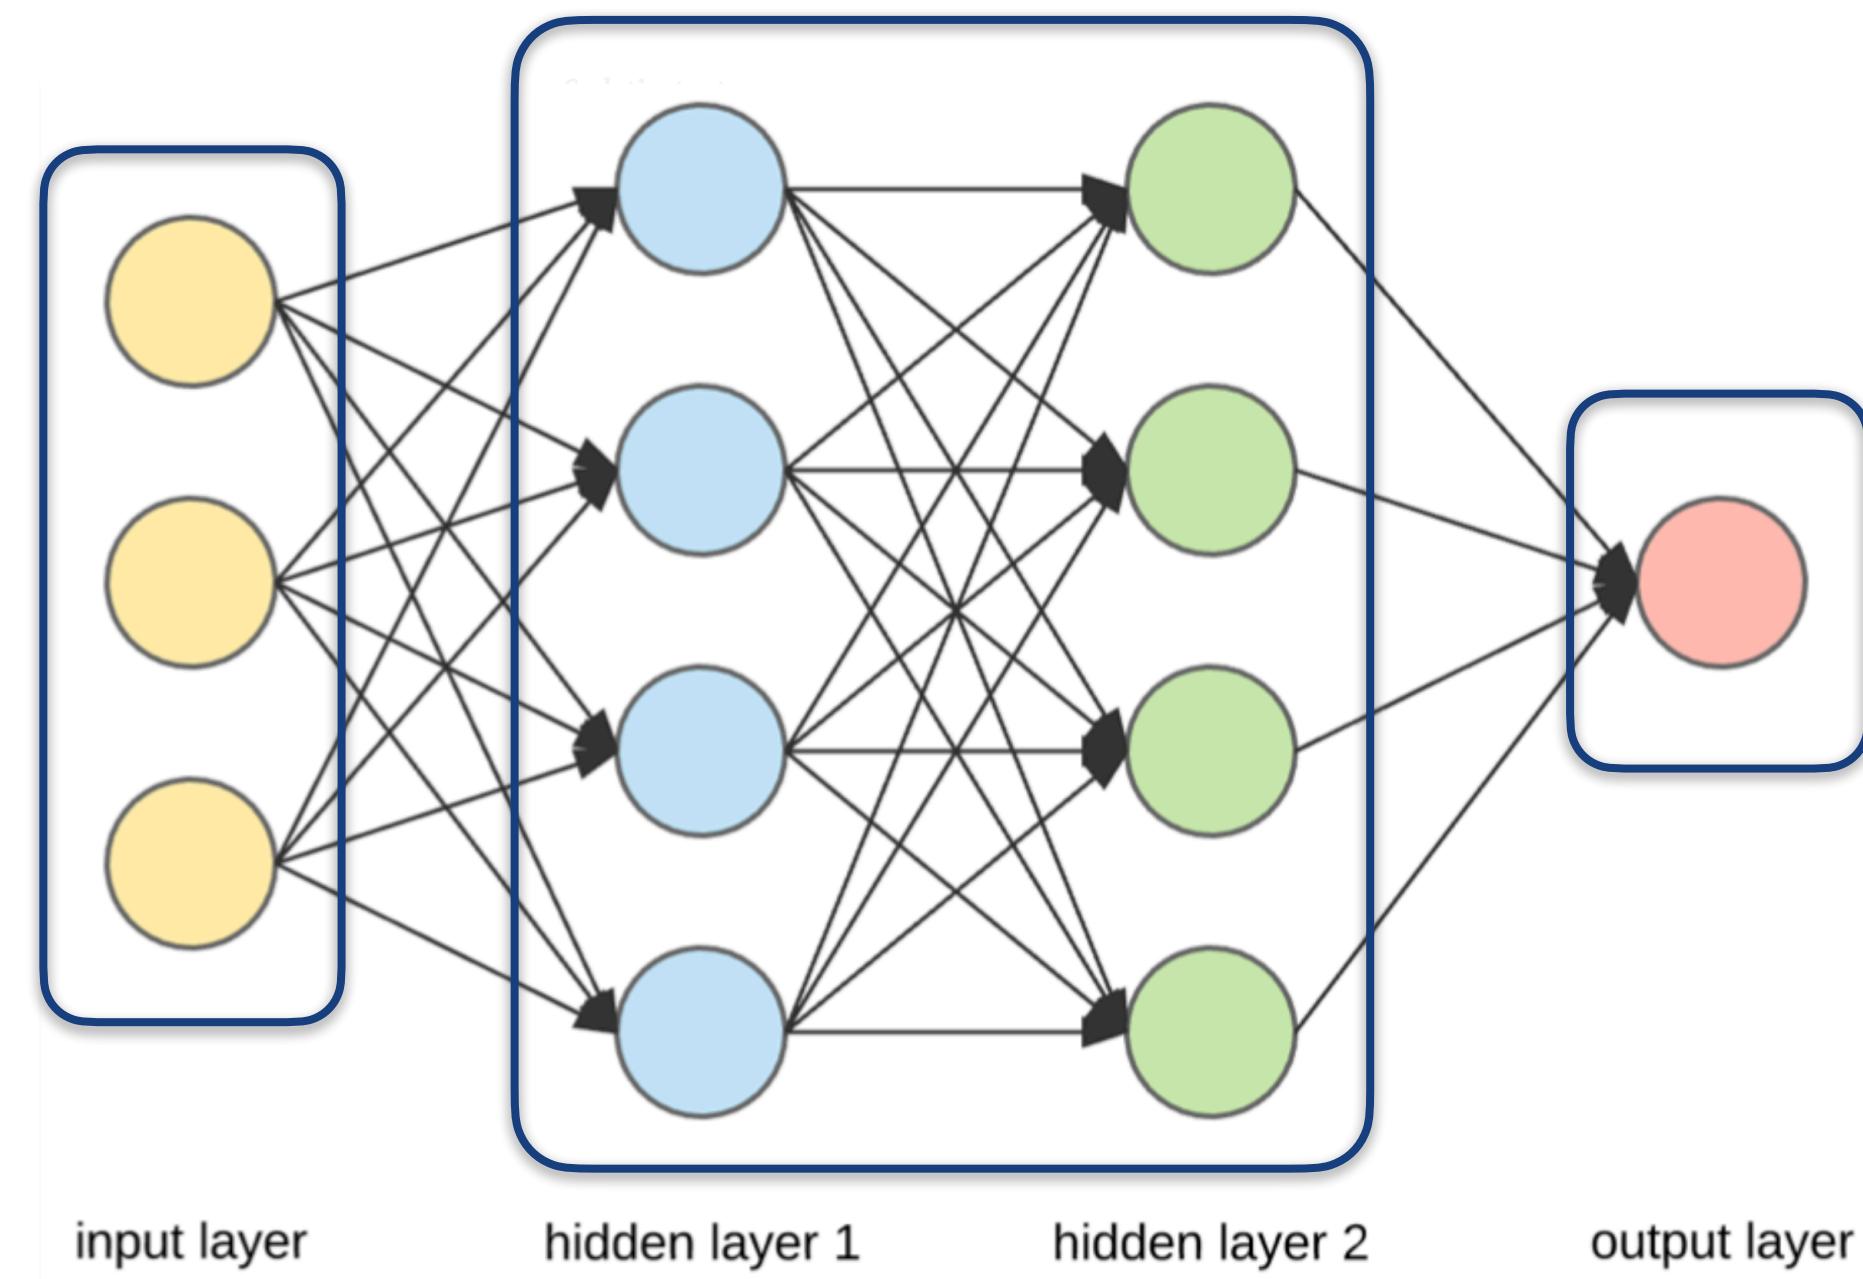
# Machine Learning Applications

## Neural Network

# Introduction

# Neural Network

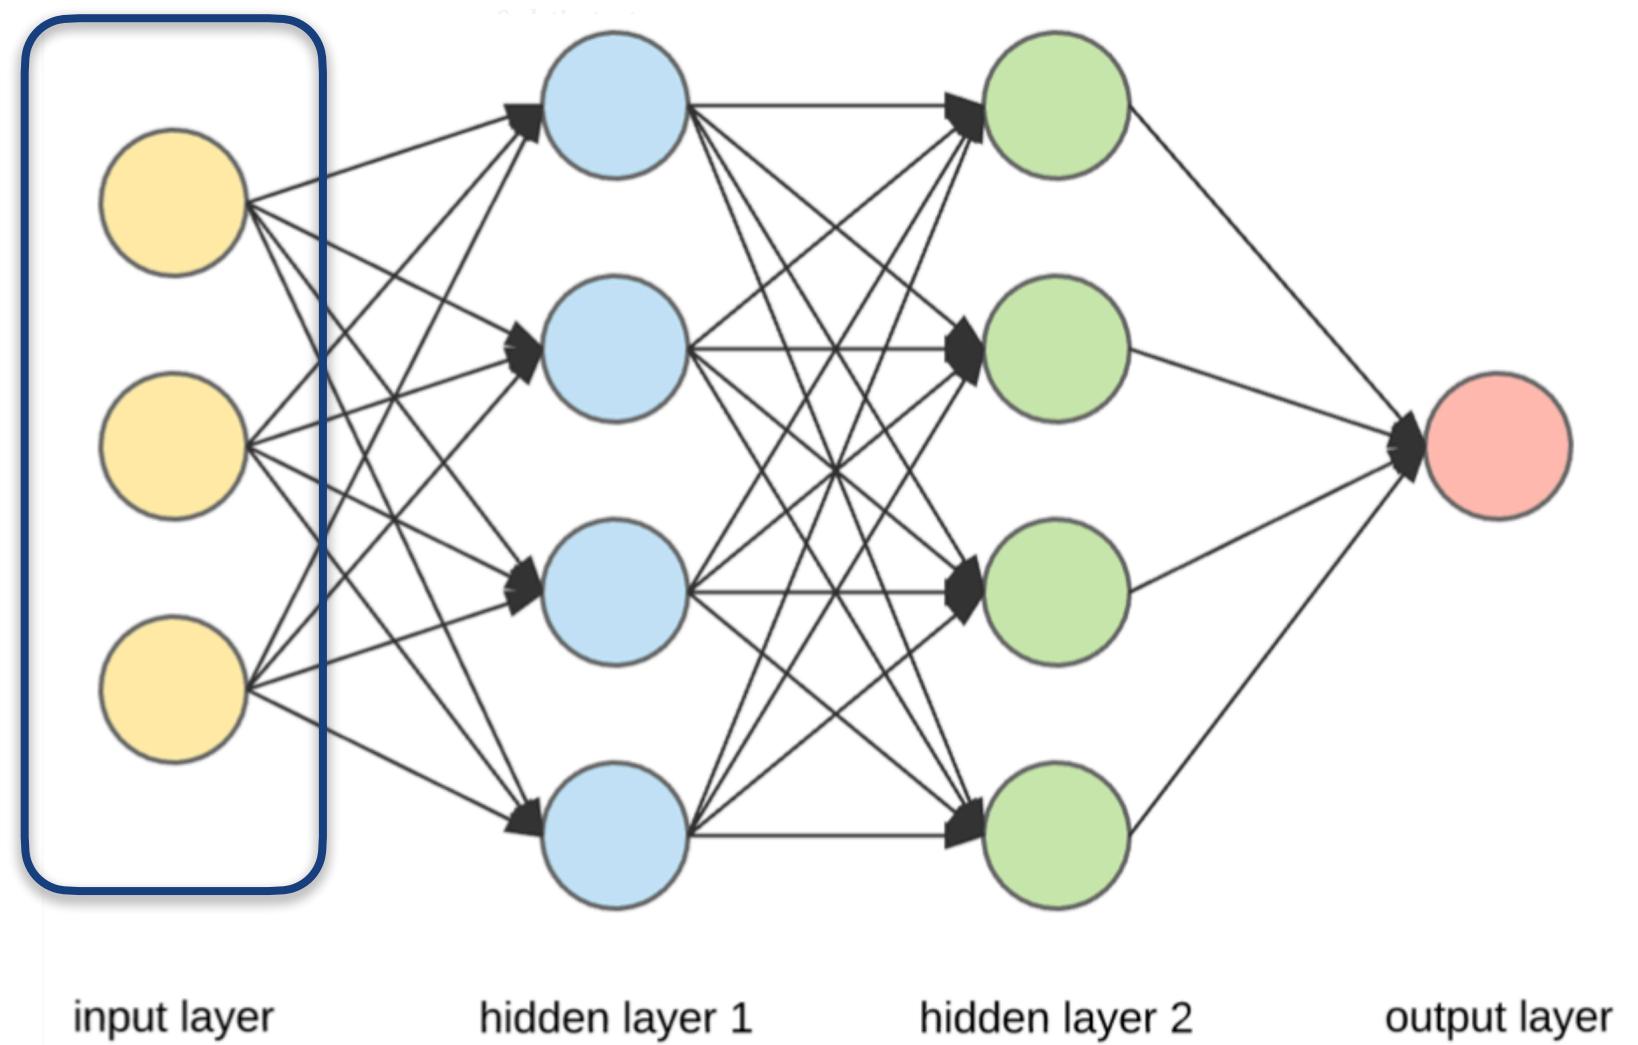
- A Neural Network is made of interconnected layers of computing units
- Common layers – Input, Hidden and Output



A Simple Neural Network

# Input Layer

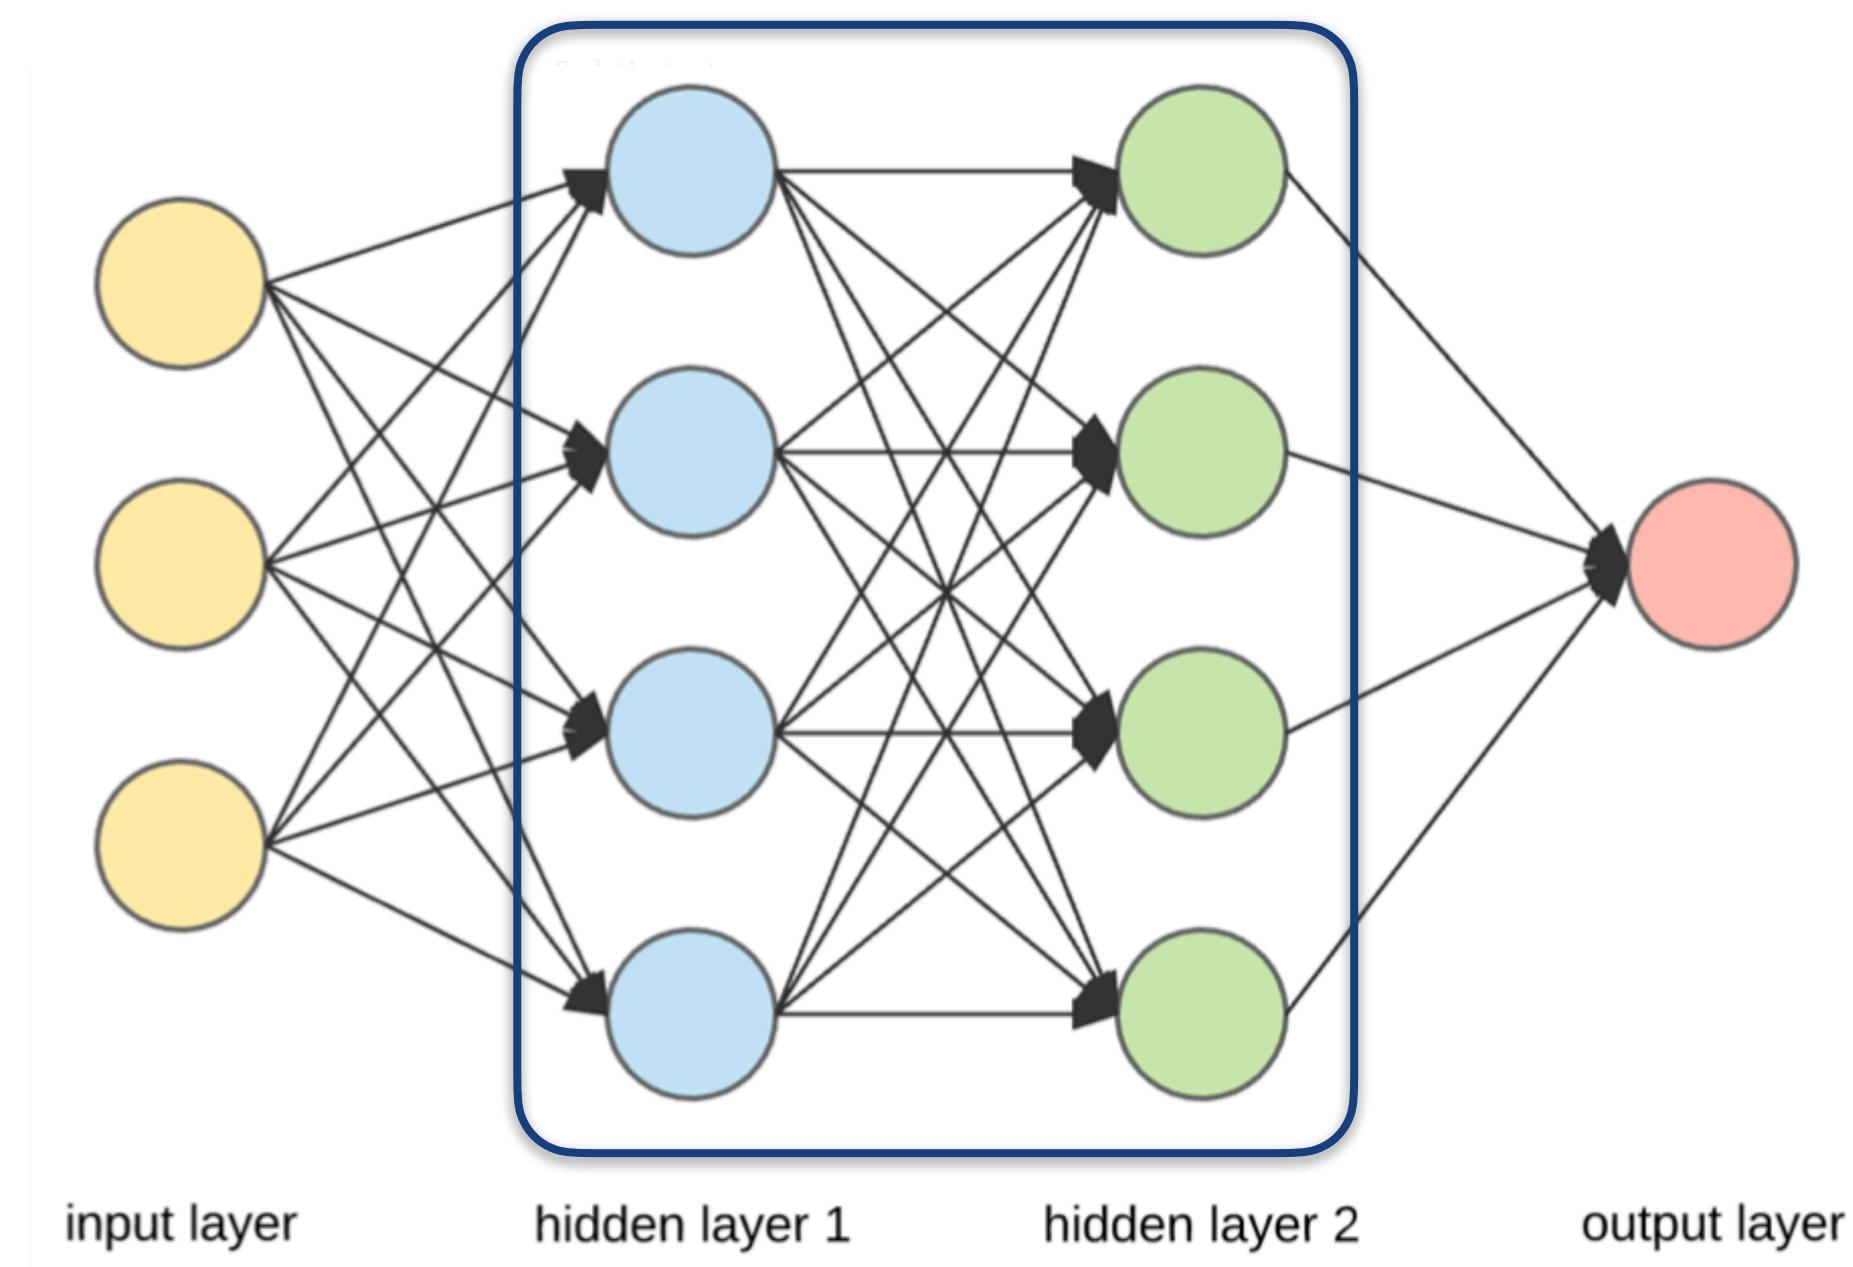
- The Input layer contains the inputs for computation
- If we are training a neural network to predict if a person is overweight, then the inputs could be Age, Weight and Height



A Simple Neural Network

# Hidden Layers

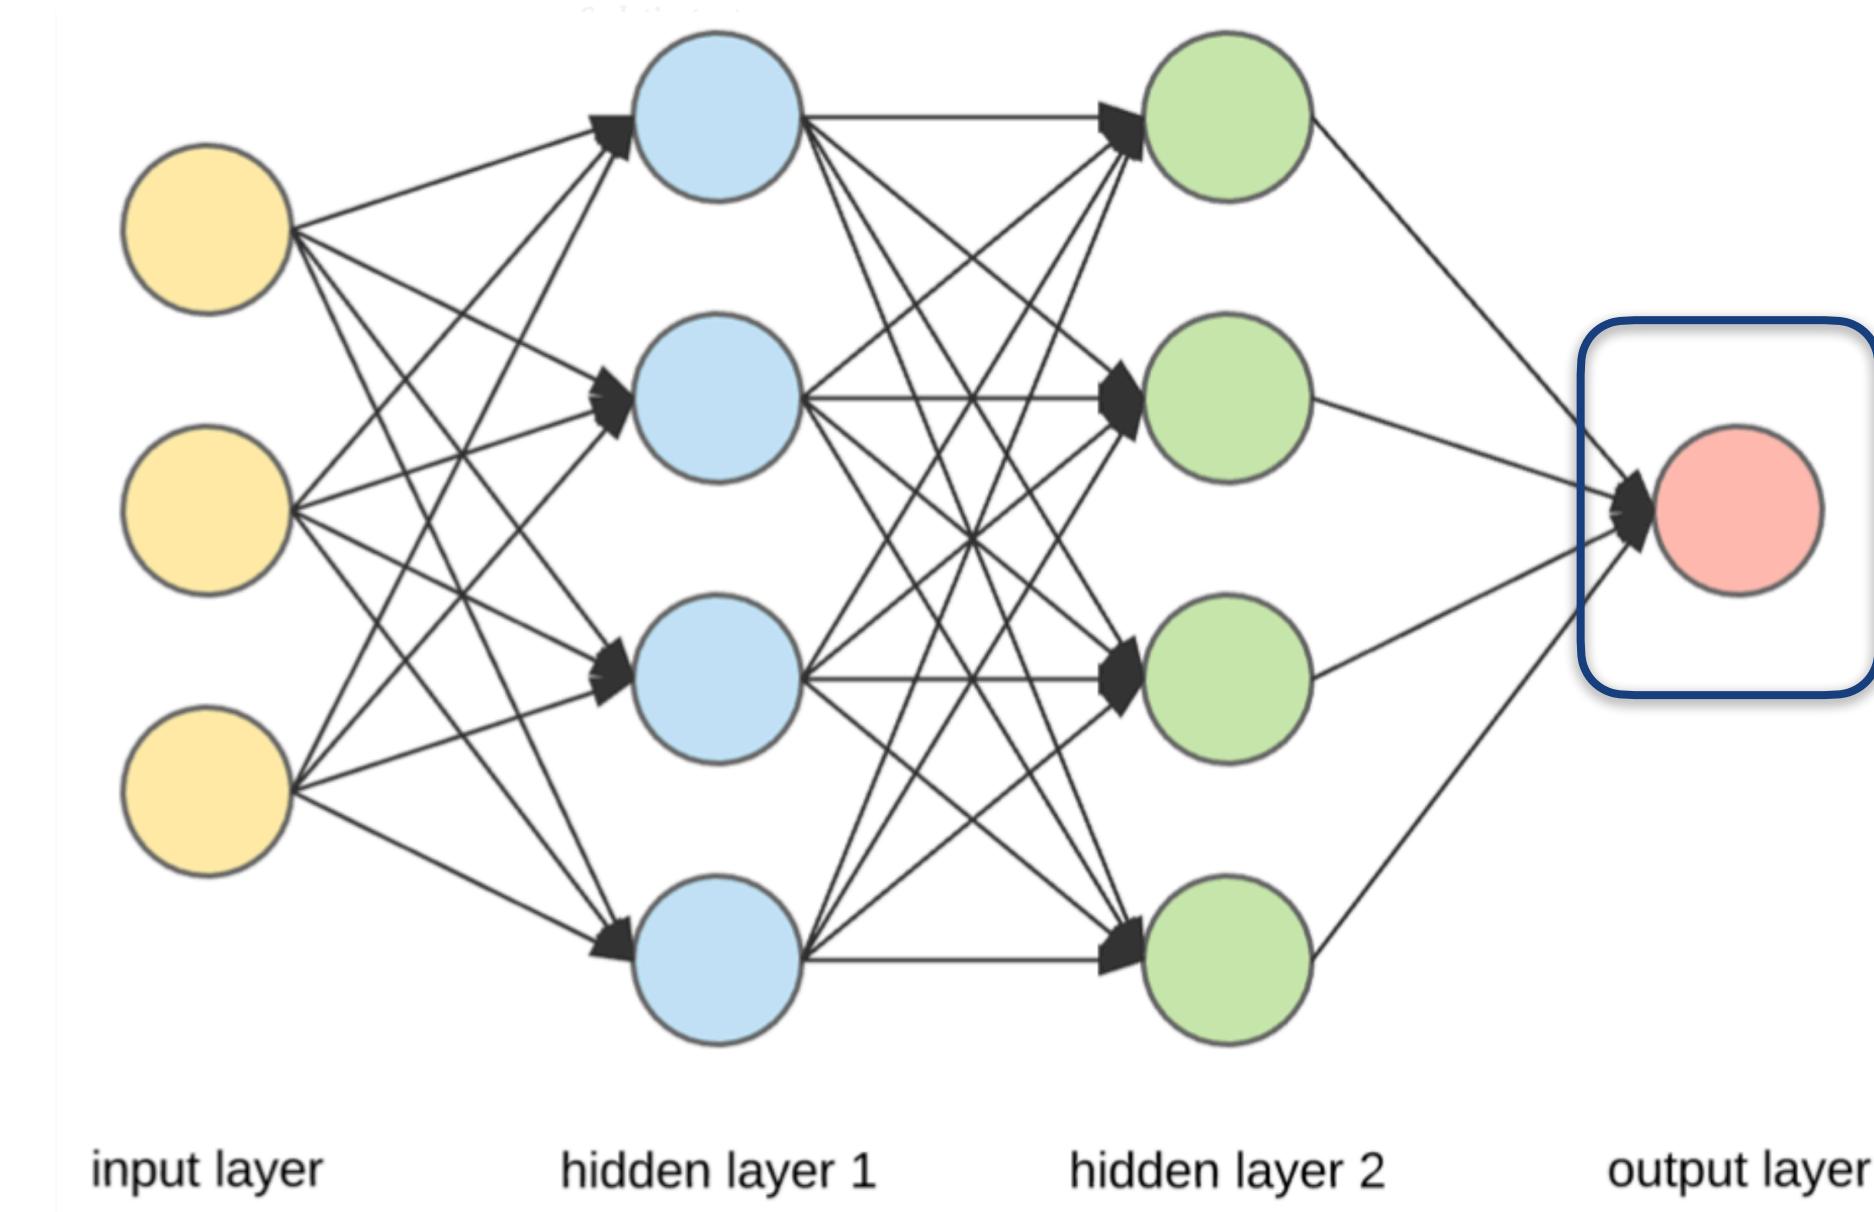
- Hidden layers are the layers between the Input and Output layers
- Bulk of computations happens in the Hidden layers



A Simple Neural Network

# Output Layer

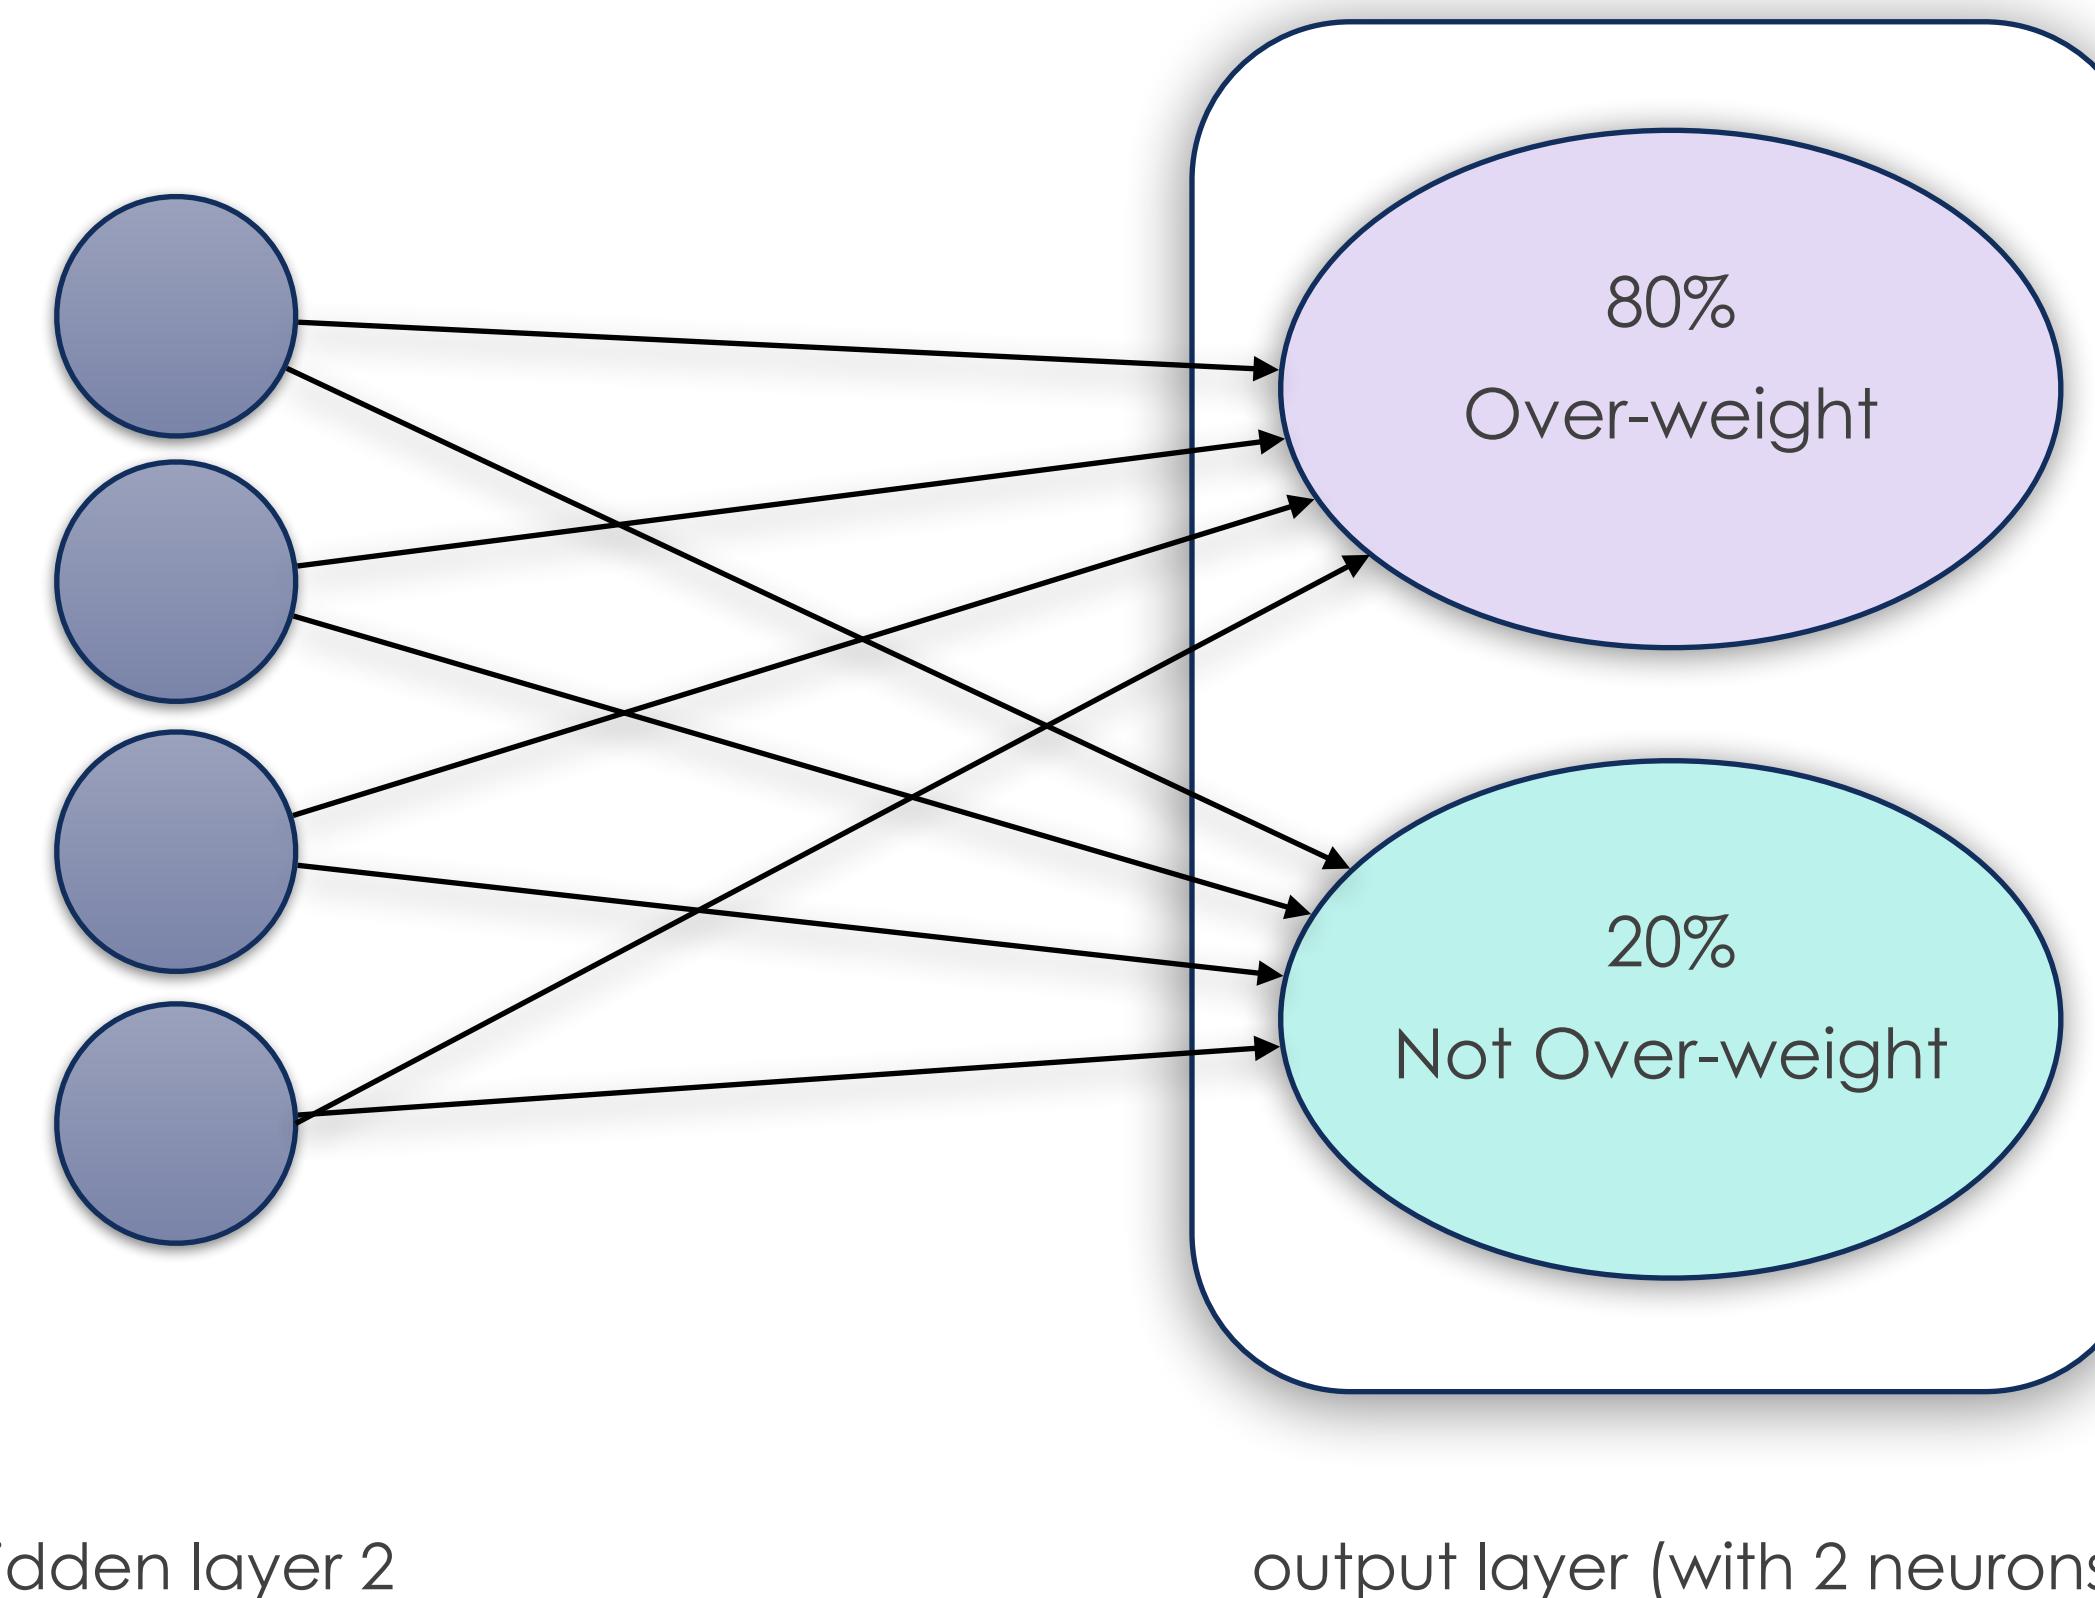
- The Output layer gives the computed prediction from a neural network
- For predicting if a person is overweight, the output can be a value between 0 and 1, where any value above 0.5 denotes the person is overweight



A Simple Neural Network

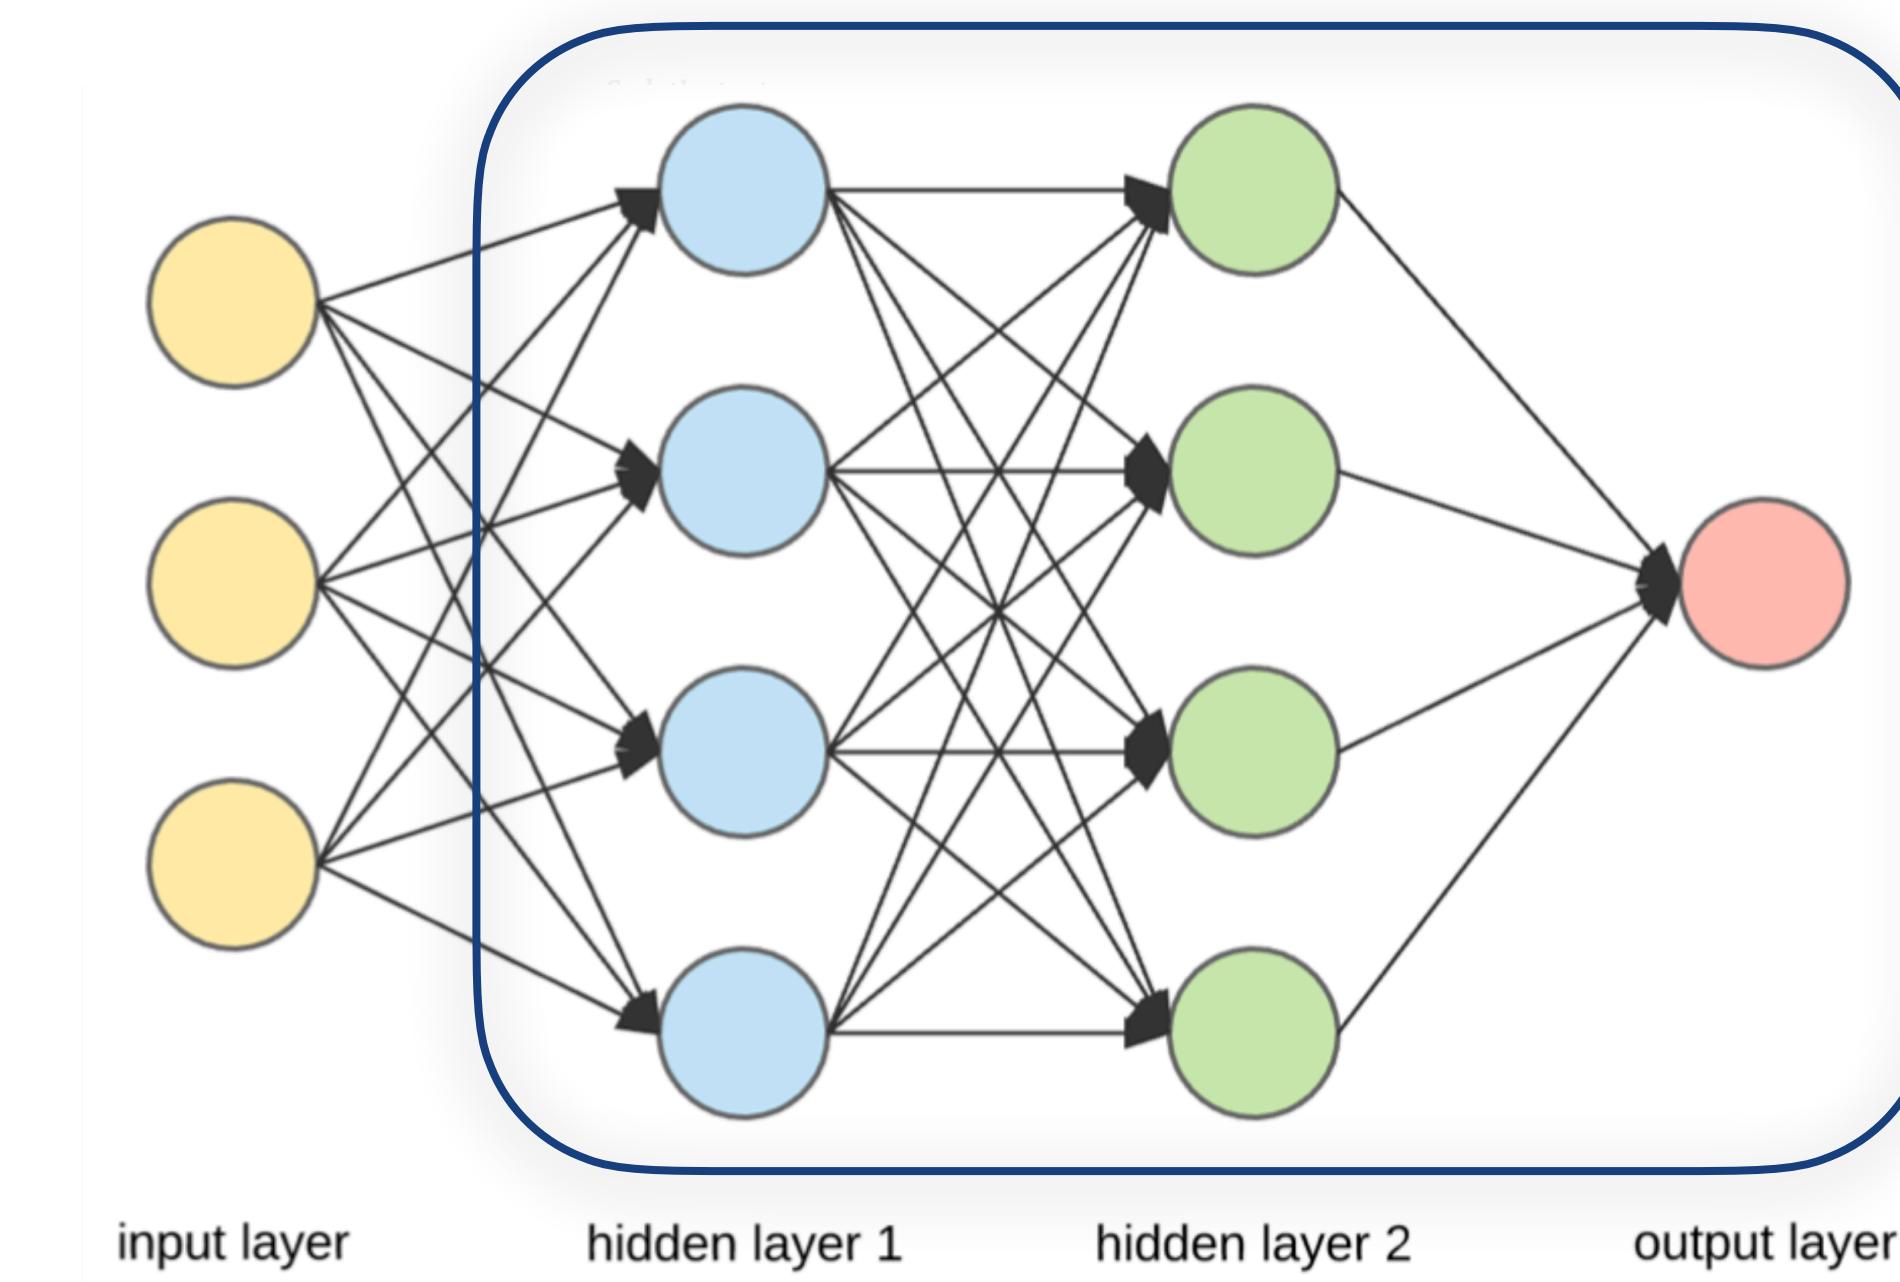
# Output Layer

- Alternatively, our Output layer could have been designed with 2 output nodes in mind, where each node carries a probability of the person being overweight



# A N-layer Neural Network

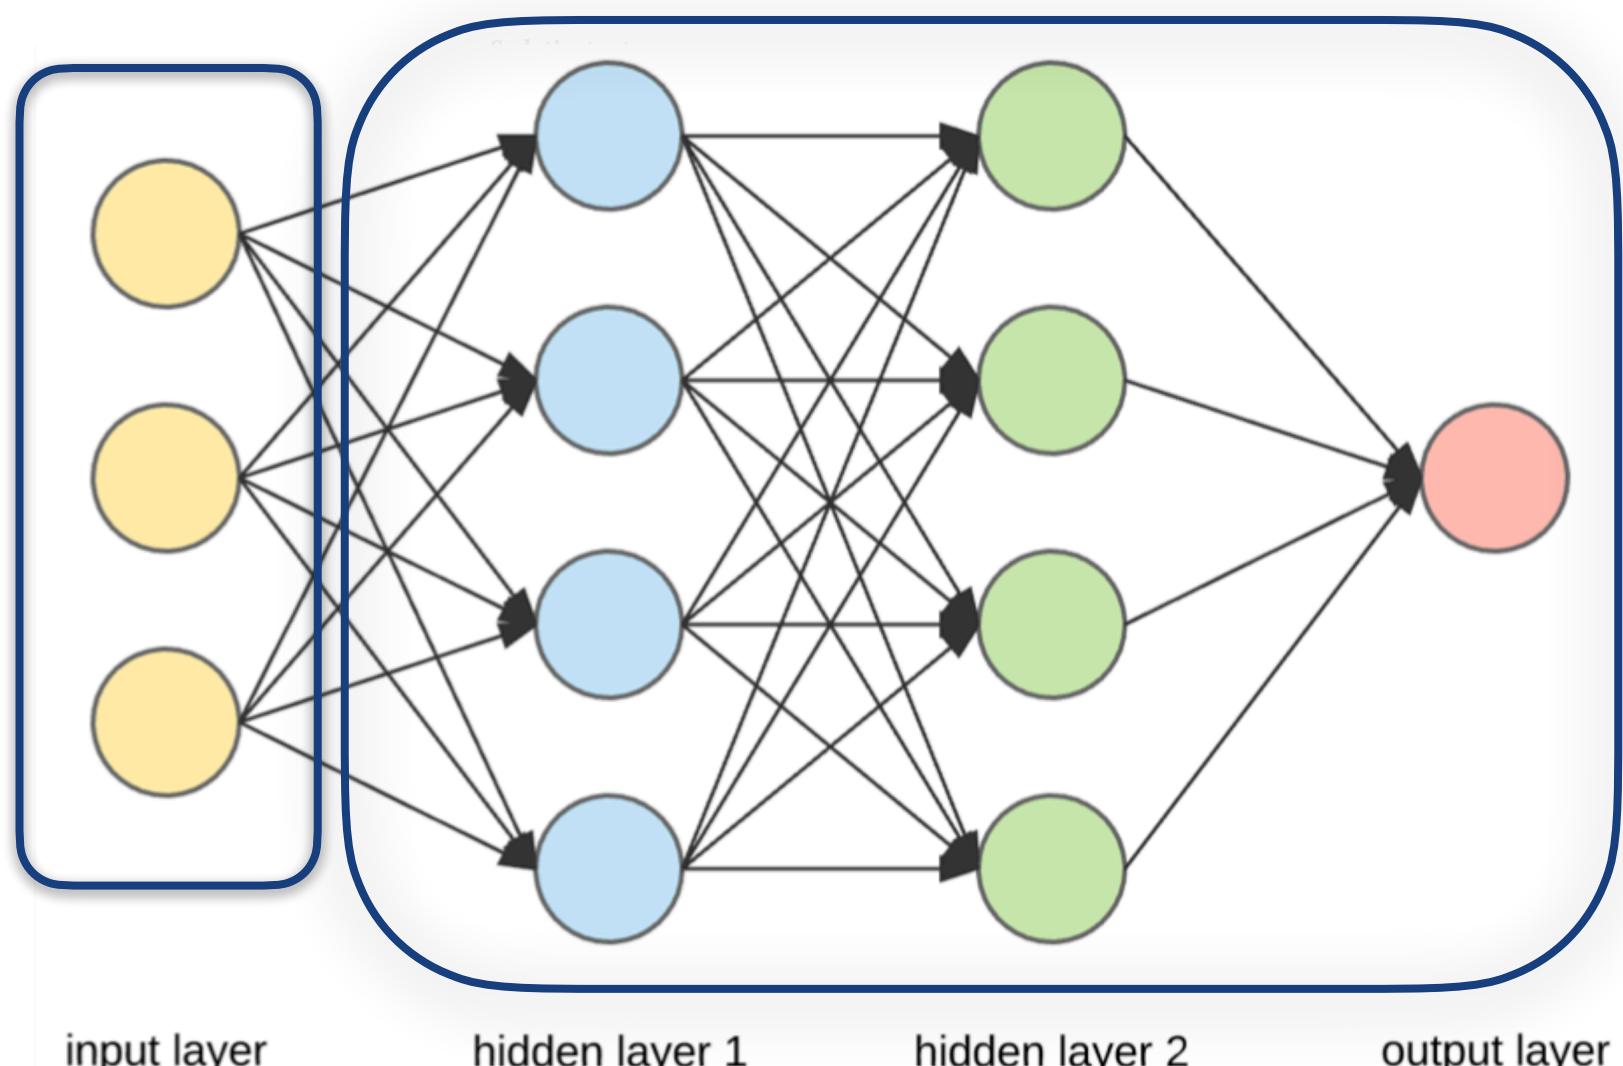
- A Neural Network is called N-layer if the sum of hidden and output layers is N
- The diagram depicts a 3-layer Neural Network



A Simple Neural Network

# Neurons

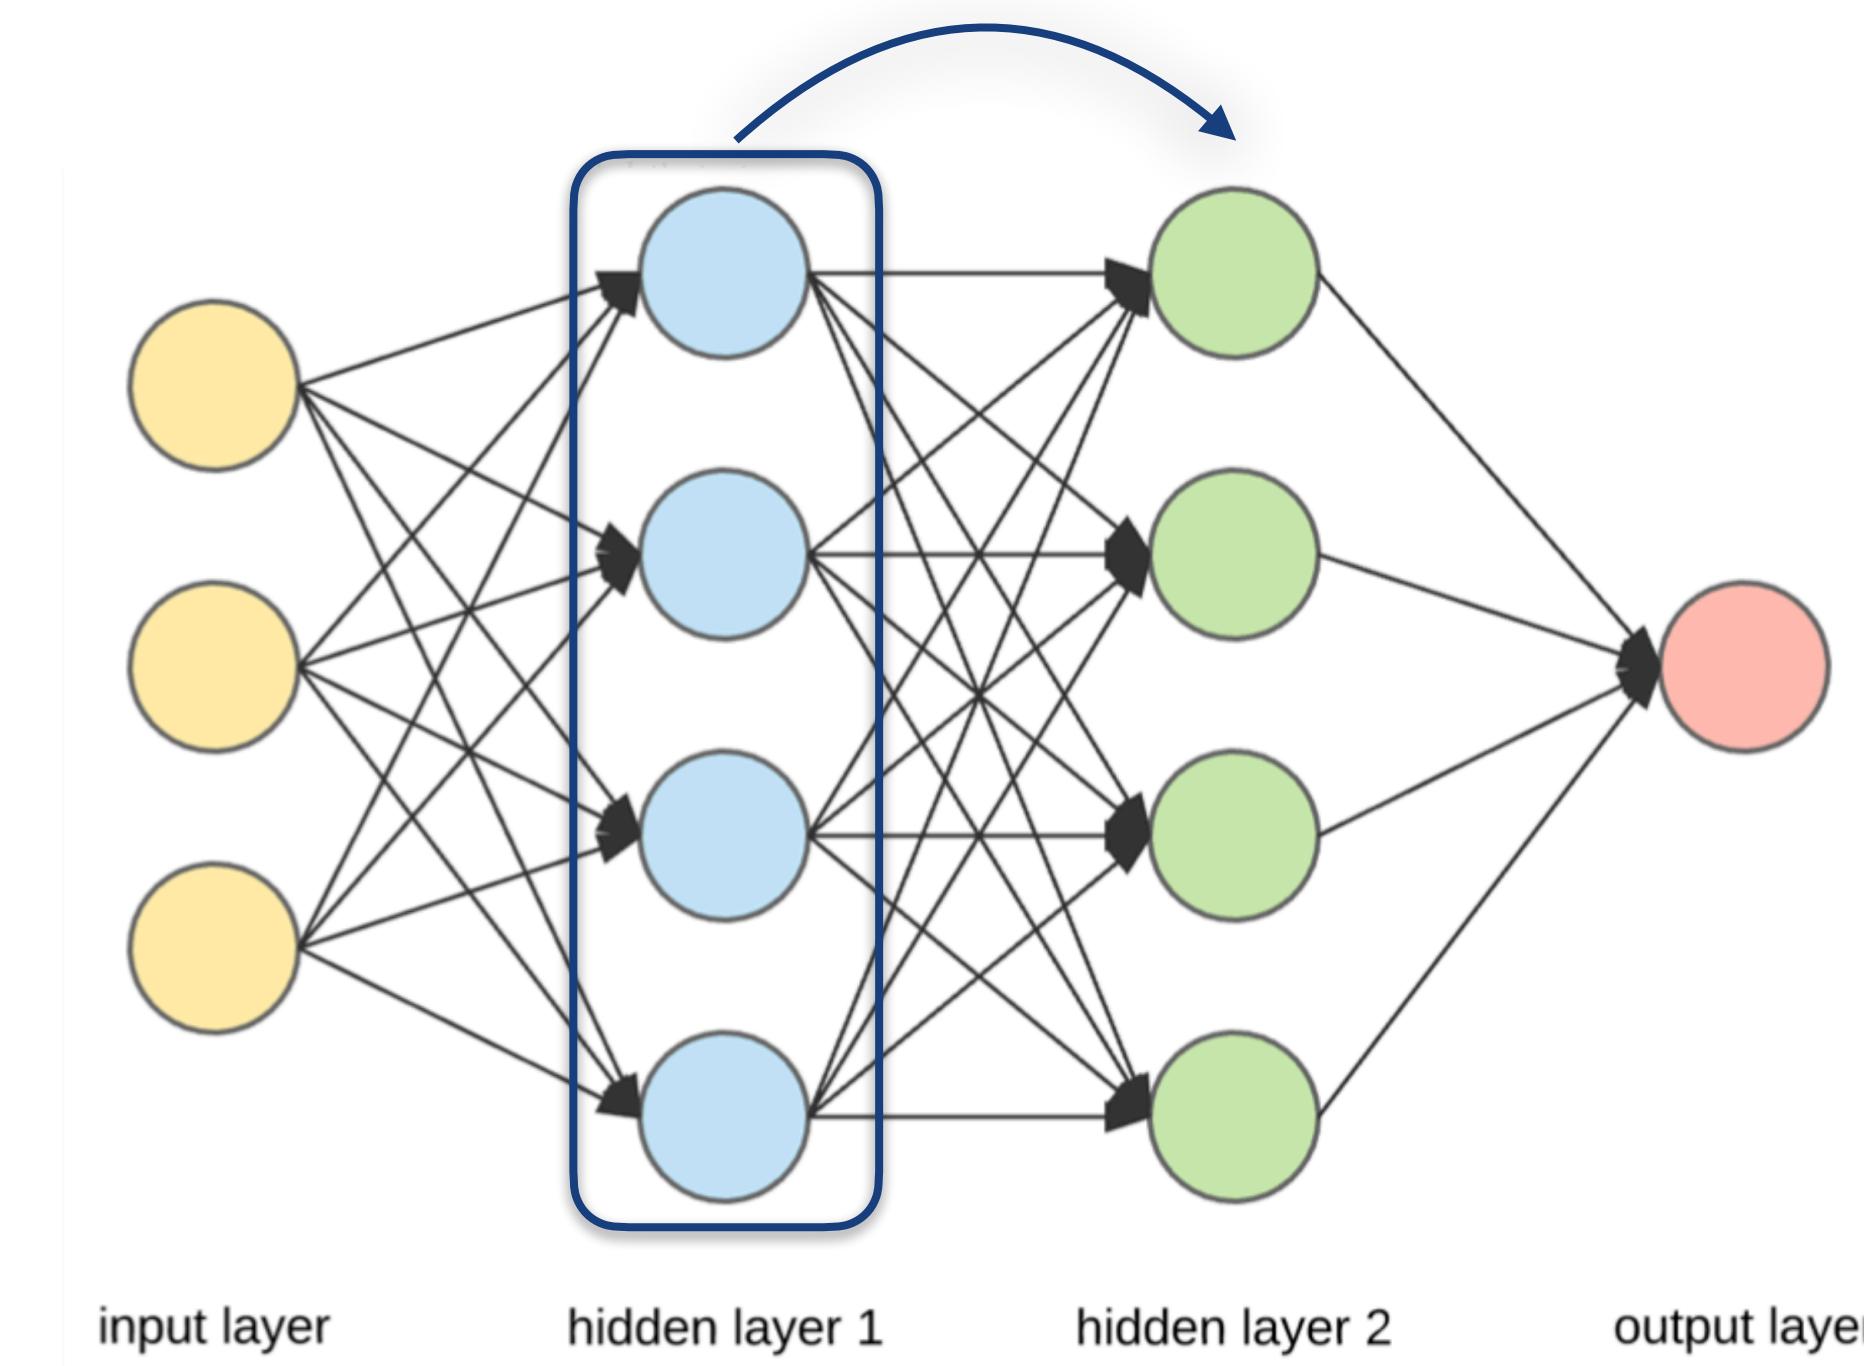
- The nodes in a neural network are called neurons
- The **blue**, **green** and **red** nodes are neurons
- The **yellow** nodes are not neurons, they are just input values



A Simple Neural Network

# Neurons

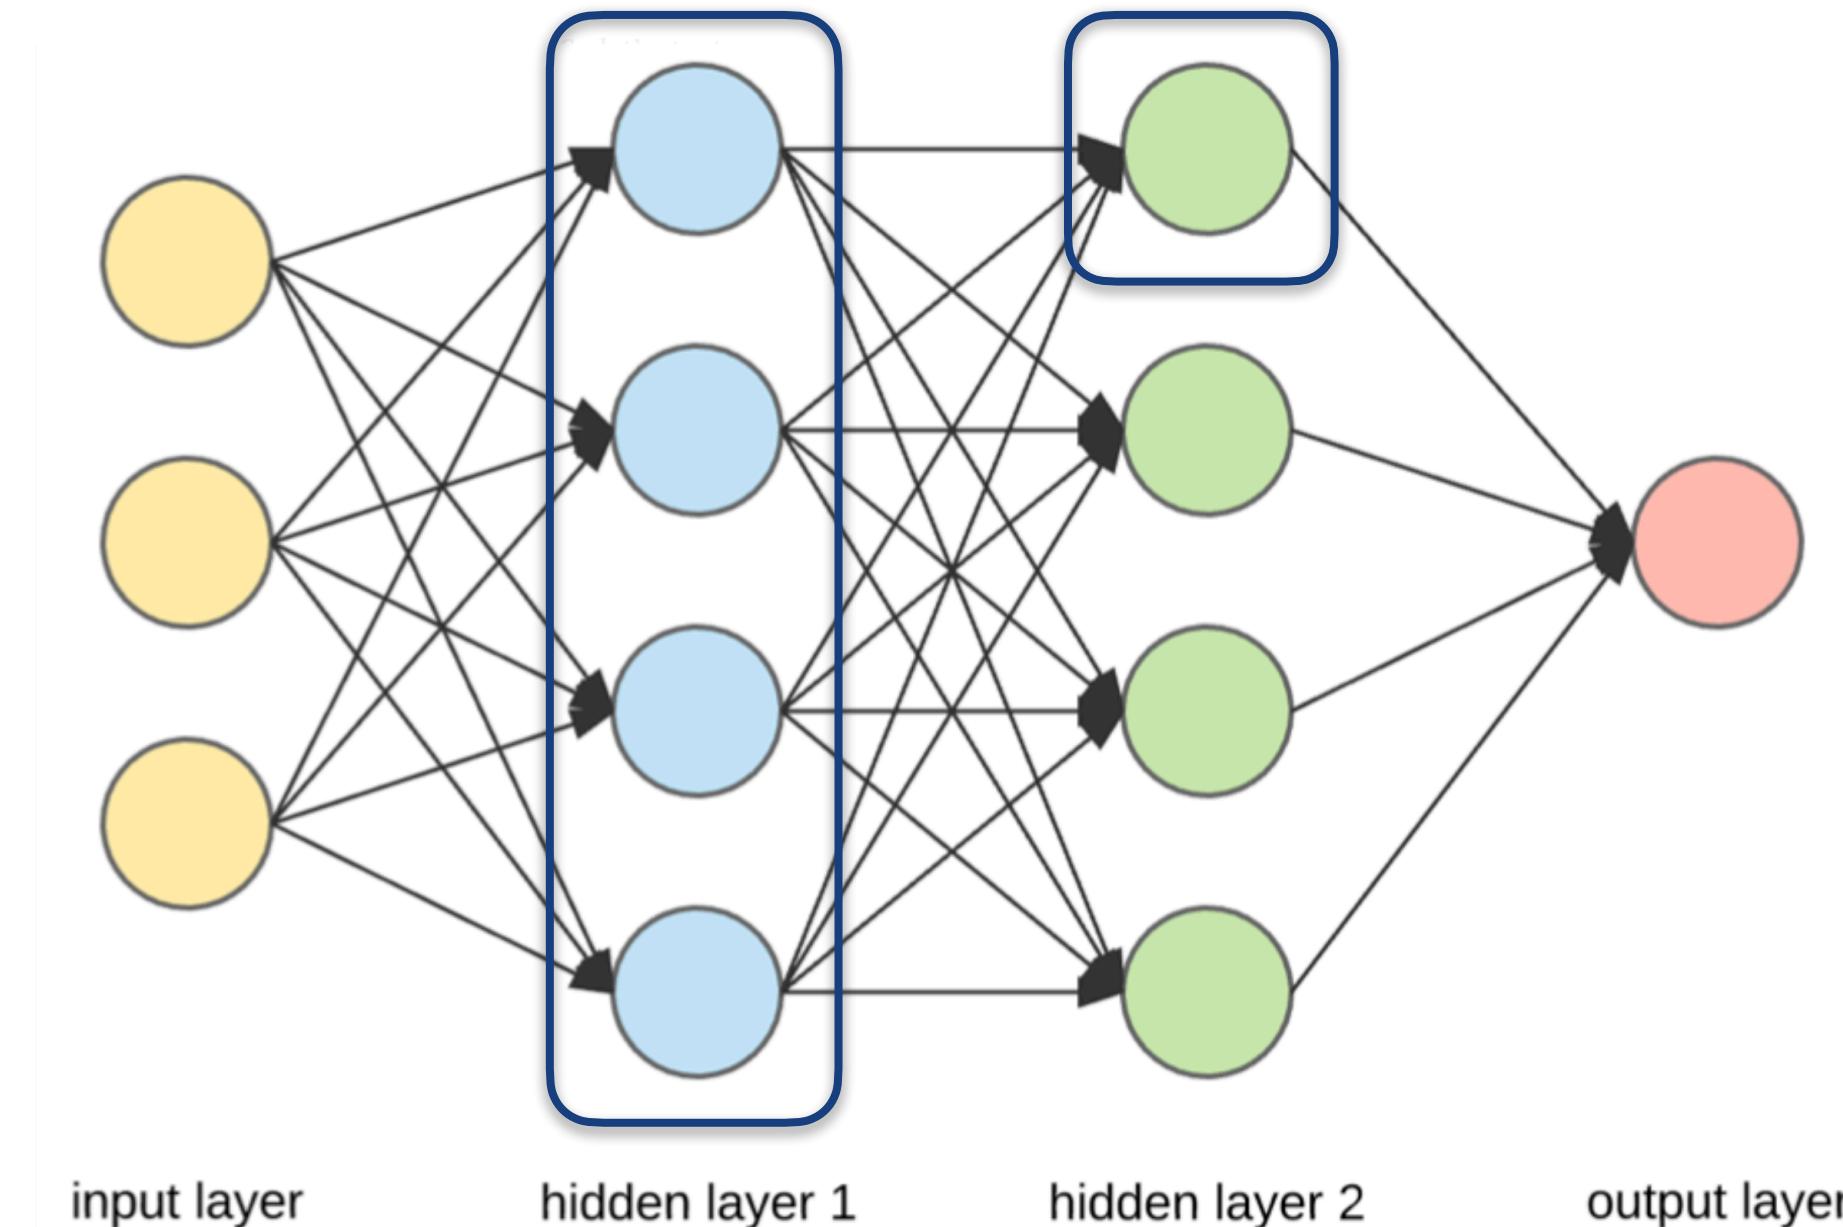
- Each neuron contains mathematical functions
- Each neuron applies its mathematical functions on values that comes from a previous layer



A Simple Neural Network

# Neurons

- Each **green** neuron applies its mathematical functions on values coming from its connected 4 **blue** neurons in the previous layer

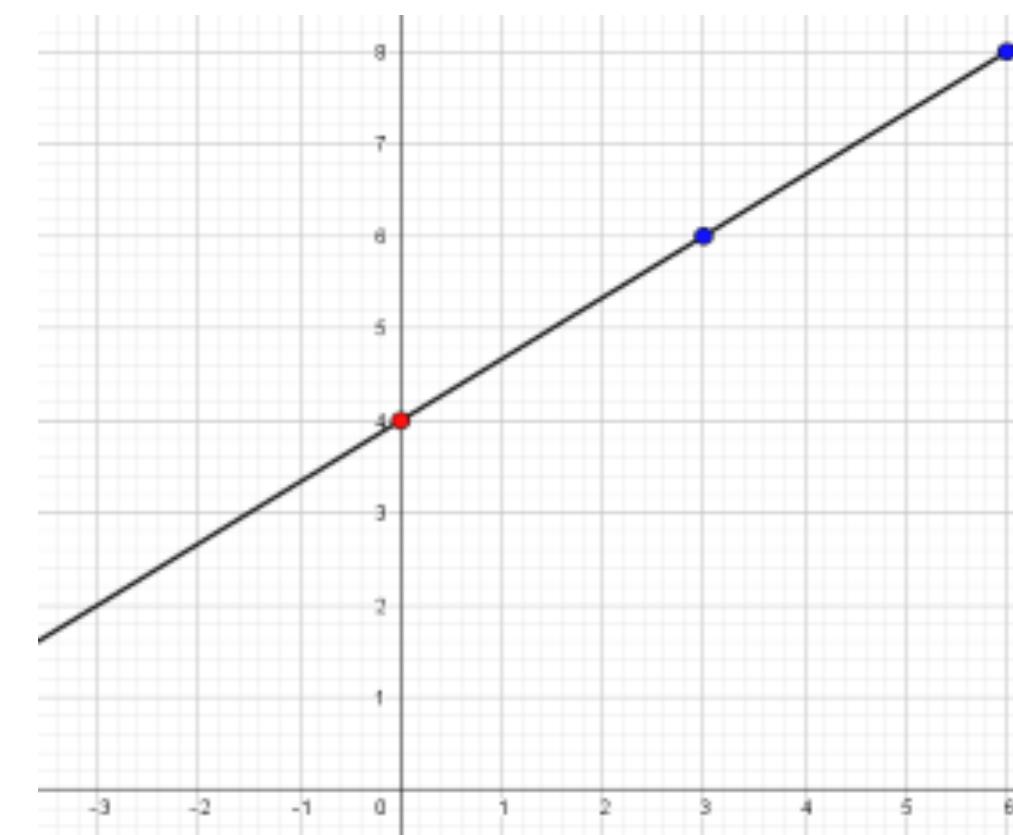


A Simple Neural Network

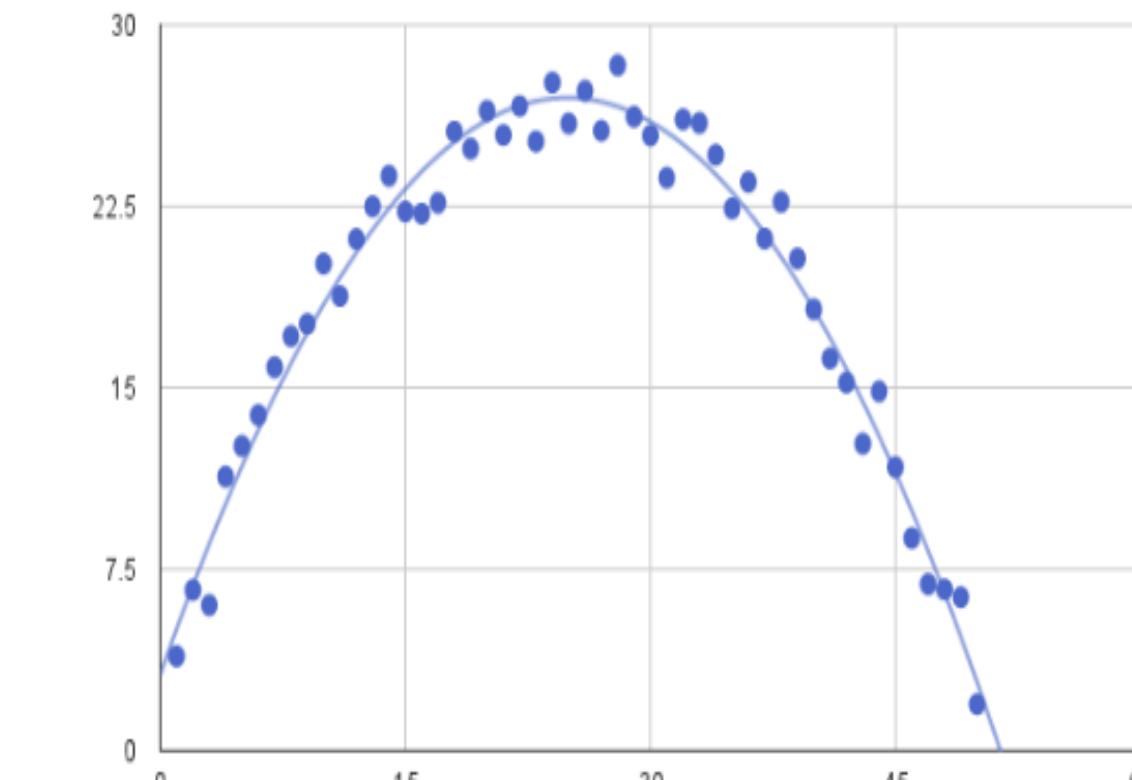
# Weights and Biases

# Universal Function Approximator

- Neural Networks are Universal Function Approximators that can approximate any mathematical function (linear or non-linear functions)



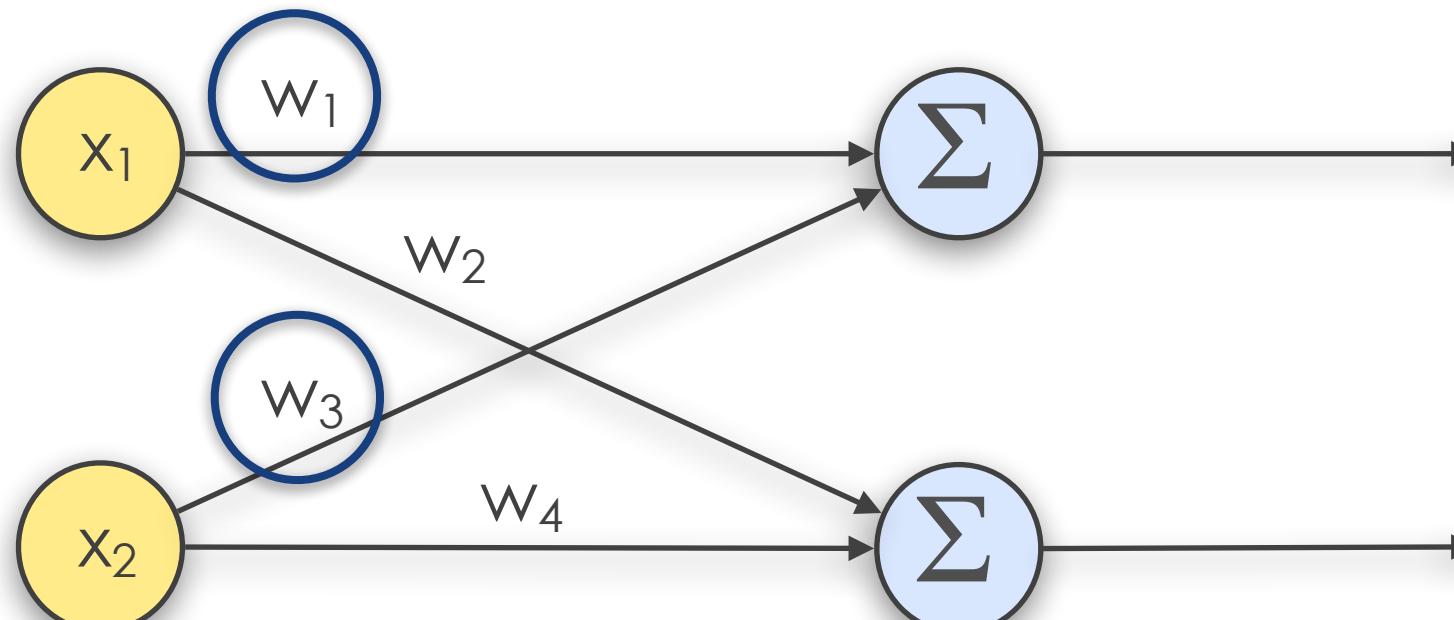
A linear equation



A non-linear equation

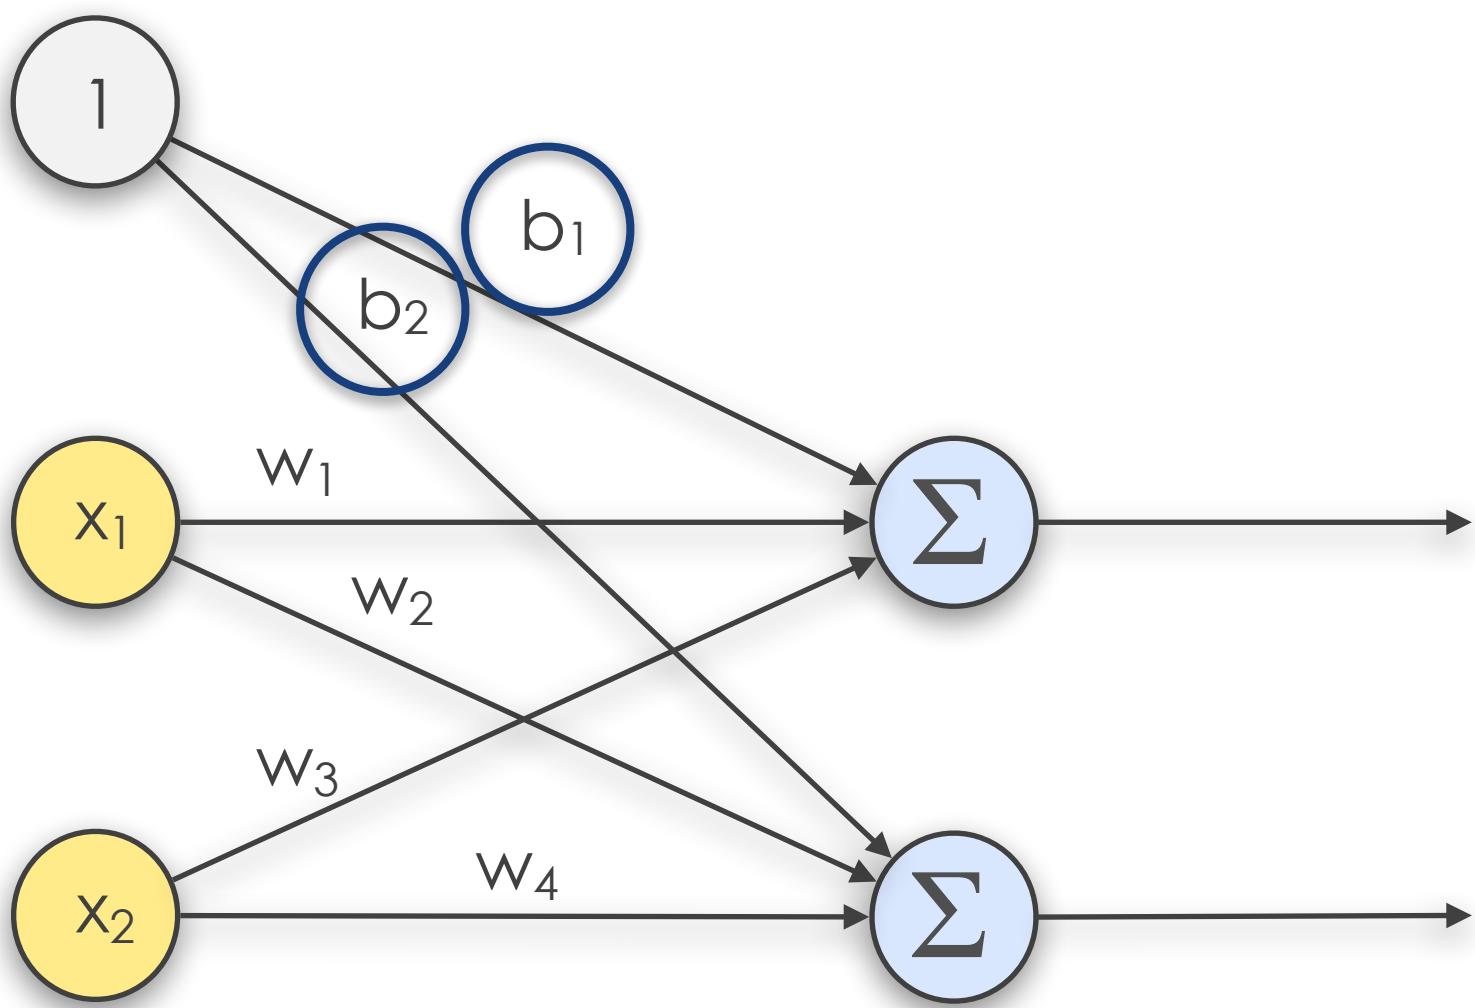
# Weights

- Weights are numerical values attached to connections between neurons, and they control the impact of incoming values to the next neuron



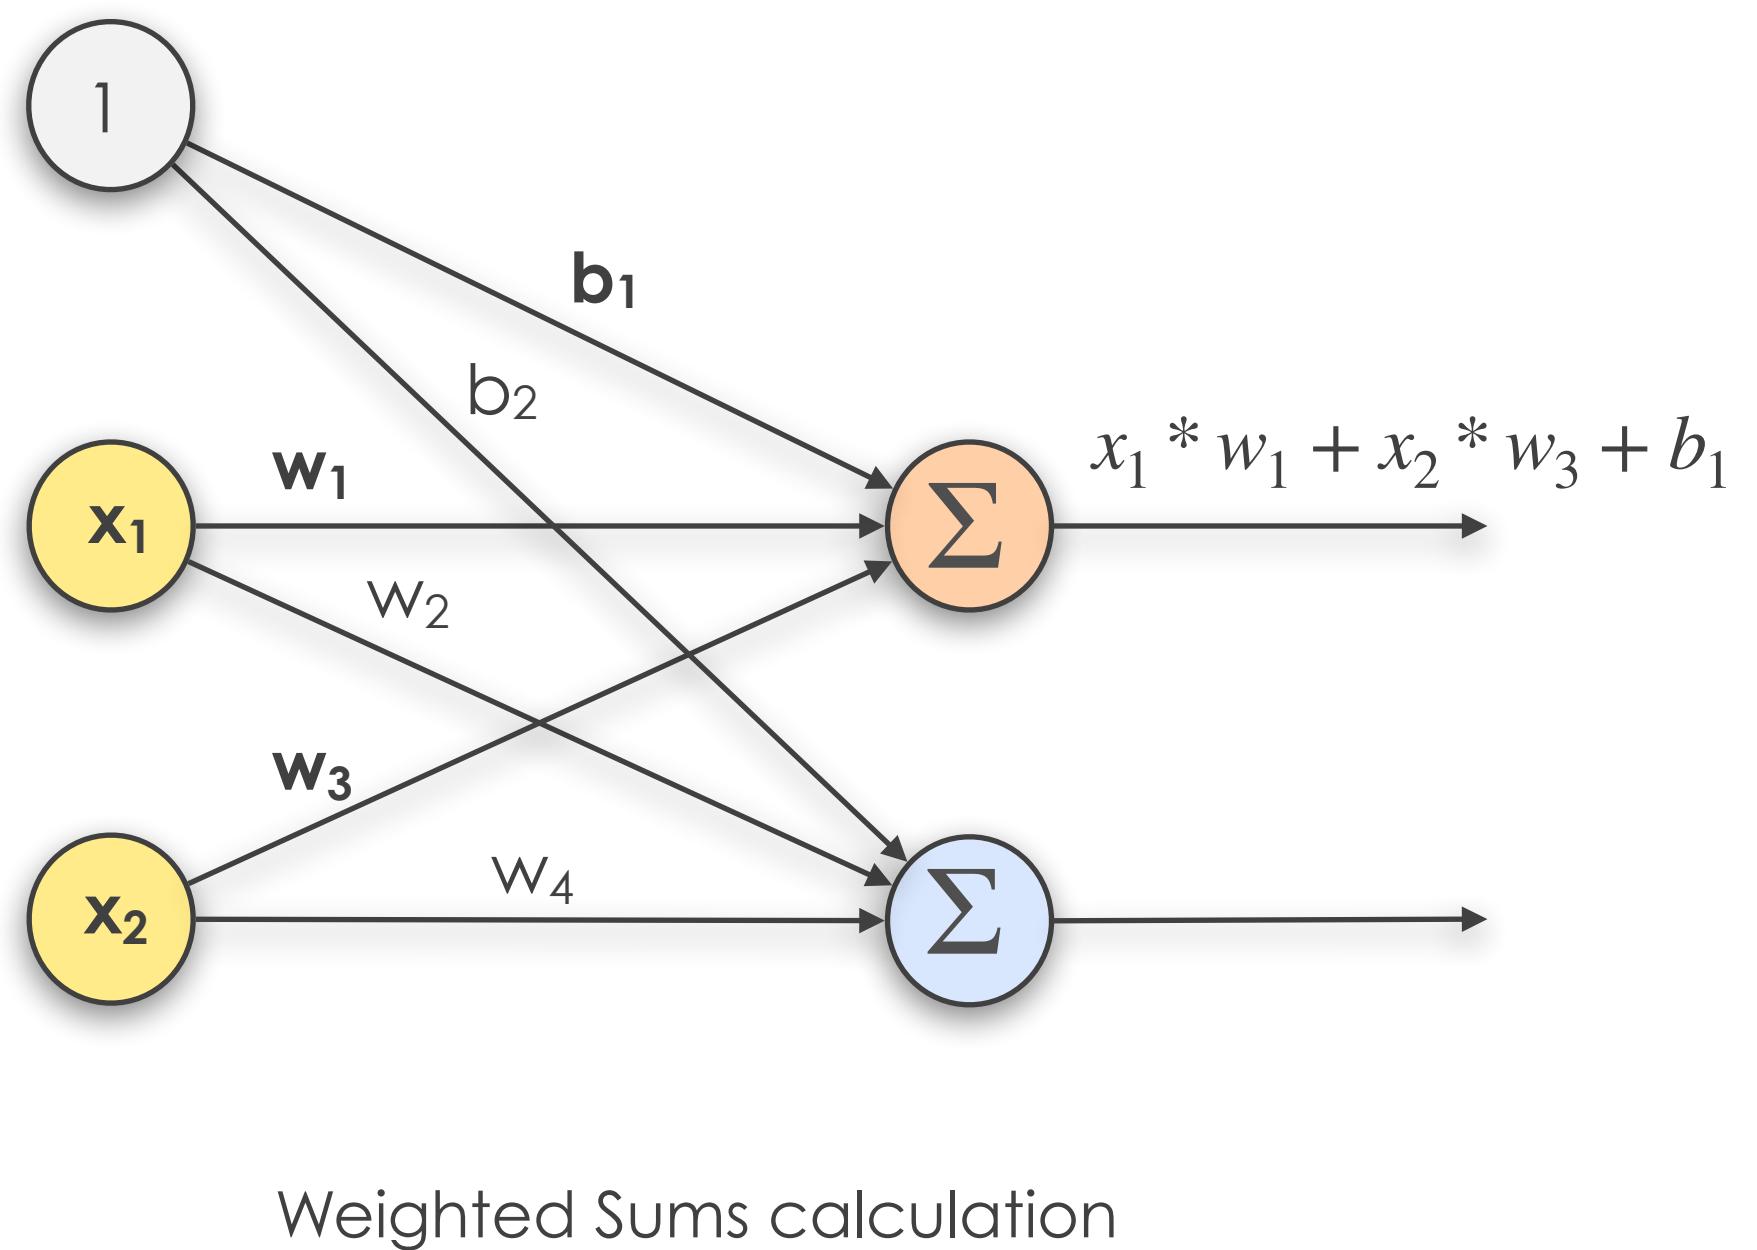
# Biases

- Biases are special weights that allow finer adjustments to computed equations



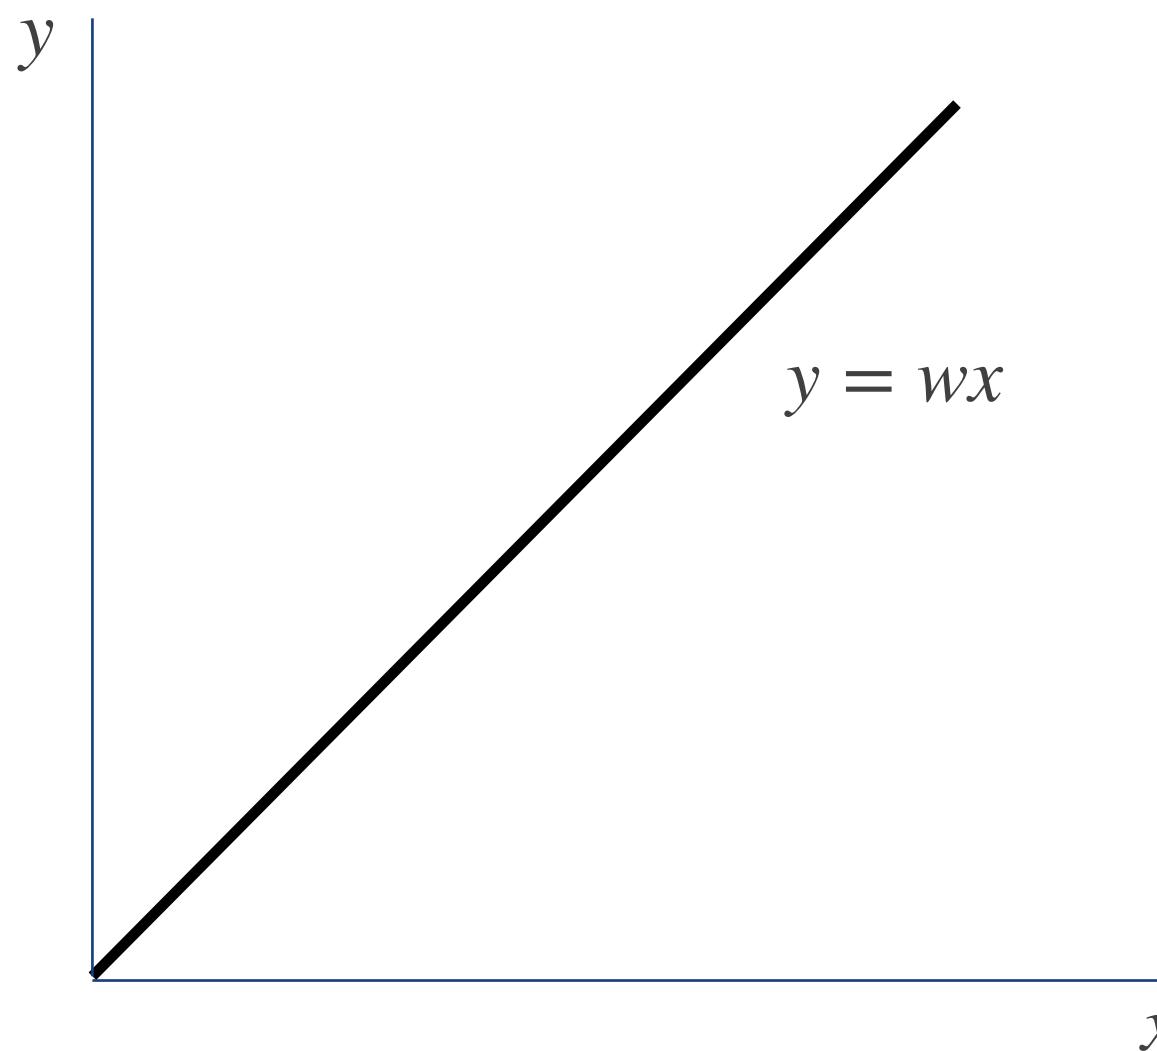
# Weighted Sum

- Weighted Sums are computed by multiplying and adding weights and biases with inputs from previous layers



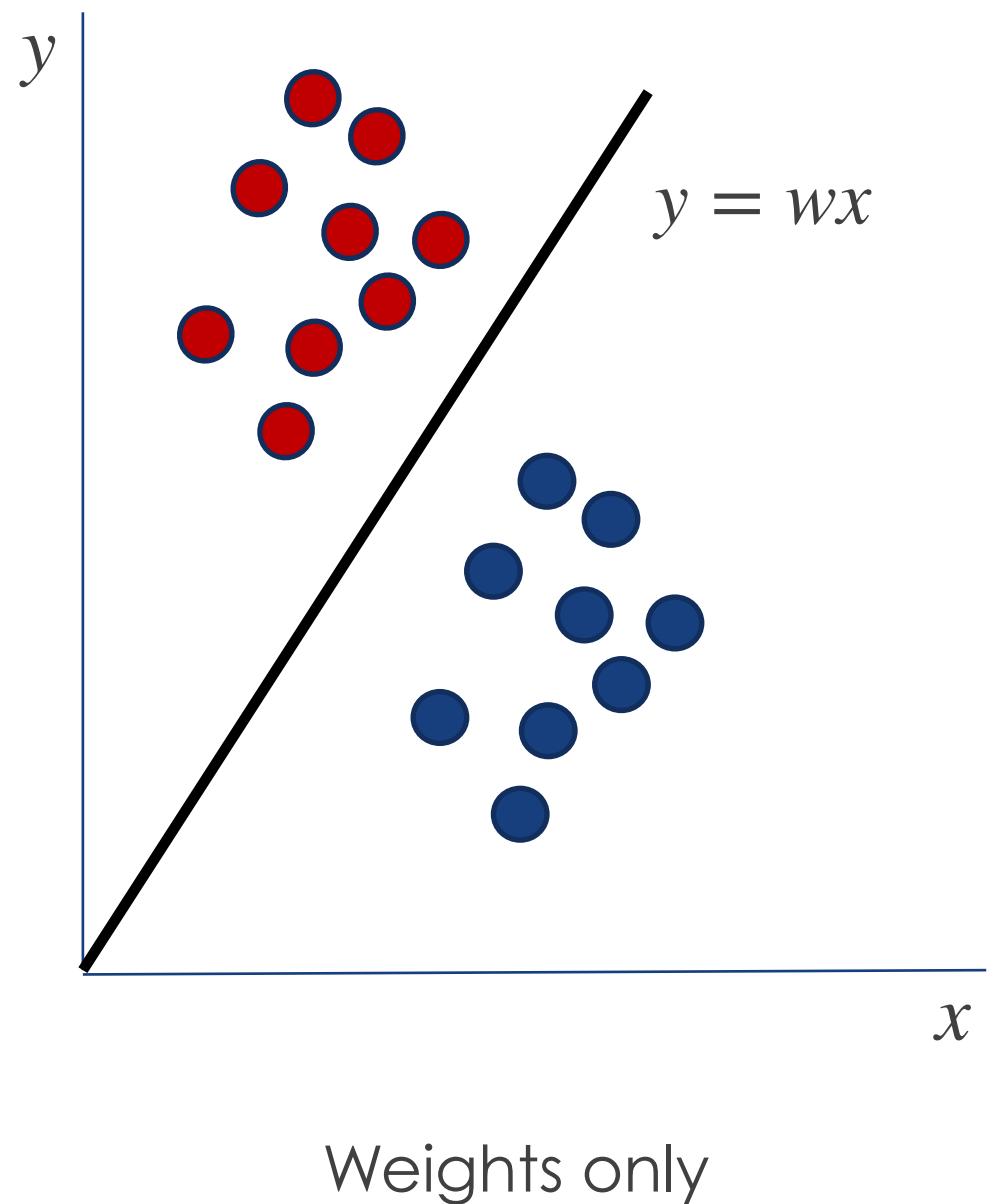
# Weights only

- Let's look at a neural network that approximates a linear function that separates a set of data
- A straight line is defined as  $y = wx$ , where  $w$  is the gradient



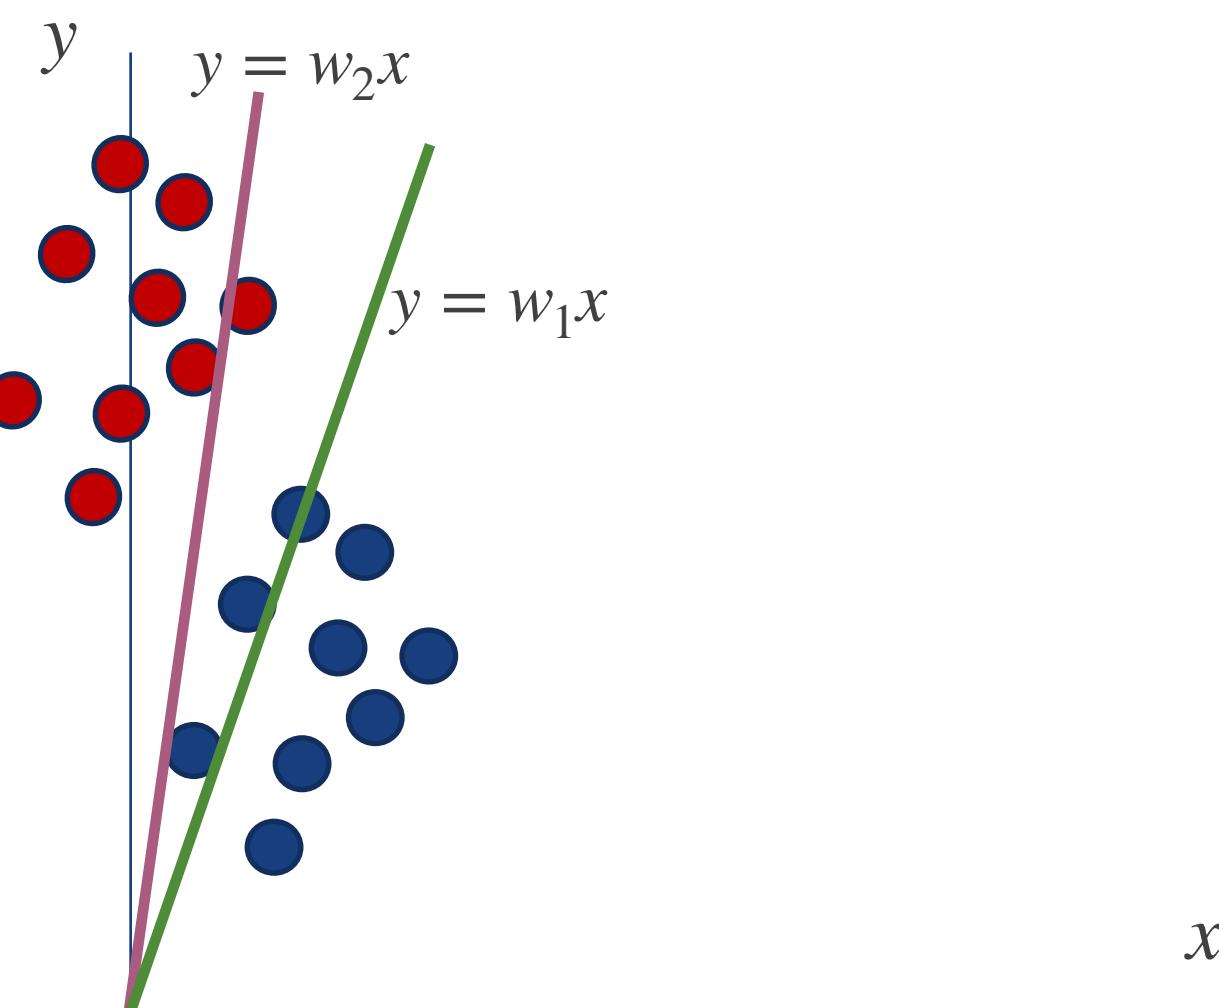
# Weights only

- We can only solve those classification problems whose data points can be separated by a line that intercepts the origin with Weights only



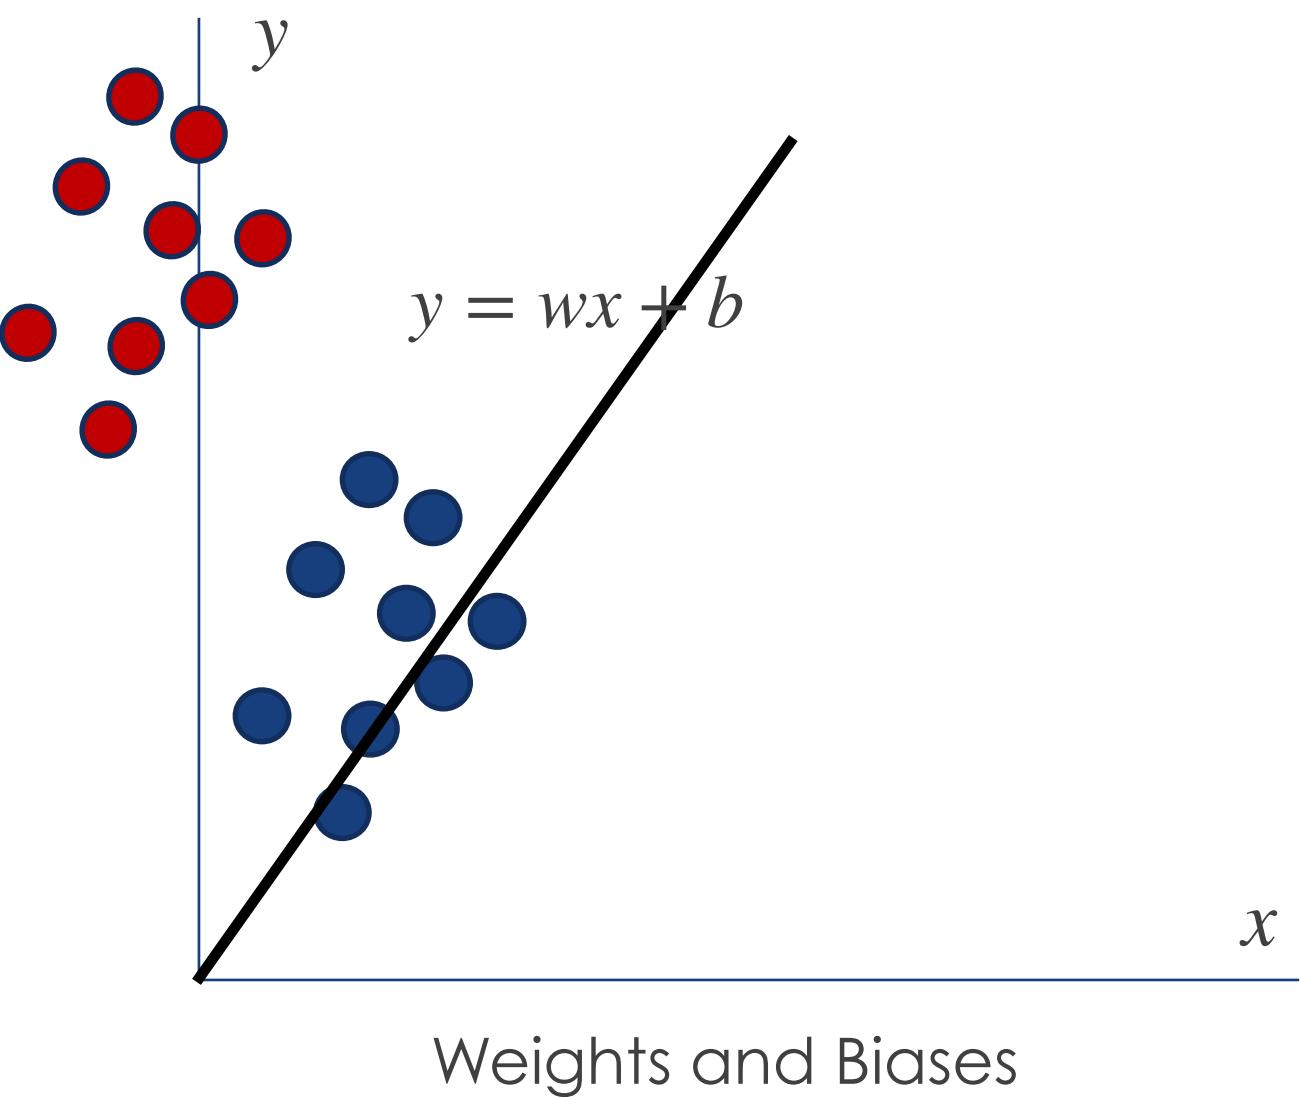
# Weights only

- These data do not allow us to separate it cleanly by adjusting the gradient alone (i.e. using only Weights in our computation)



# Weights and Biases

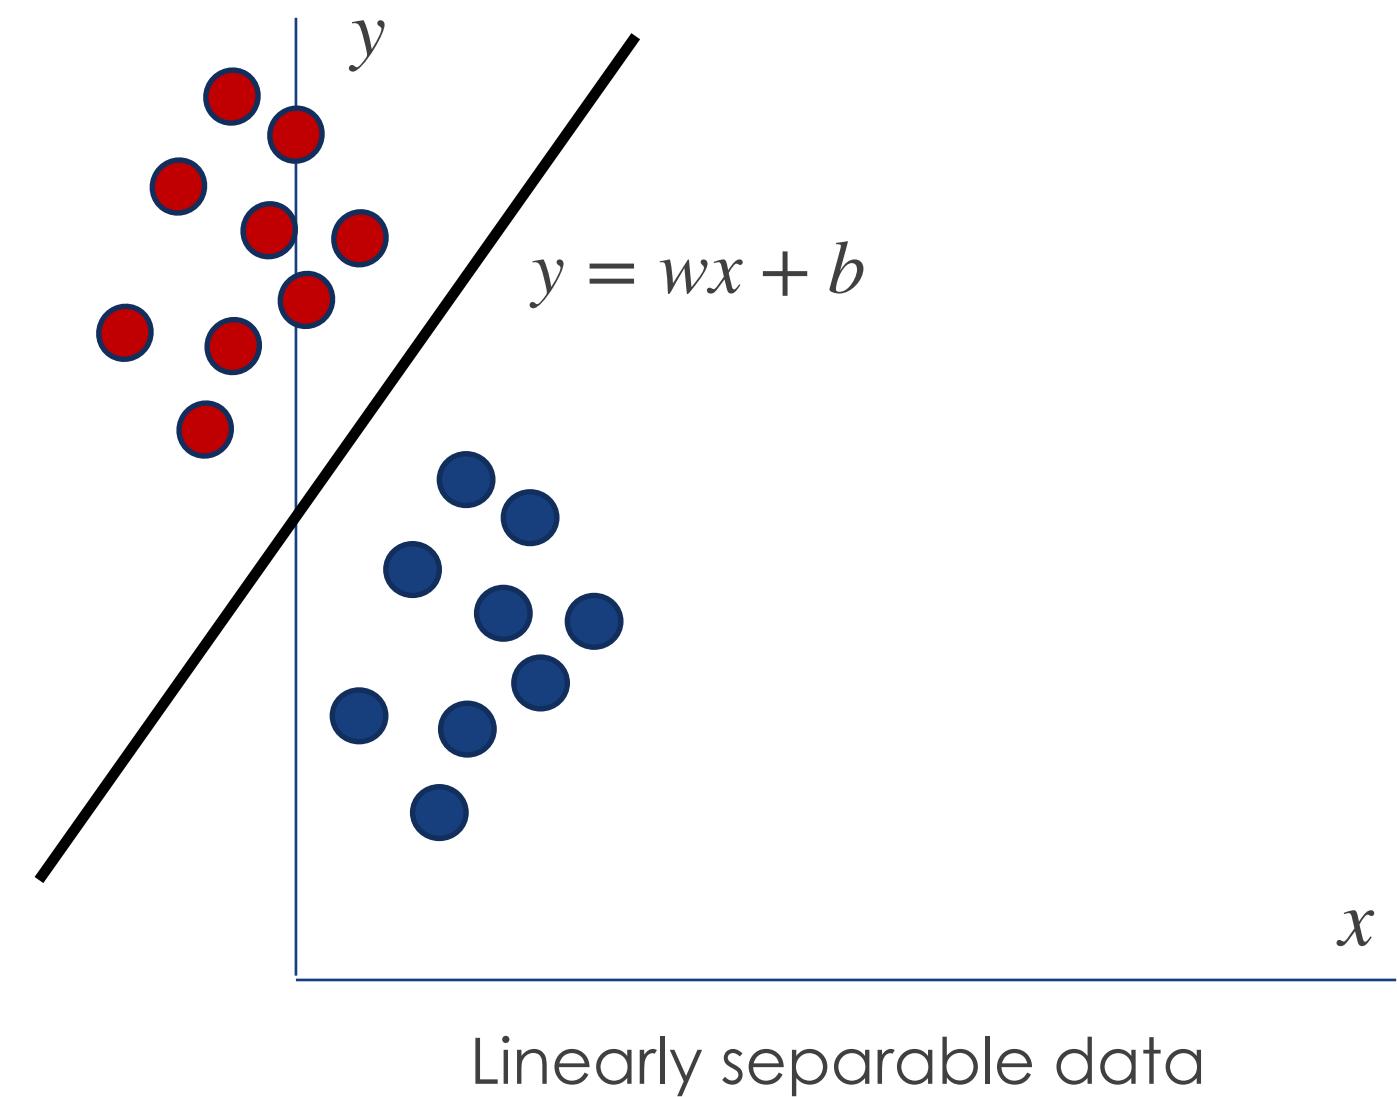
- Bias gives us the ability to move away from the origin
- The diagram shows the same problem solved by adding a bias



# Activation functions

# Weighted Sums only

- Linear solutions, such as the one provided by Weighted Sums, can solve problems that are linearly separable

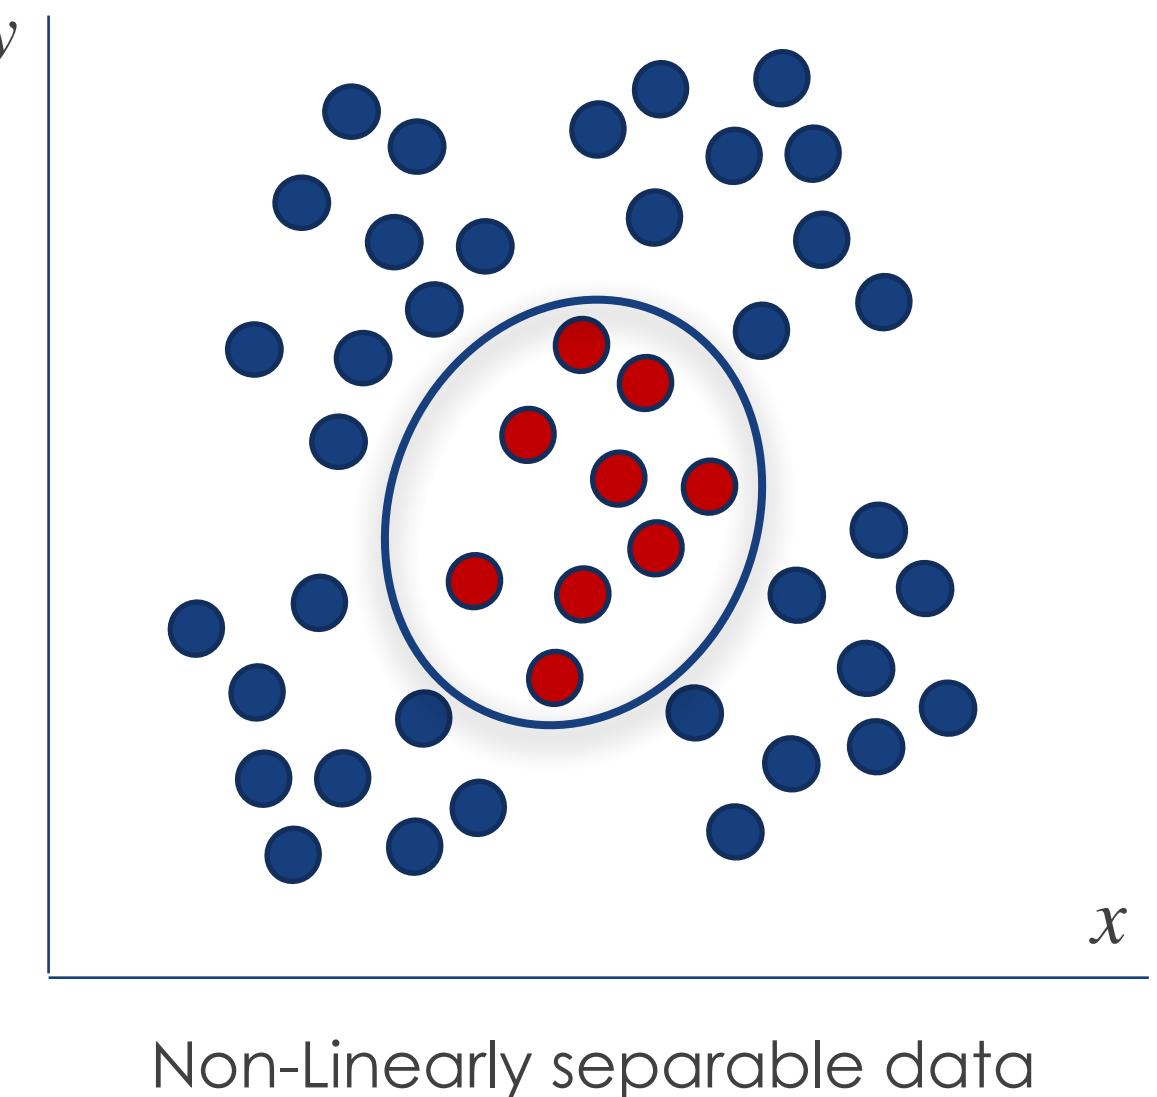


# Weighted Sums only

- A Weighted Sum operation yields a linear equation:  $y = Wx + b$
- As the output of a weighted sum operation in one layer is the input to a weighted sum operation of the next, a neural network with only weighted sum operations can only produce linear approximations
- Example:
$$Y_1 = 2X + 1$$
$$Y_2 = 3Y_1 + 4$$
$$Y_2 = 3(2X + 1) + 4$$
$$Y_2 = 6X + 7 \text{ (still a linear equation)}$$

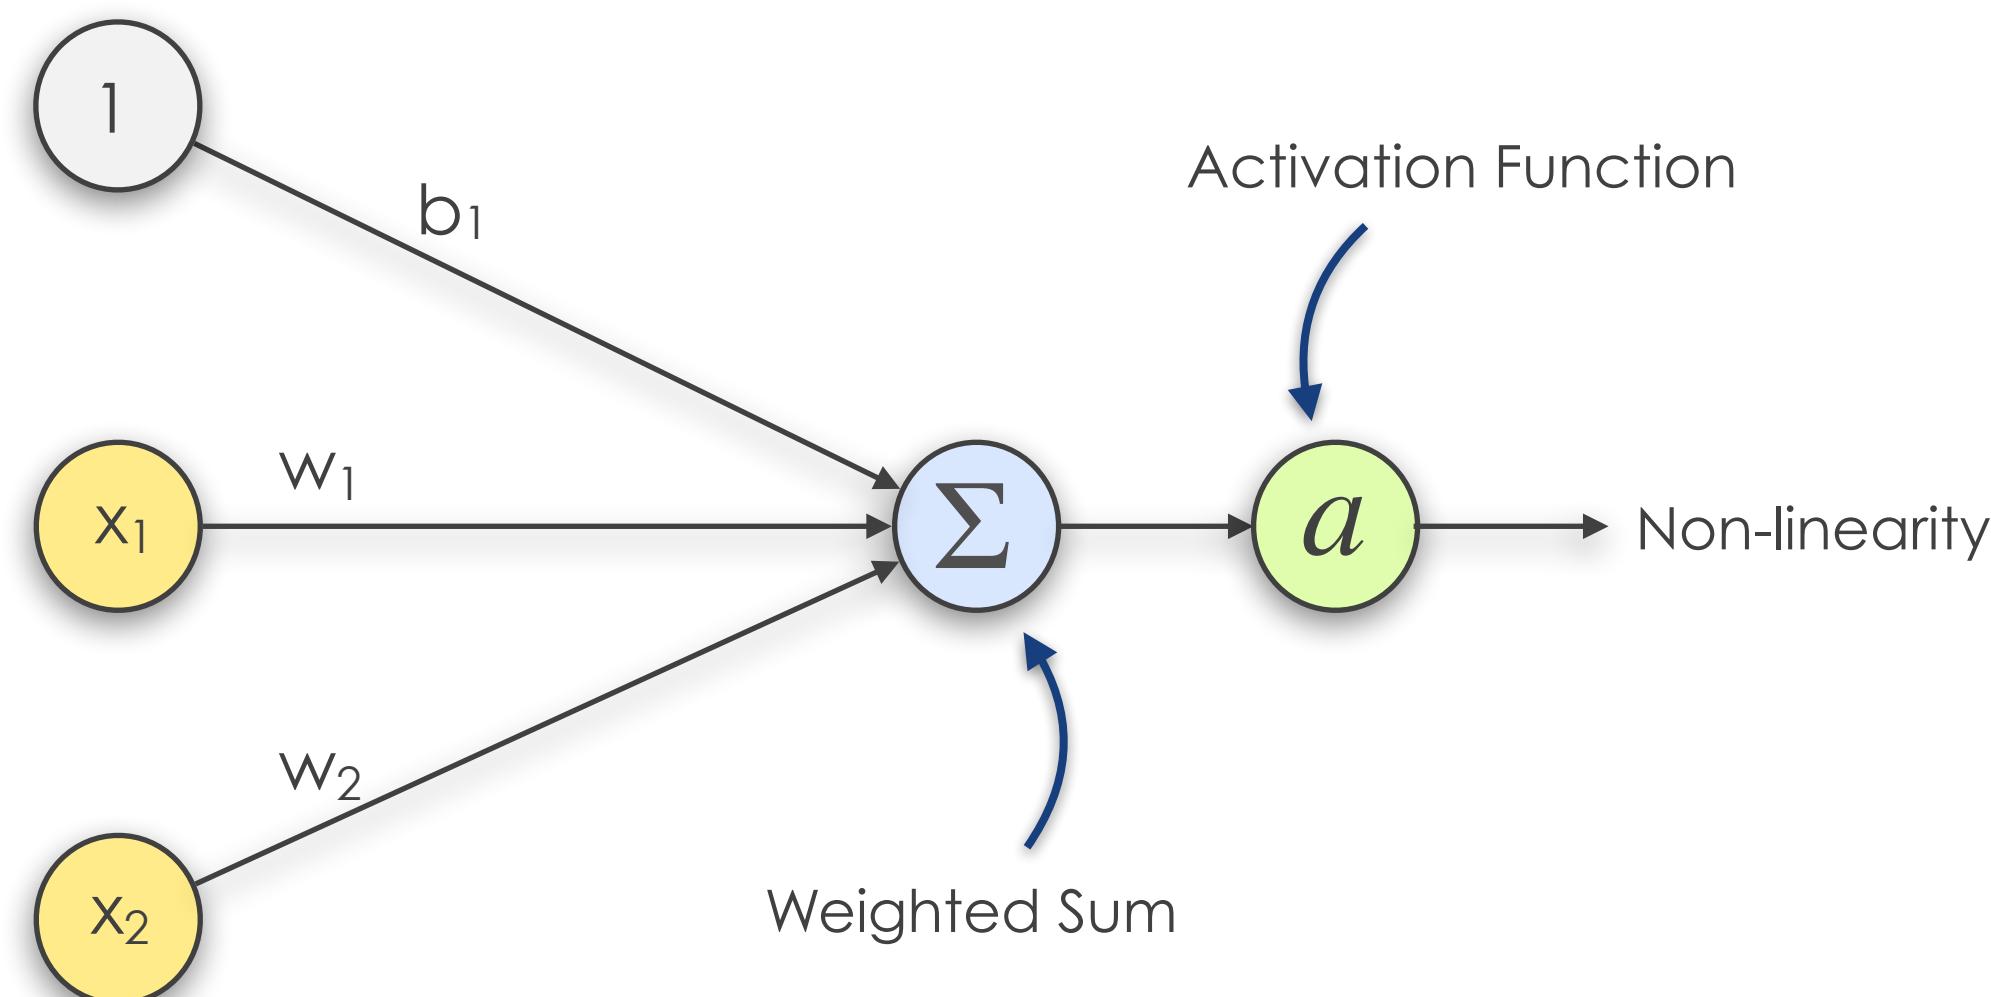
# Weighted Sums only

- Data that are non-linearly separable requires the use of non-linear equations
- We shall see how Activation Functions add non-linearity to our models to solve problems like such



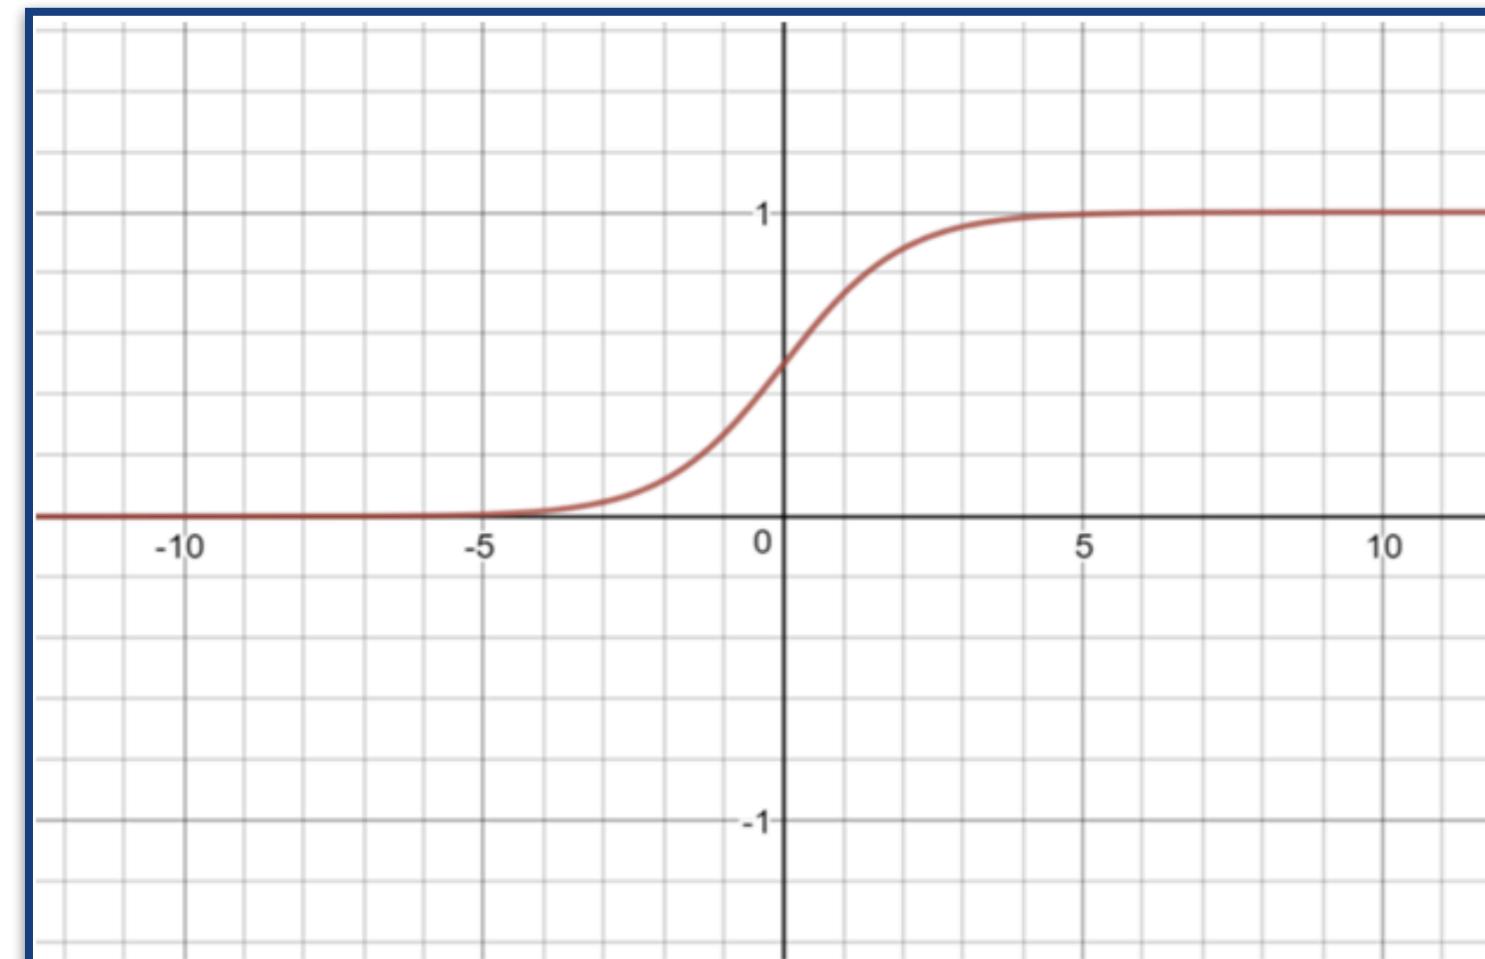
# Non-linear transformation

- A Weighted Sum operation, followed by a Non-Linear Transformation applied by an Activation Function



# Sigmoid

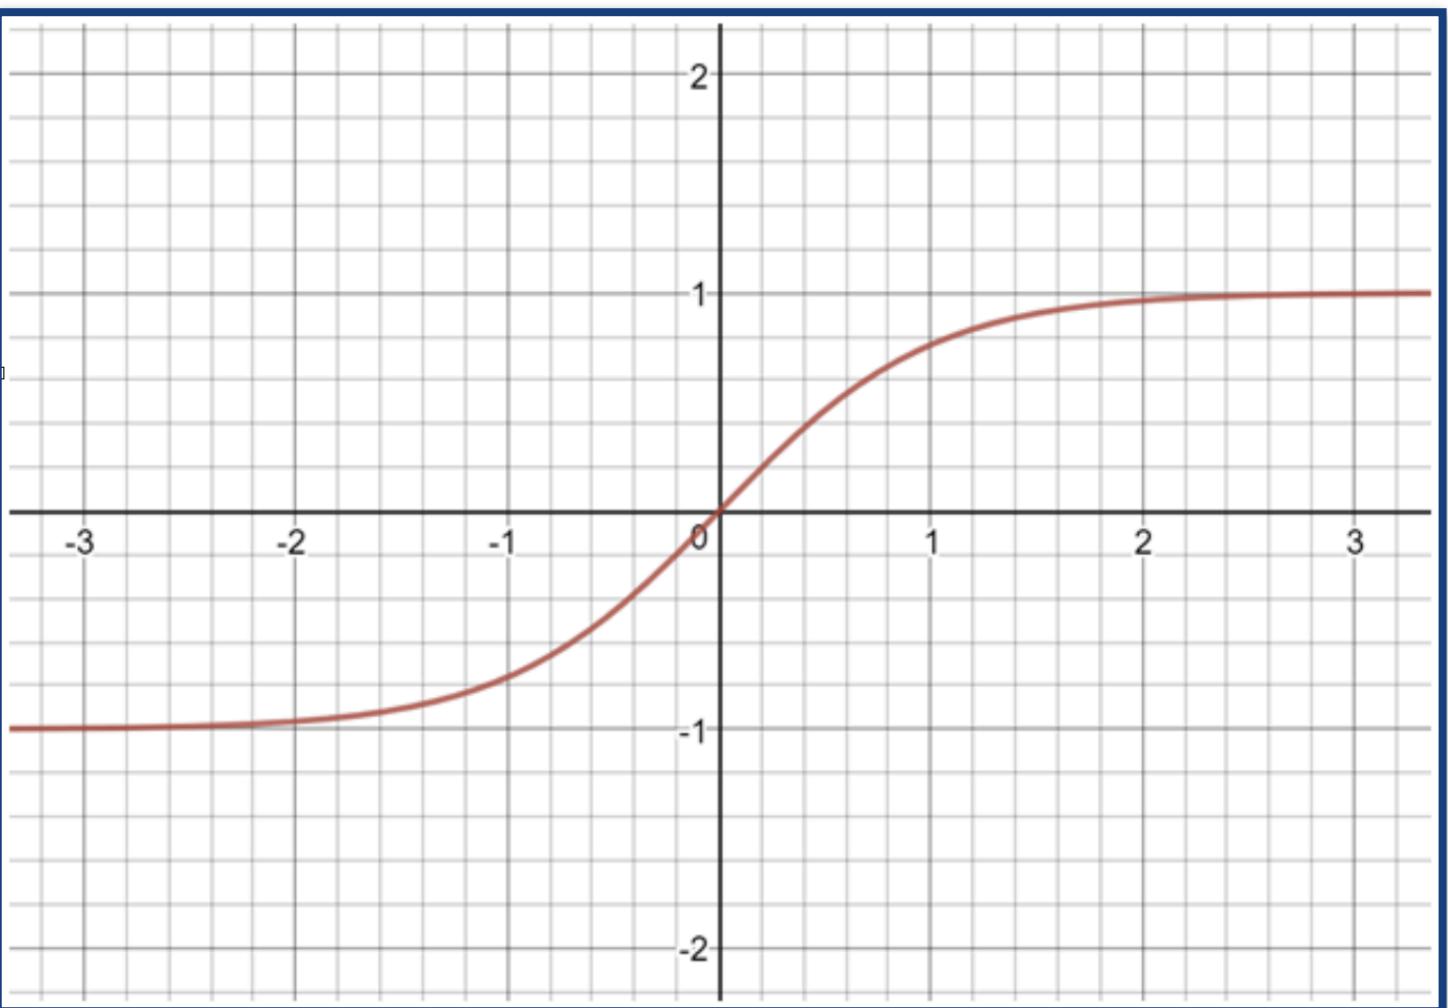
- The Sigmoid activation function, also called the Logistic function, transforms inputs into a value between 0 and 1
- Useful for models that predict a probability as an output



$$f(x) = \frac{1}{1 + e^{(-x)}}$$

# Tanh

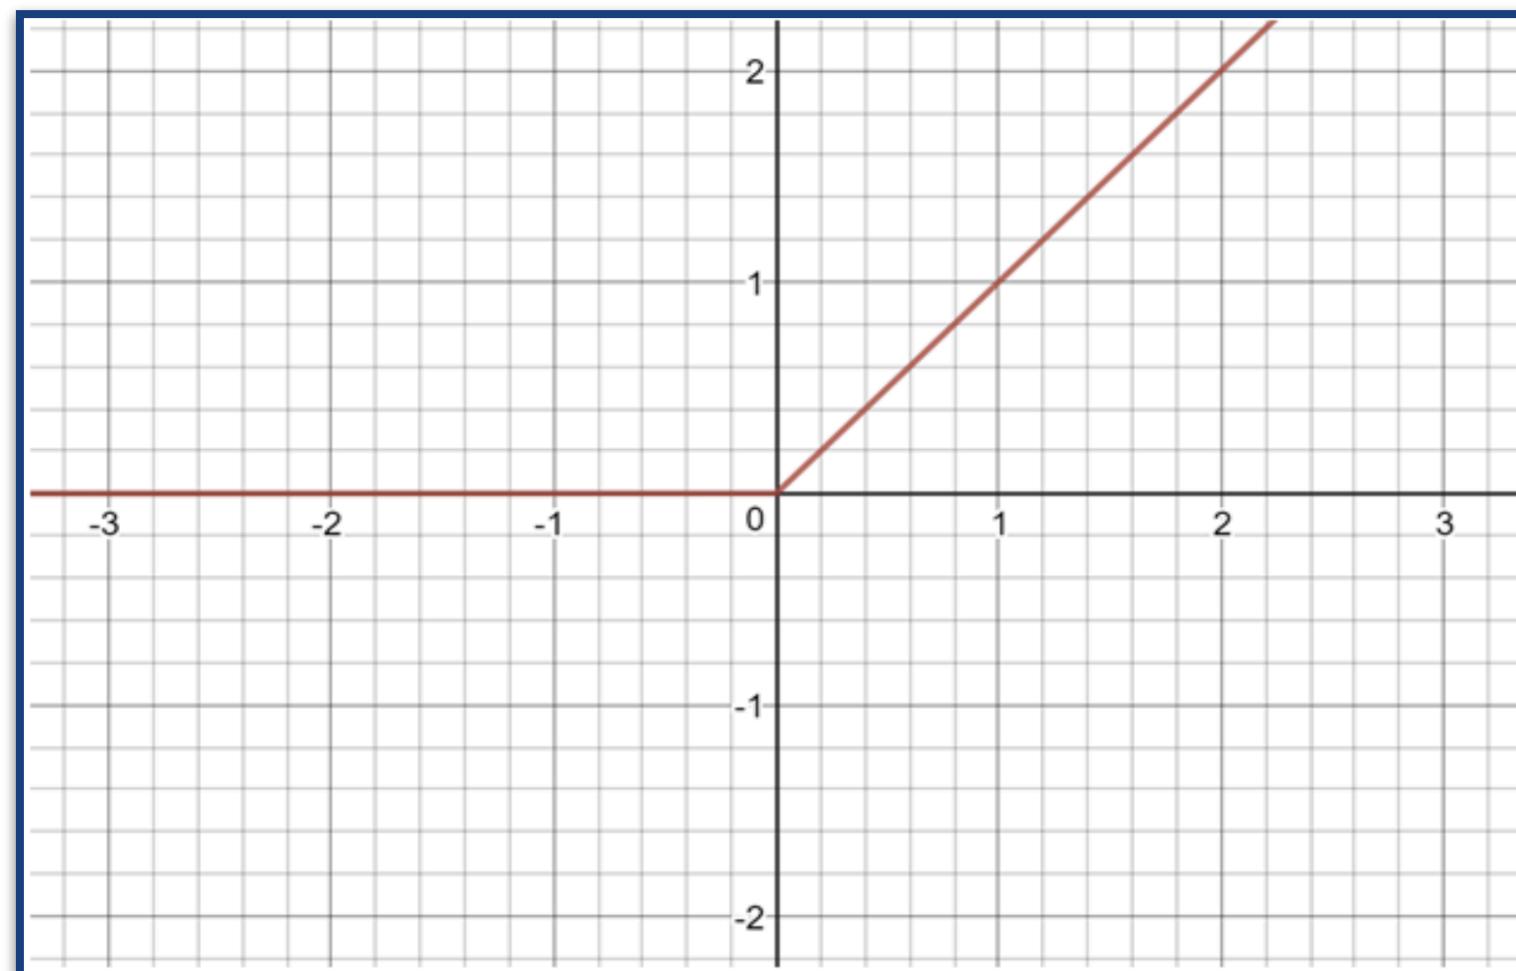
- The Tanh activation function transforms inputs to values between -1 and 1
- Useful for models that yields normalized positive and negative outputs



$$f(x) = \frac{2}{1 + e^{-2x}} - 1$$

# ReLU

- The ReLU activation function transforms inputs with negative values to 0, but leave other values unchanged



$$f(x) = \begin{cases} 0, & \text{if } x < 0 \\ x, & \text{if } x \geq 0 \end{cases}$$

# Softmax

- The Softmax activation function takes a final computed vector and outputs a vector of probabilities as predictions

$$\sigma(\vec{z})_i = \frac{e^{z_i}}{\sum_{j=1}^N e^{z_j}}$$

Diagram illustrating the Softmax function:

- a vector of logits
- softmax function
- compute probability for  $i^{th}$  element of the logits vector
- value of  $i^{th}$  element of the logits vector
- value of every element in the logits vector
- no. of classes the model is trained to predict

# Loss Functions

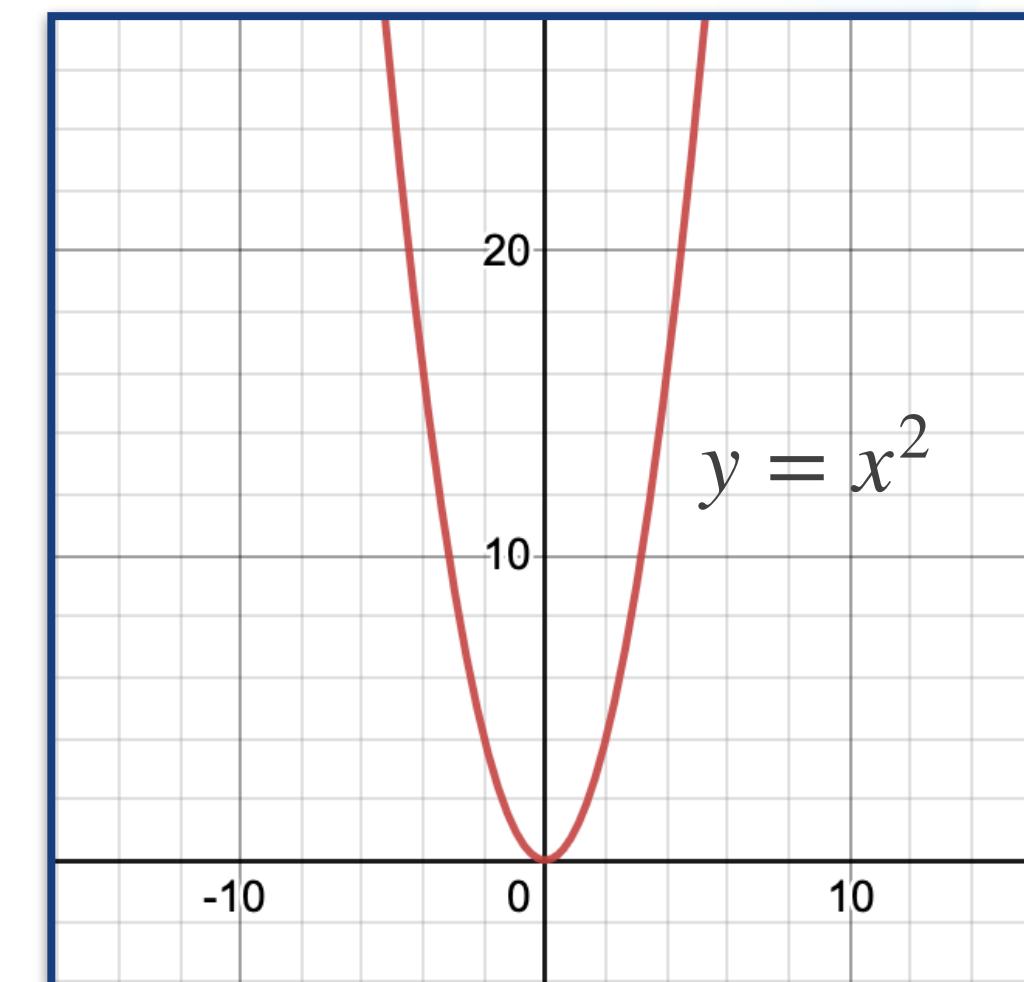
- When training, a loss function helps to determine the error of our neural network
- The error tell us how far our predicted value is from the actual value
- We want to minimize the error in our neural network such that its predictions are accurate
- Loss functions are only used in model training

# Loss Function

- Mean Square Error (aka L<sub>2</sub> loss) is a loss function that is commonly used

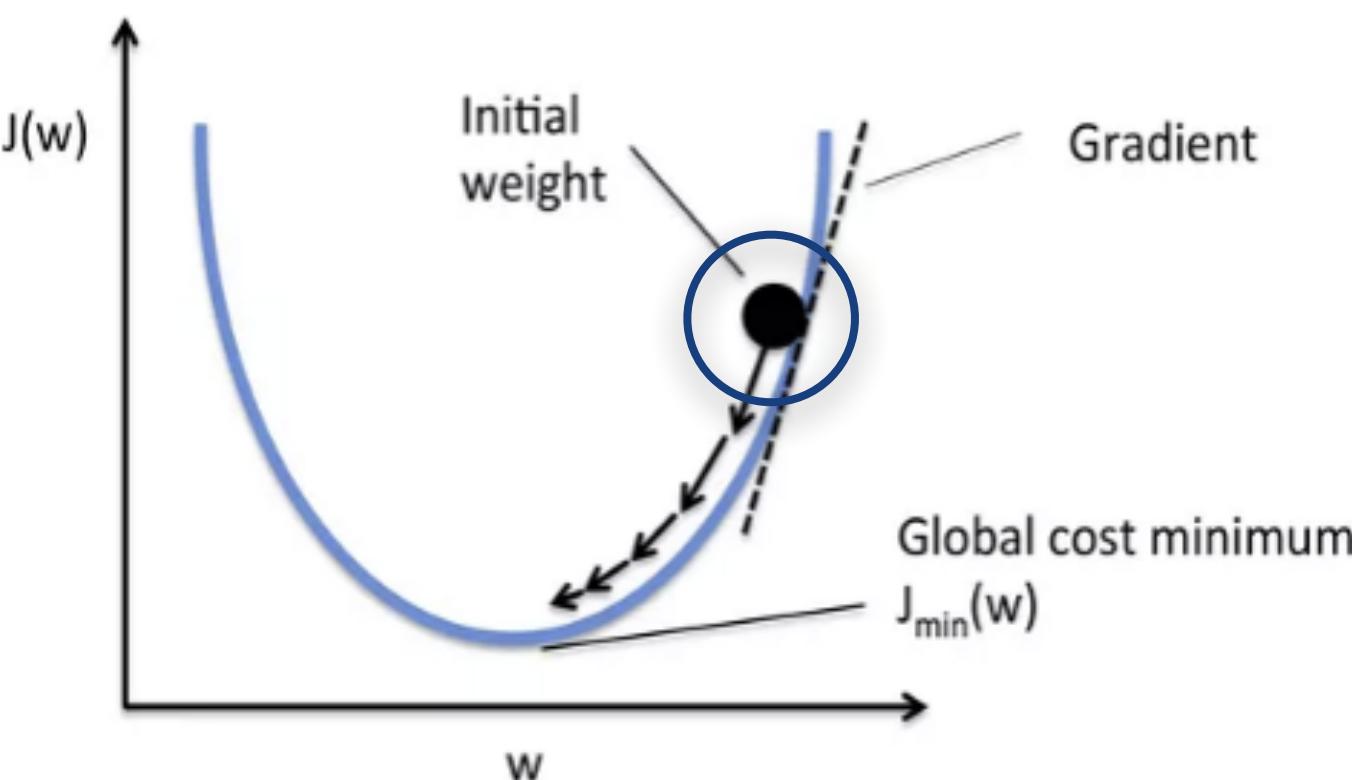
$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

Actual Value (Ground Truth)
Batch Size
Predicted Value



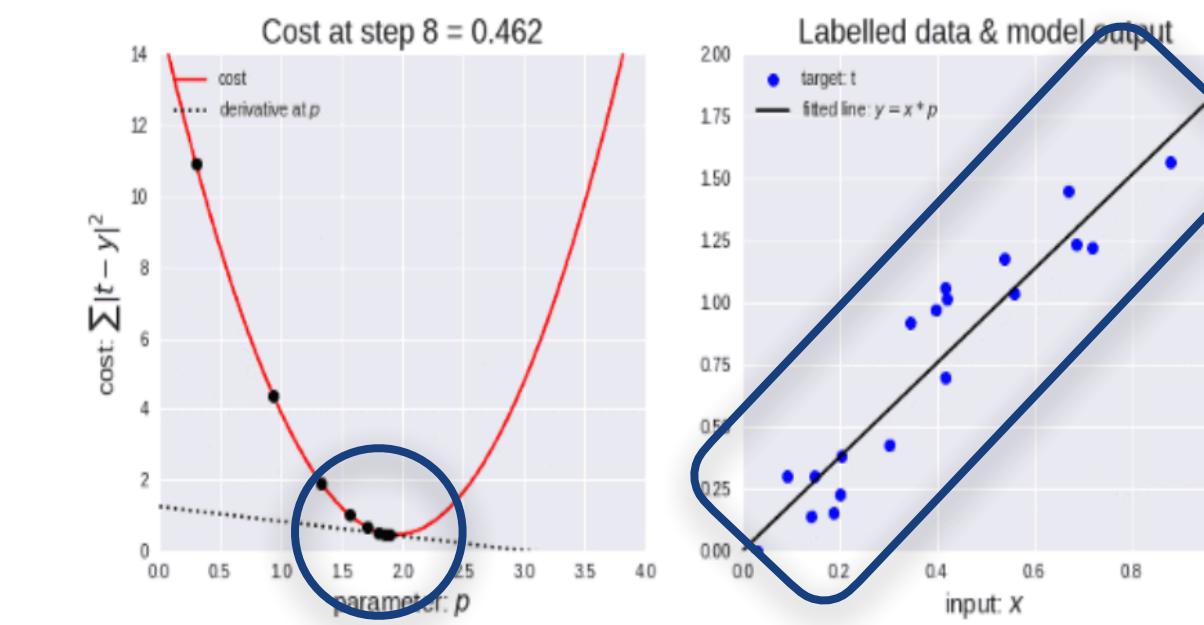
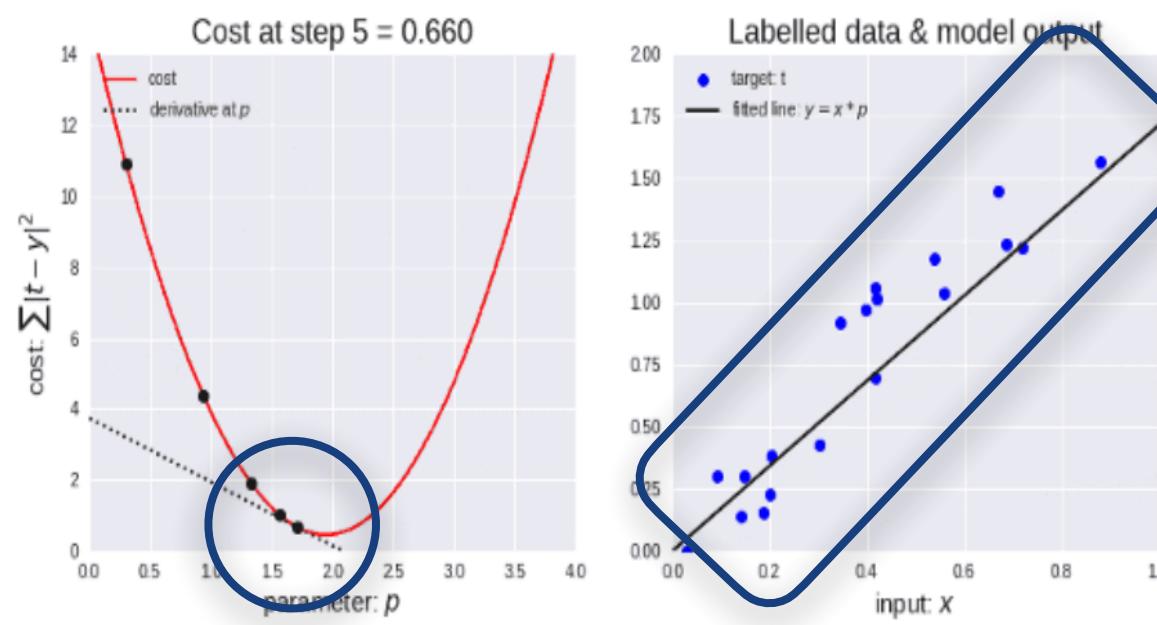
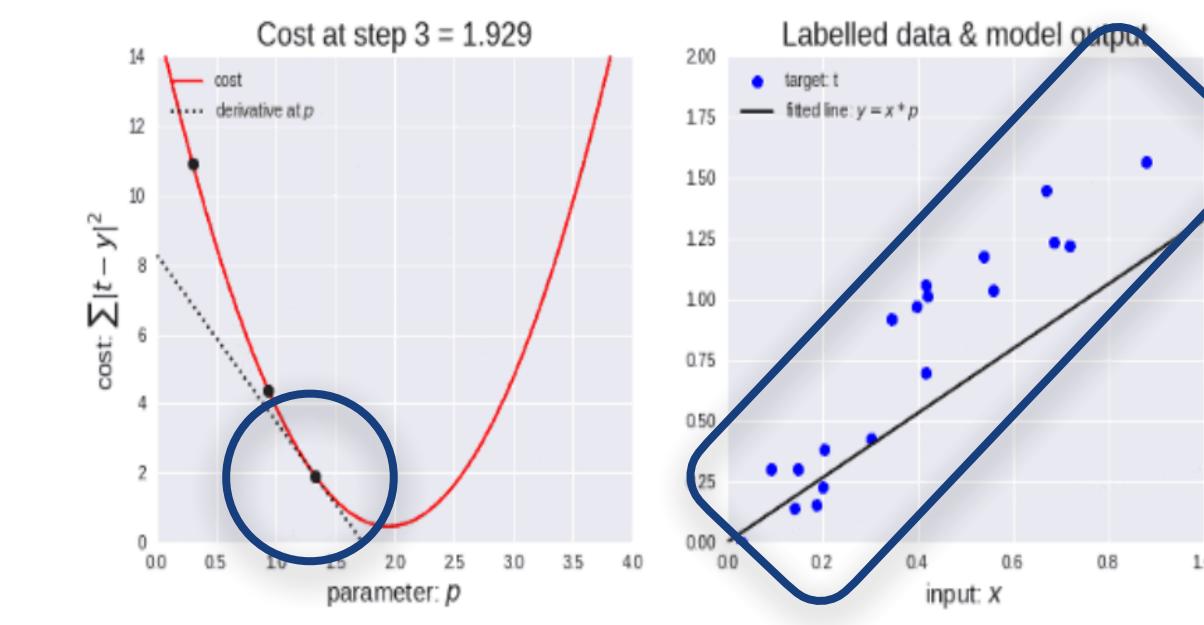
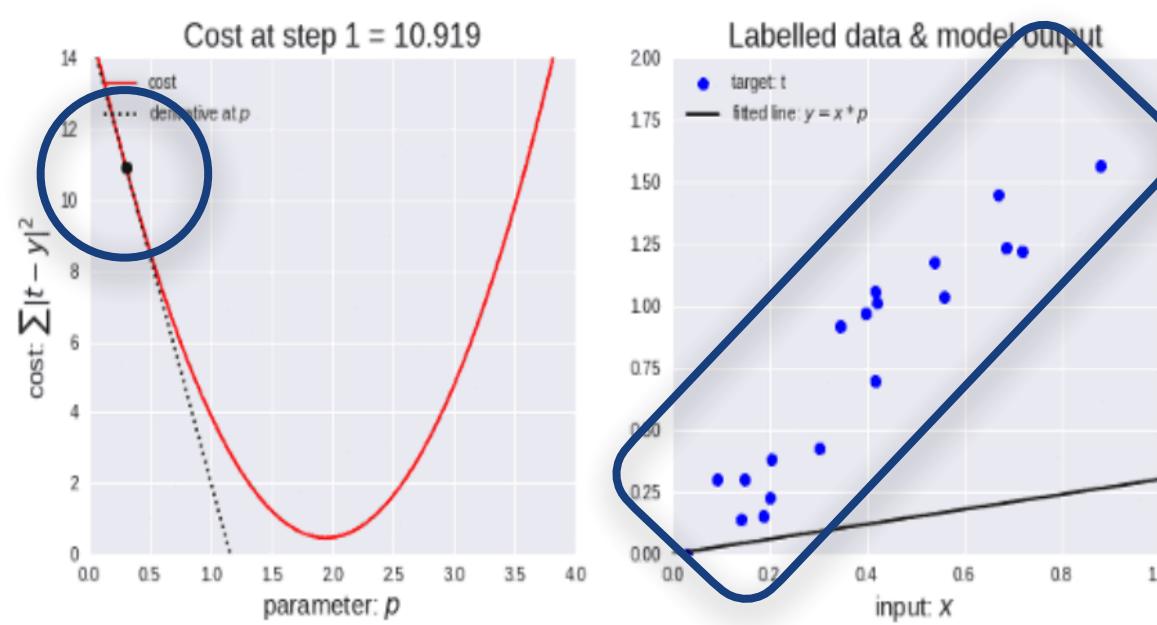
# Gradient Descent

- Gradient Descent is an optimization technique to minimize a loss function by iteratively moving in the direction of the loss function's gradient
- It determines the amount of changes to be made to the Weights and Biases for a better prediction (i.e. smaller error) in the next iteration



# Loss Function

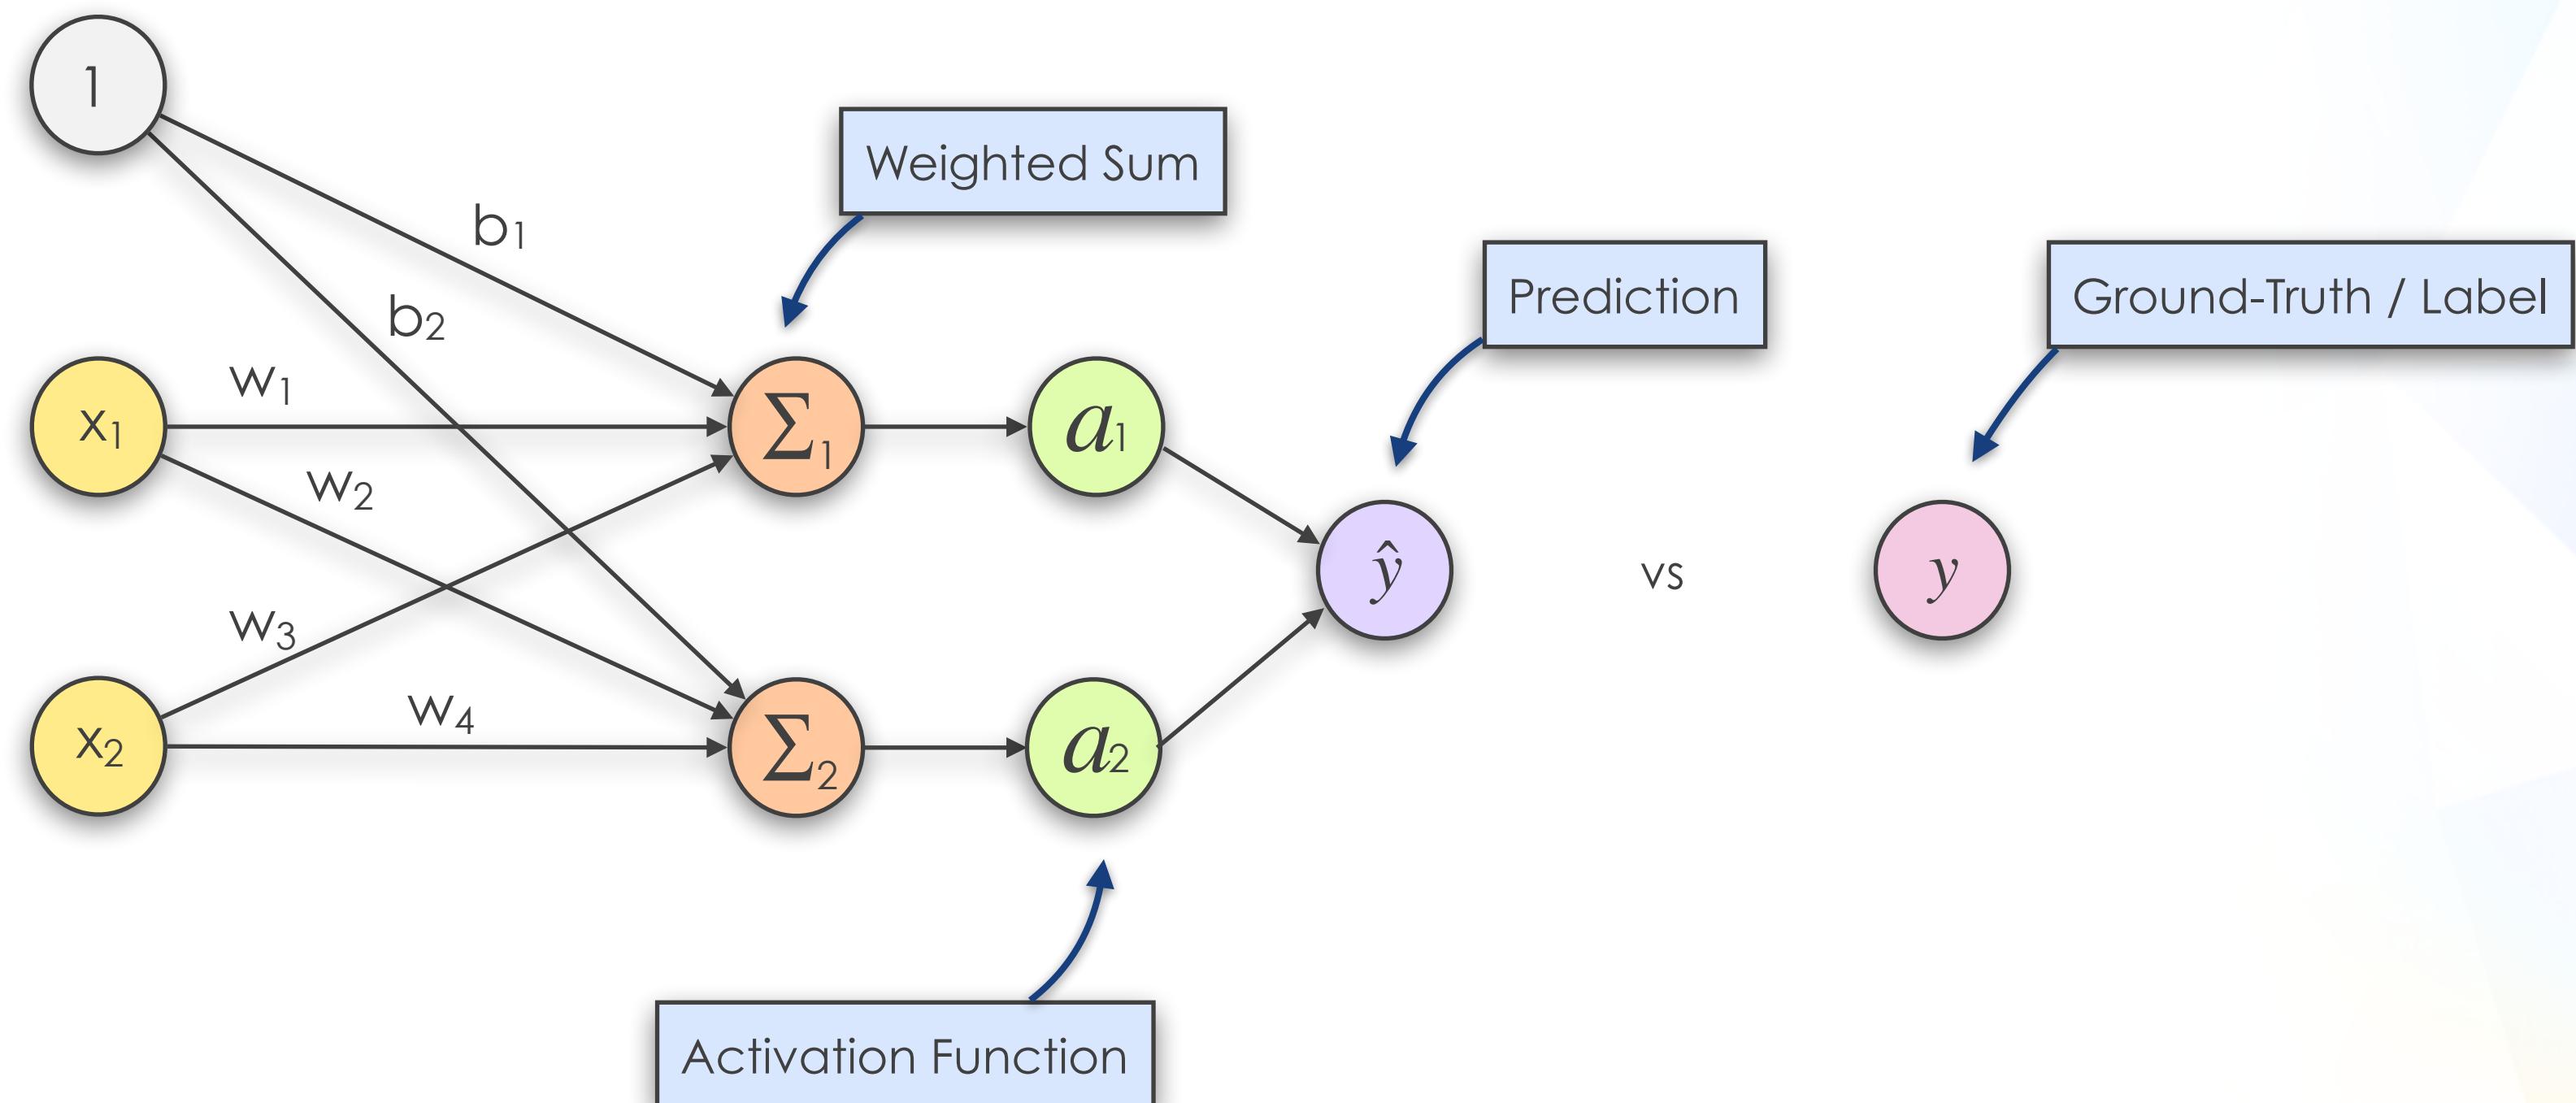
- An illustration to find a line that best fits the set of points by minimizing the MSE loss function



# Backpropagation

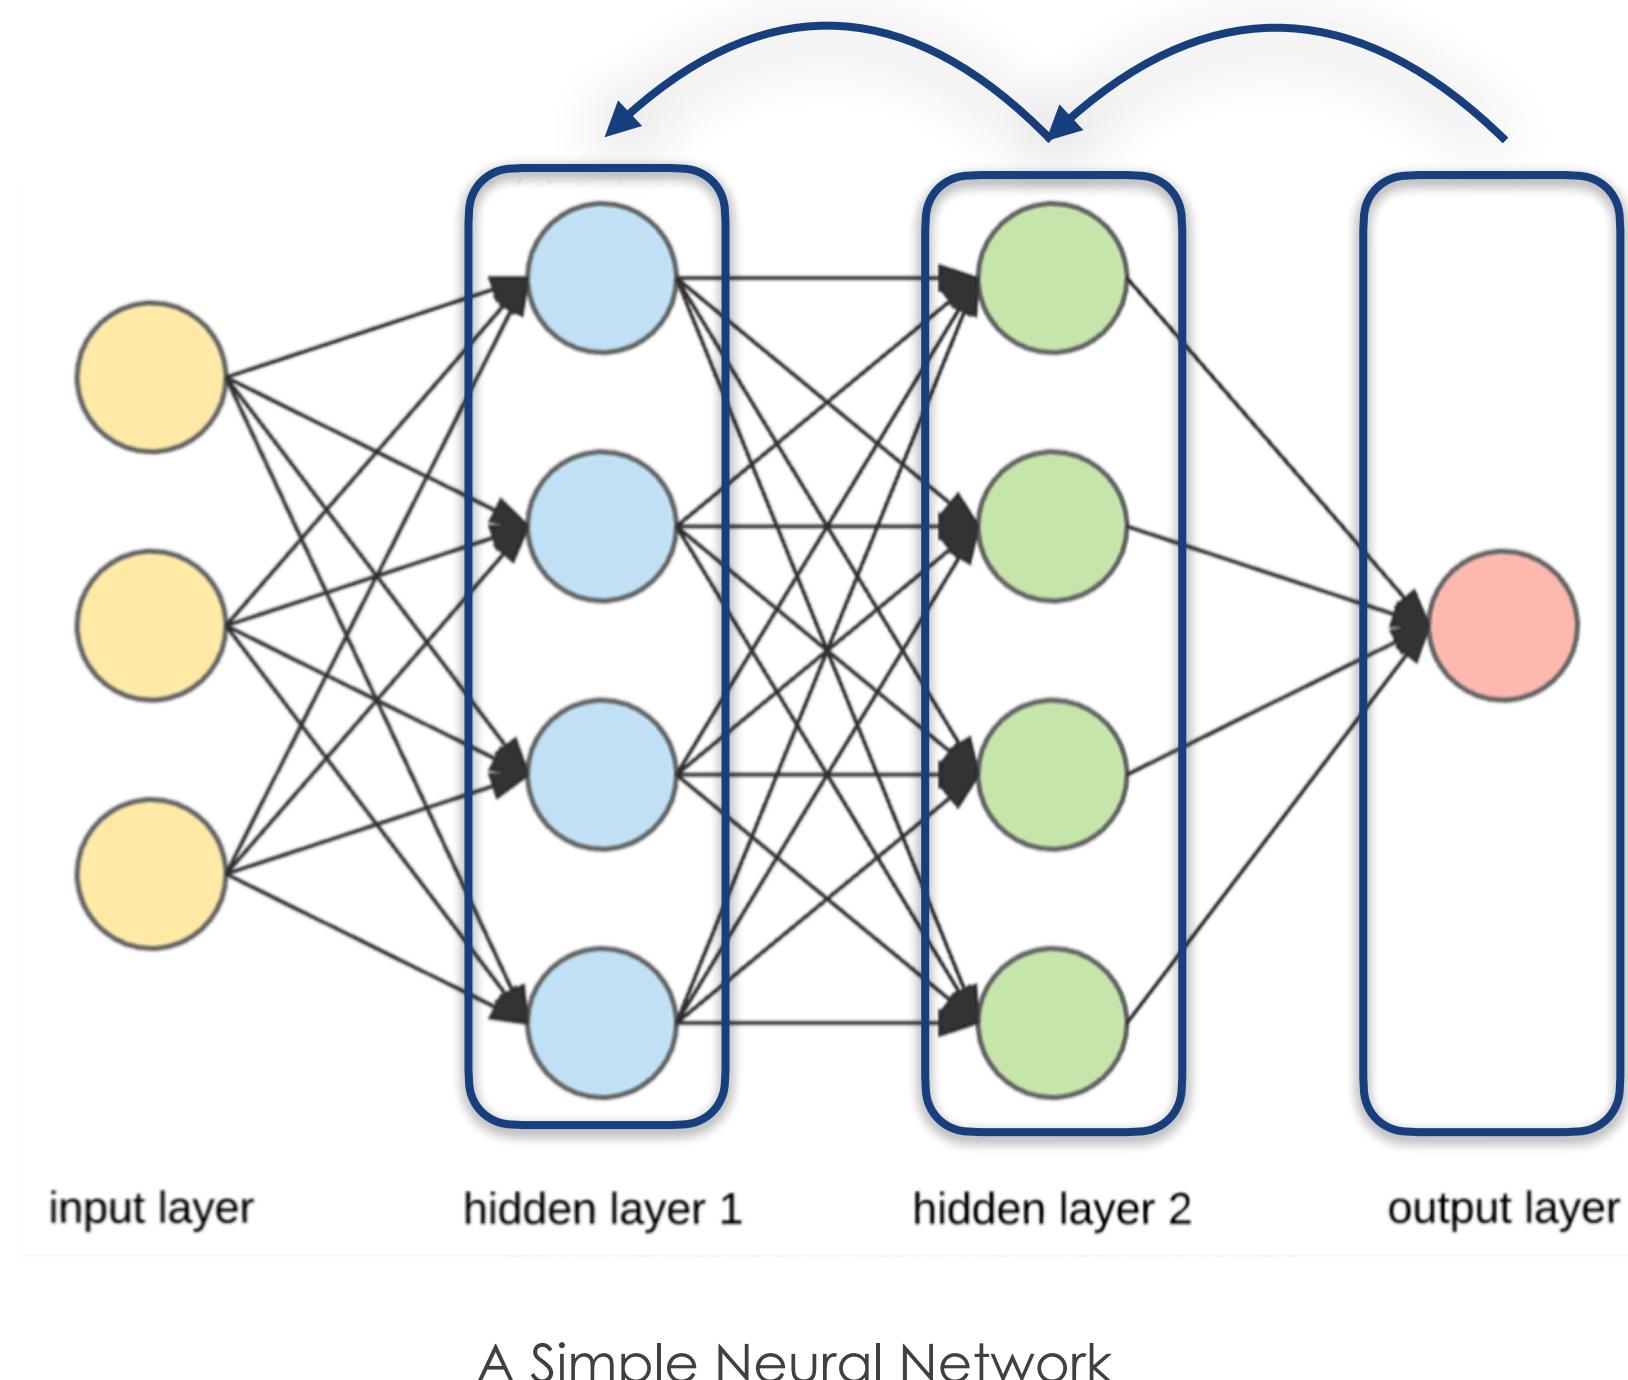
# Error

- The difference between the predicted value and the ground-truth is known as the Error



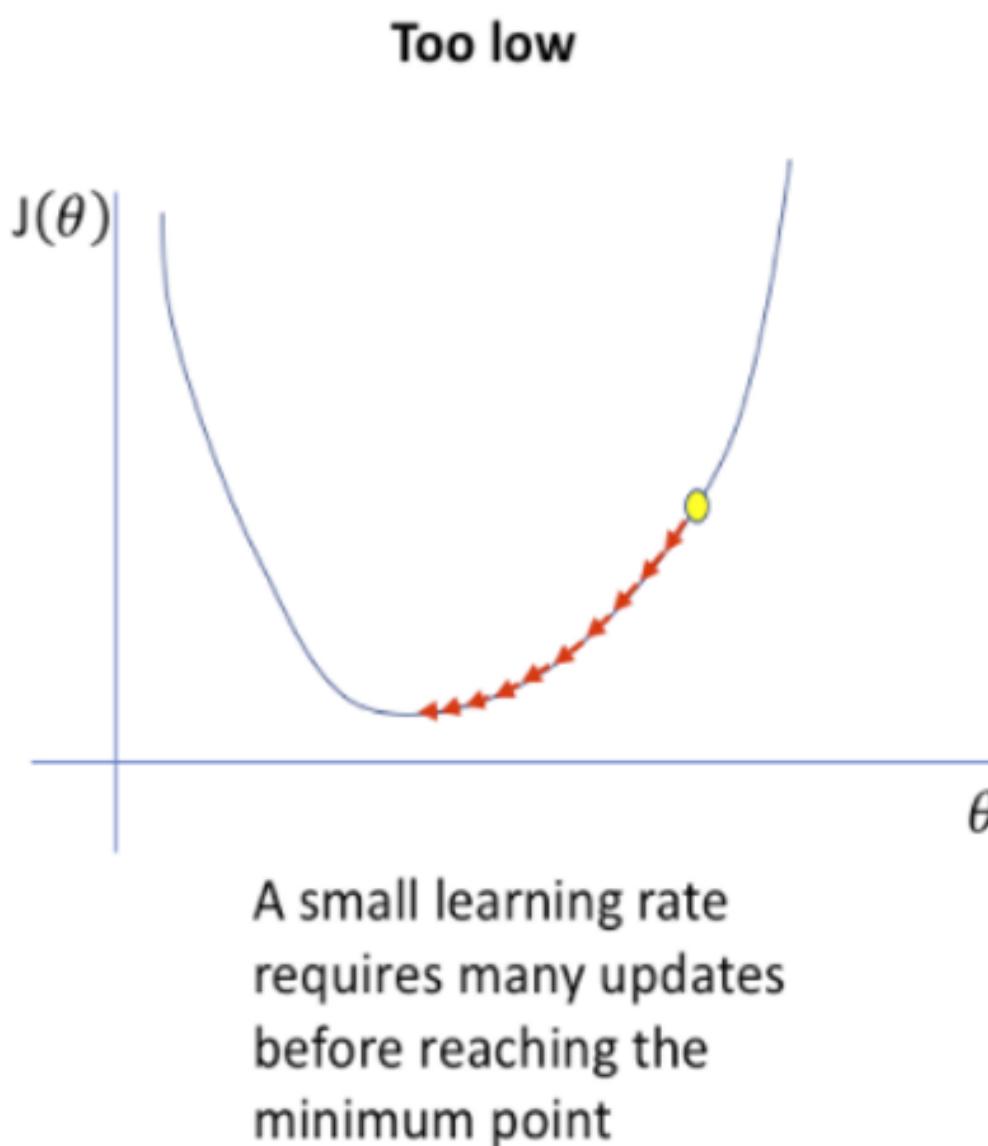
# Backpropagation

- Backpropagation is the backward propagation of small changes in weights and biases within a Neural Network
- Backpropagation is when “learning” takes place, and only happens when a Neural Network is being trained



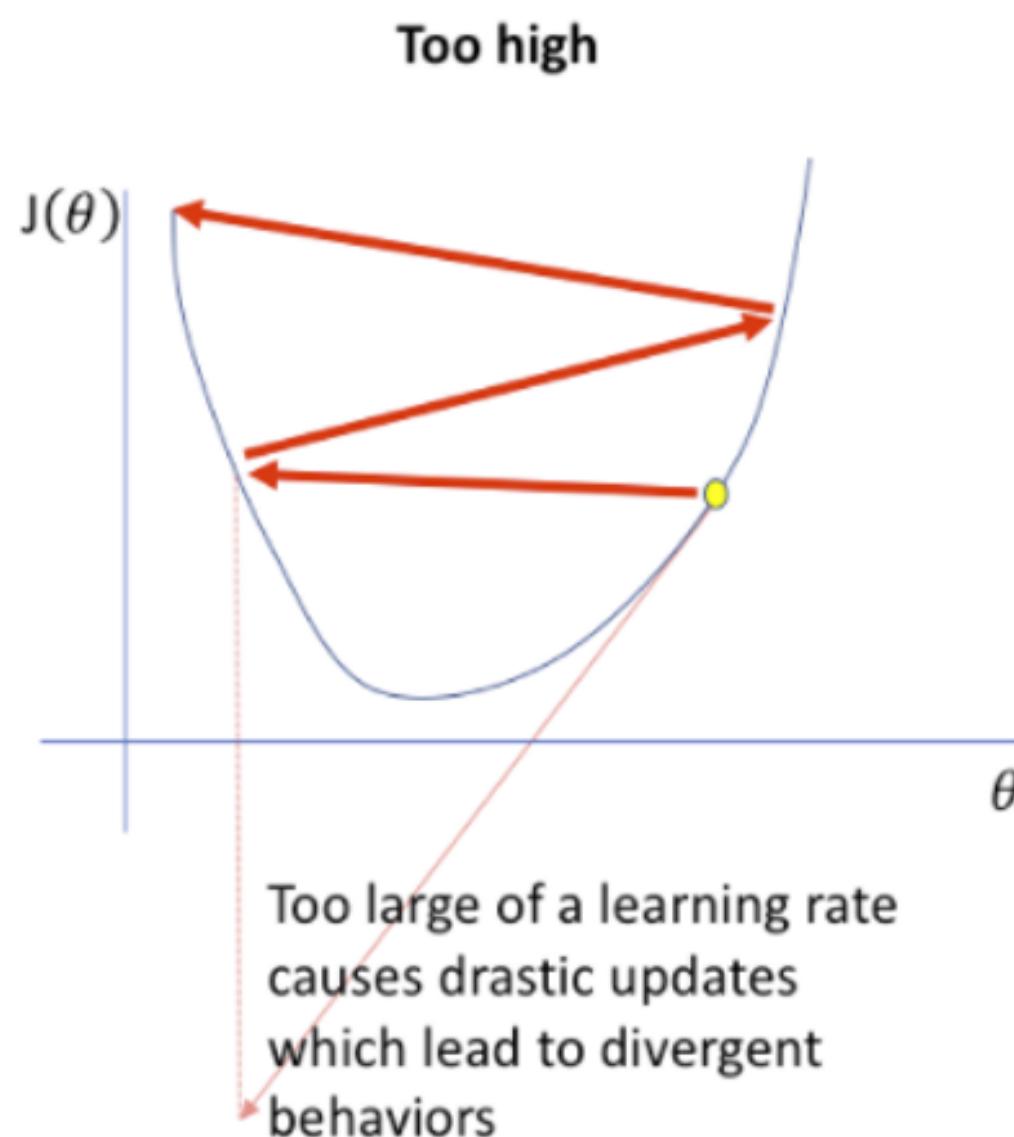
# Learning Rate

- The learning rate controls the amount of adjustments made to our weights and biases in each backpropagation
- A learning rate that is too small causes a neural network to learn slowly



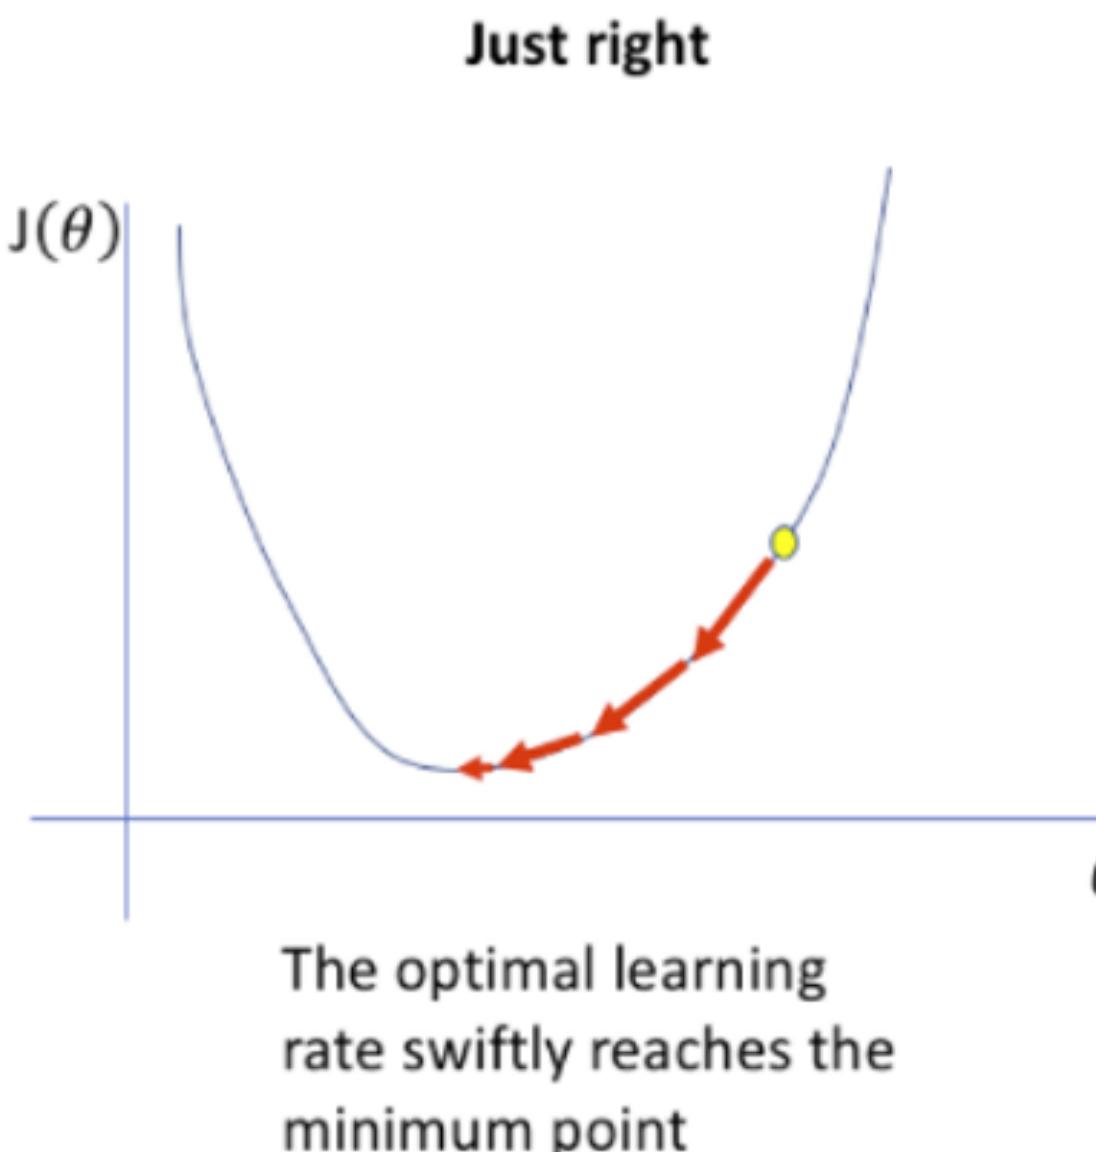
# Learning Rate

- A learning rate that is too large causes a neural network to yield losses that swing wildly from one iteration to the next



# Learning Rate

- An optimal learning rate allows a neural network to learn steadily and swiftly



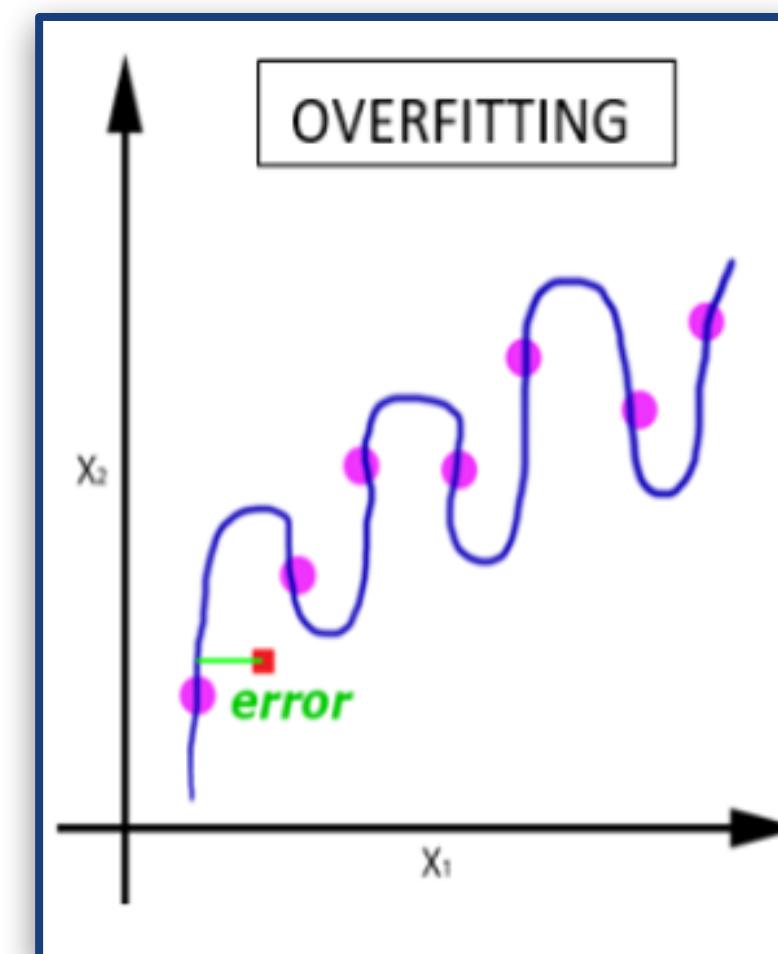
# Fitting a Model

# Data

- Training Data is used to train our model over multiple epochs
  - Epochs are the number of times a model has seen the entire set of data
  - Model updates its weights and biases from its prediction errors using these data (back-propagation takes place)
- Validation and Test Data are used to access the accuracy of our model
  - Model does not learn using these data (no back-propagation during validation and testing phase)

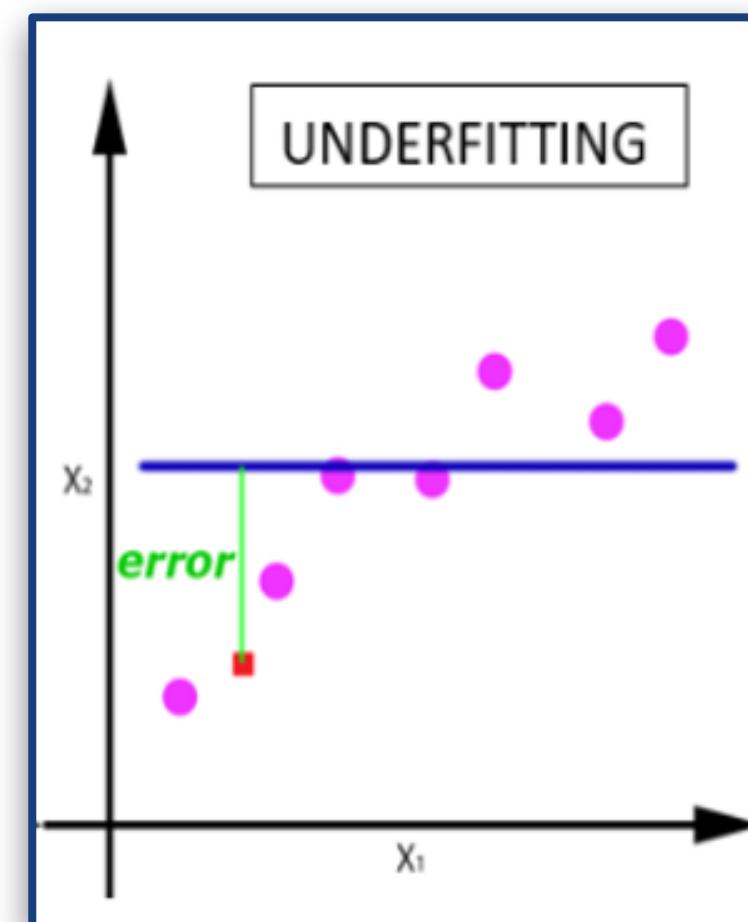
# Over-Fitting

- If our model gives good predictions on training data, but bad predictions on validation and/or test data, then it is over-fitting
- To resolve this, use more and better training data, or apply regularization (add an error term) during backpropagation



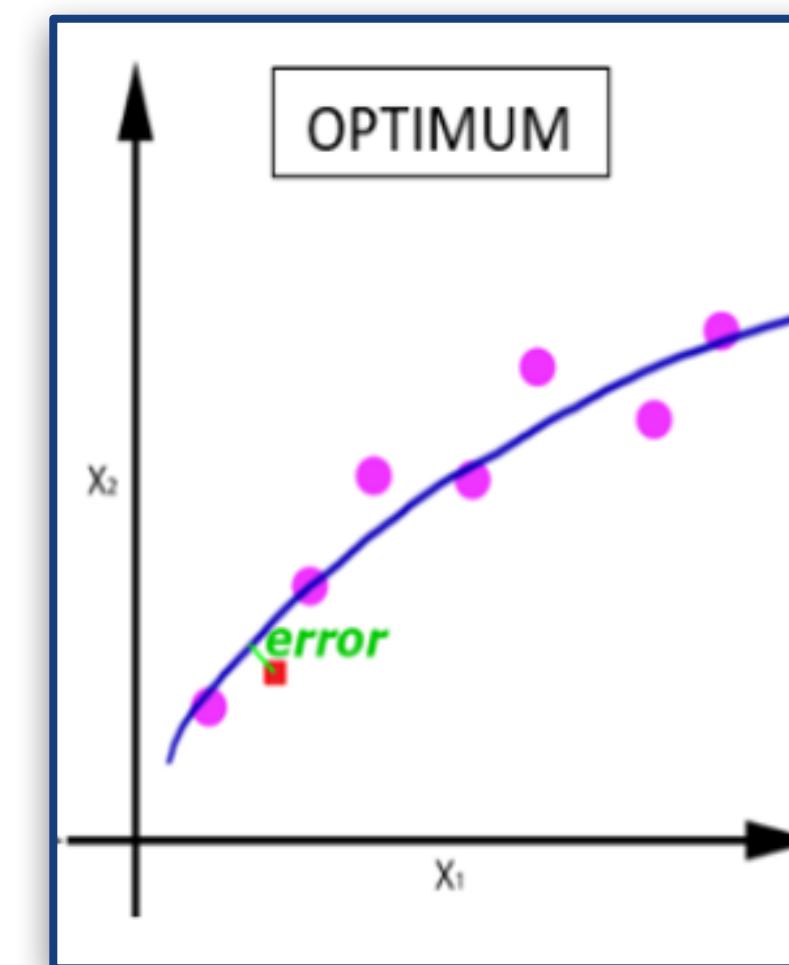
# Under-Fitting

- If our model predicts badly in training, validation and/or test data, it is under-fitting
- To resolve this, train our model longer, add more hidden layers to it or feed it with more and better training data



# Optimal Fitting

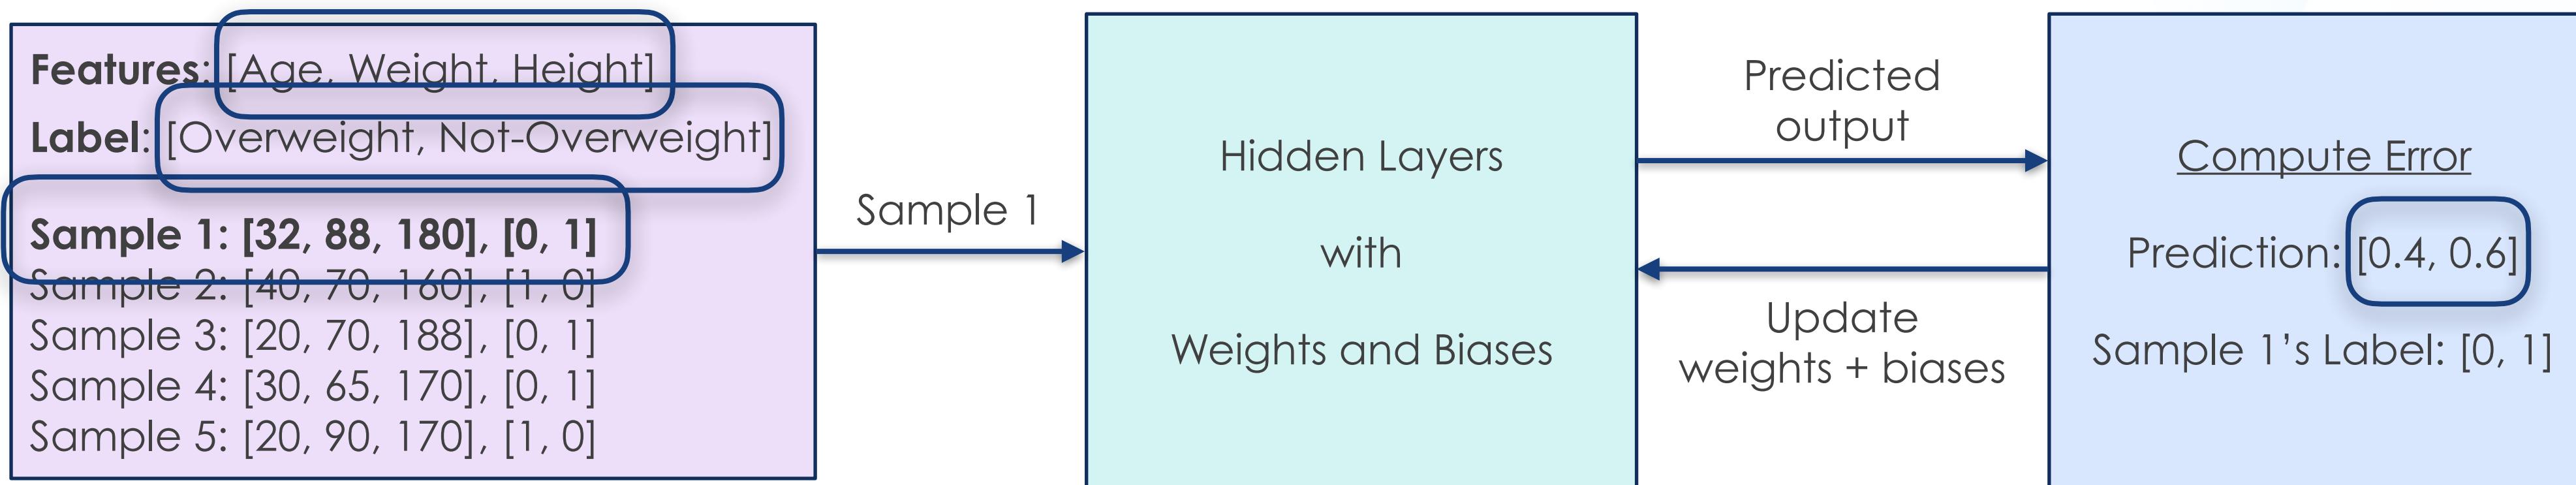
- If a model predicts correctly for its training, validation and/or test data, it has achieved optimal fitting
- Our model is general enough to accommodate new points that follows the distribution of the training data



# Training and Testing

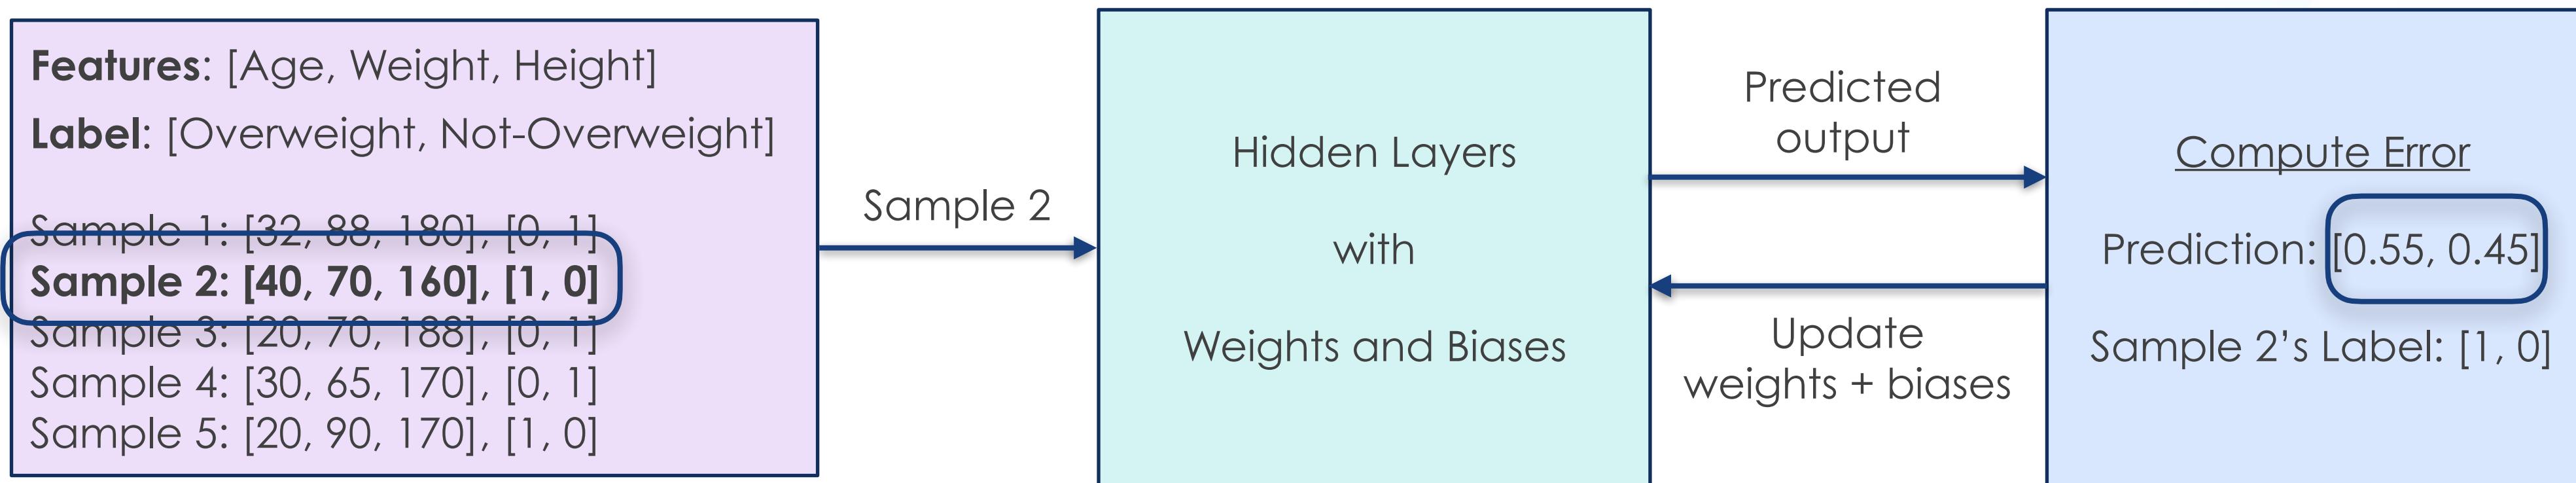
# Training a Neural Network

- Train a Neural Network to predict if a person is overweight starting with Training Sample 1



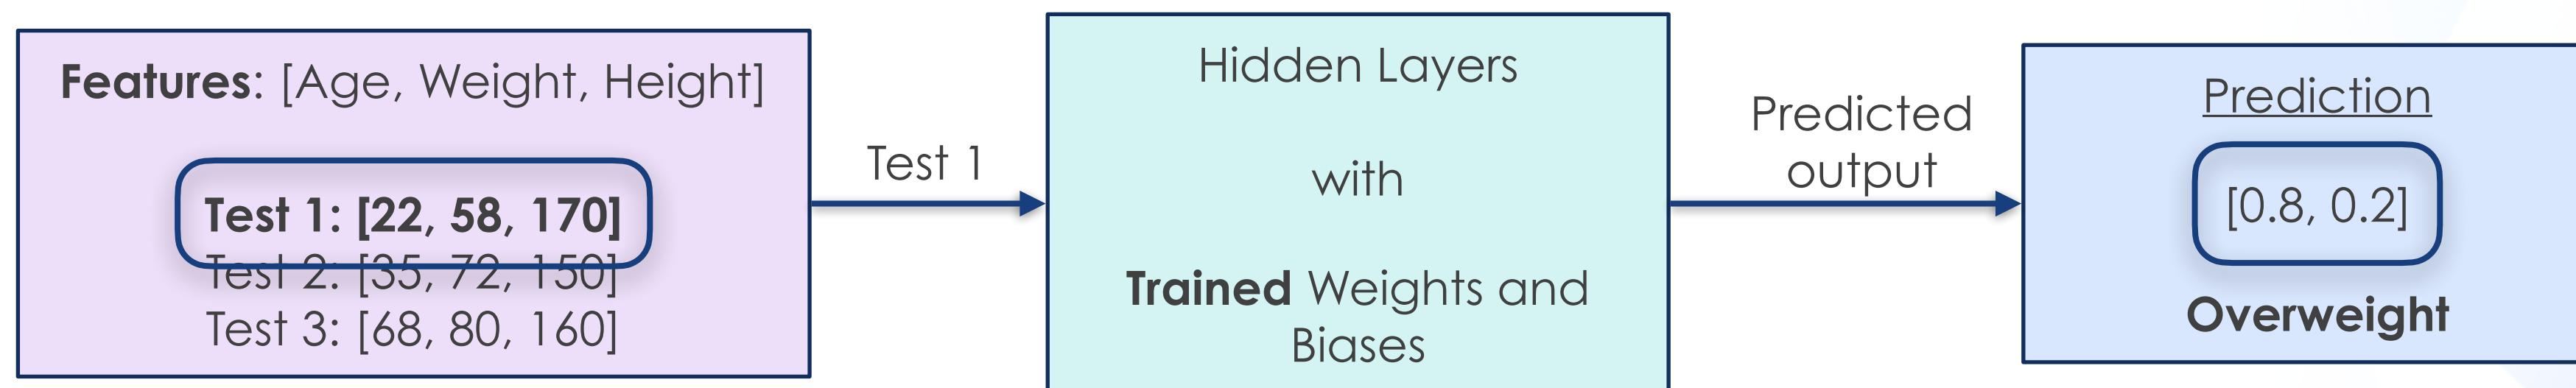
# Training a Neural Network

- Train a Neural Network to predict if a person is overweight using the next training-sample and so on



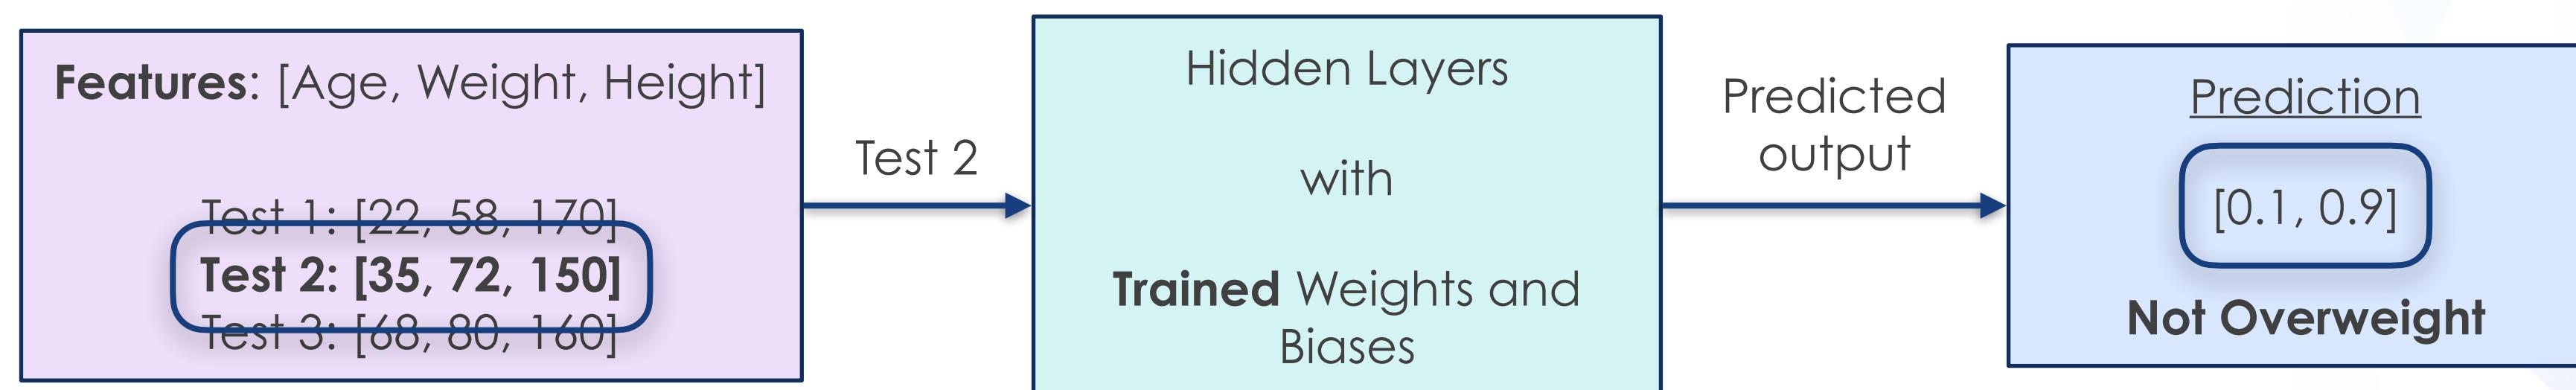
# Using a trained Neural Network

- Using the trained Neural Network to predict if a person is overweight using Test Sample 1



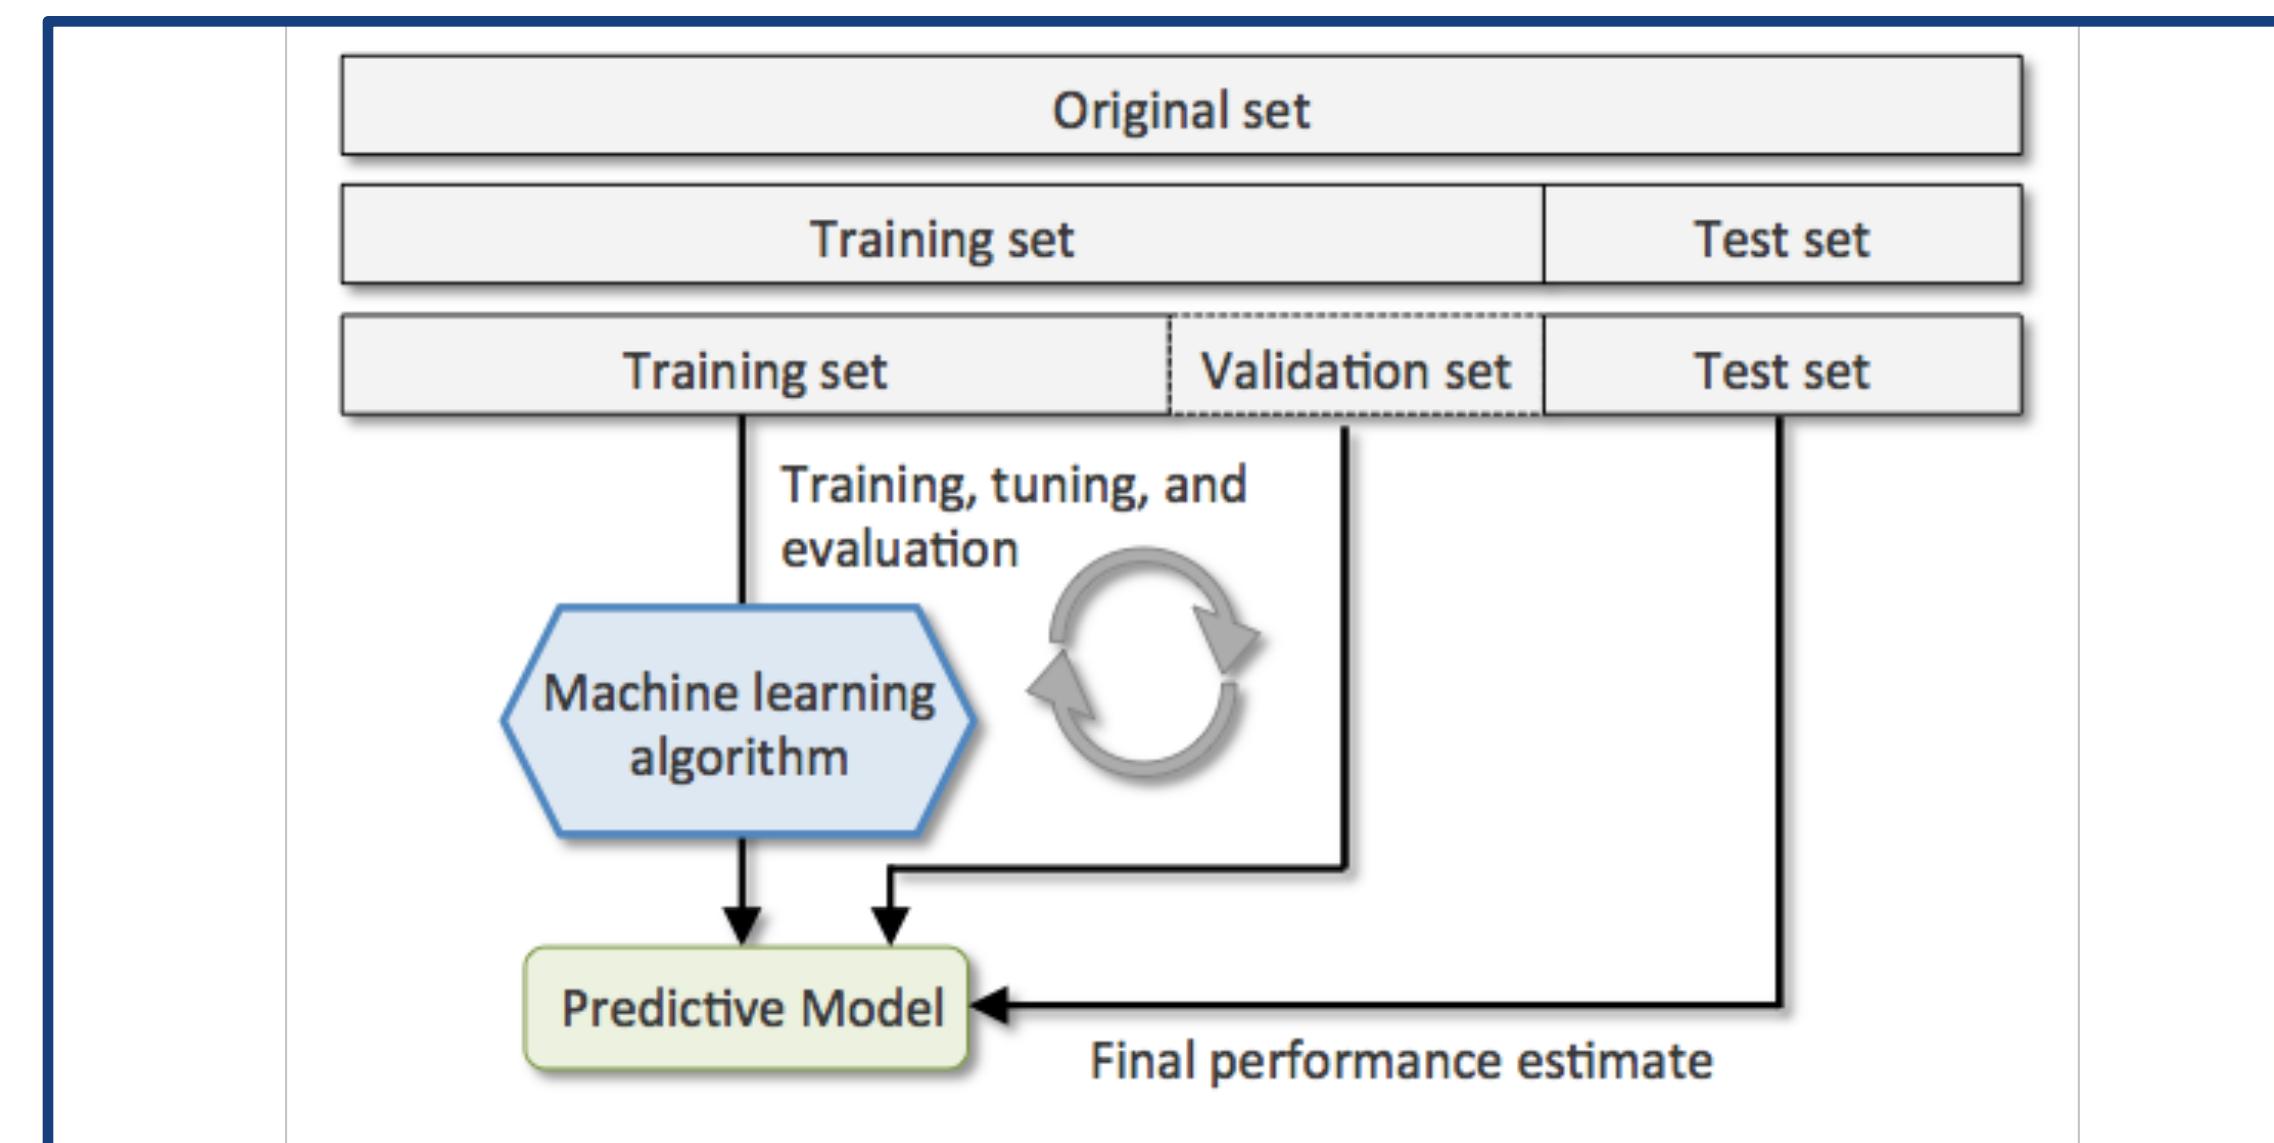
# Using a trained Neural Network

- Using the trained Neural Network to predict if a person is overweight using subsequent test samples



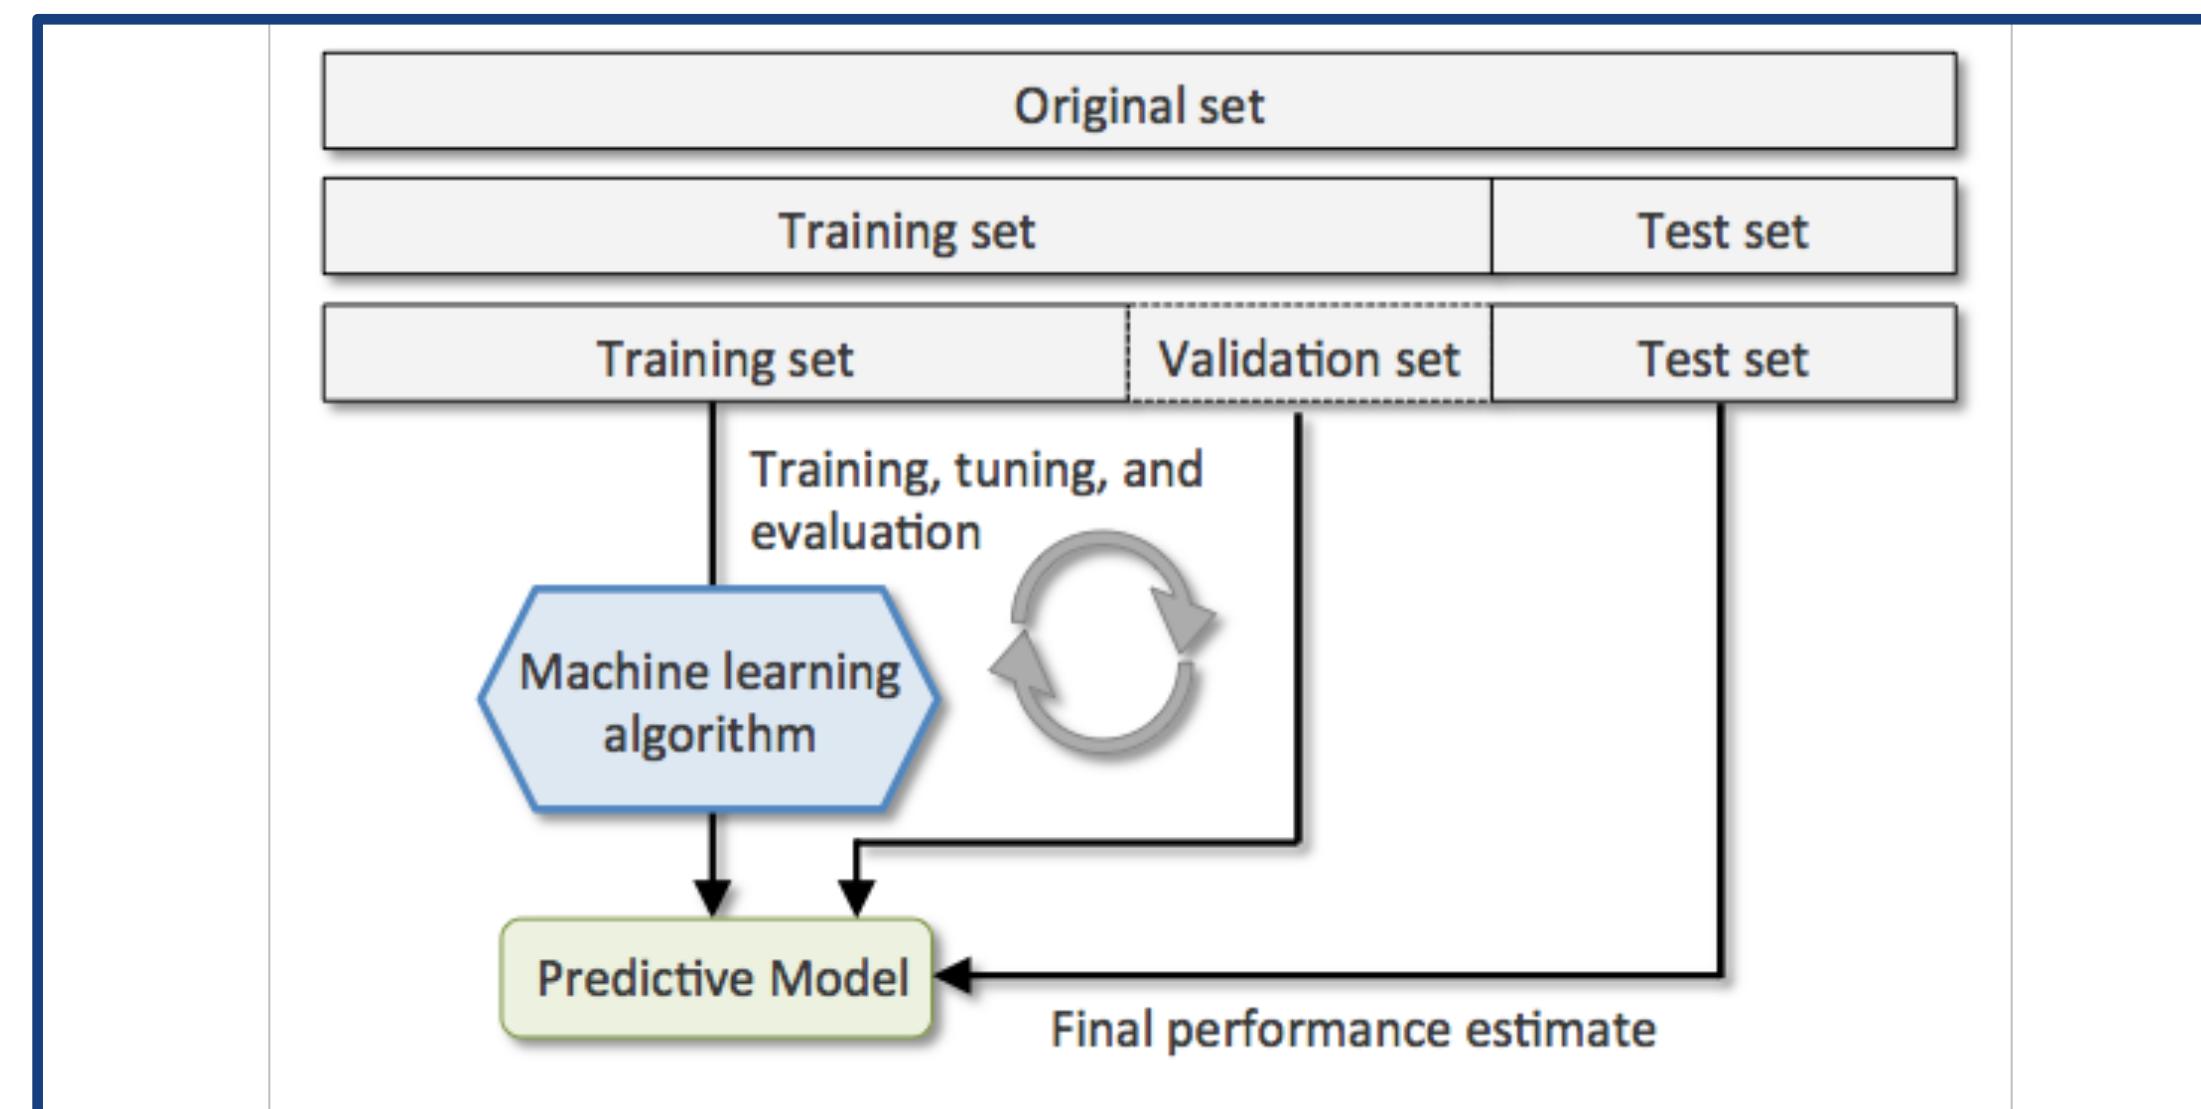
# Training Data

- Training Data is used to train our Neural Network to give accurate predictions over multiple epochs (an epoch is when the Neural Network has seen all training data once)
- Weights and Biases are updated during backpropagation (learning)
- Objective is to minimize error (recall “loss function”) and train a set of final weights and biases that gives accurate predictions



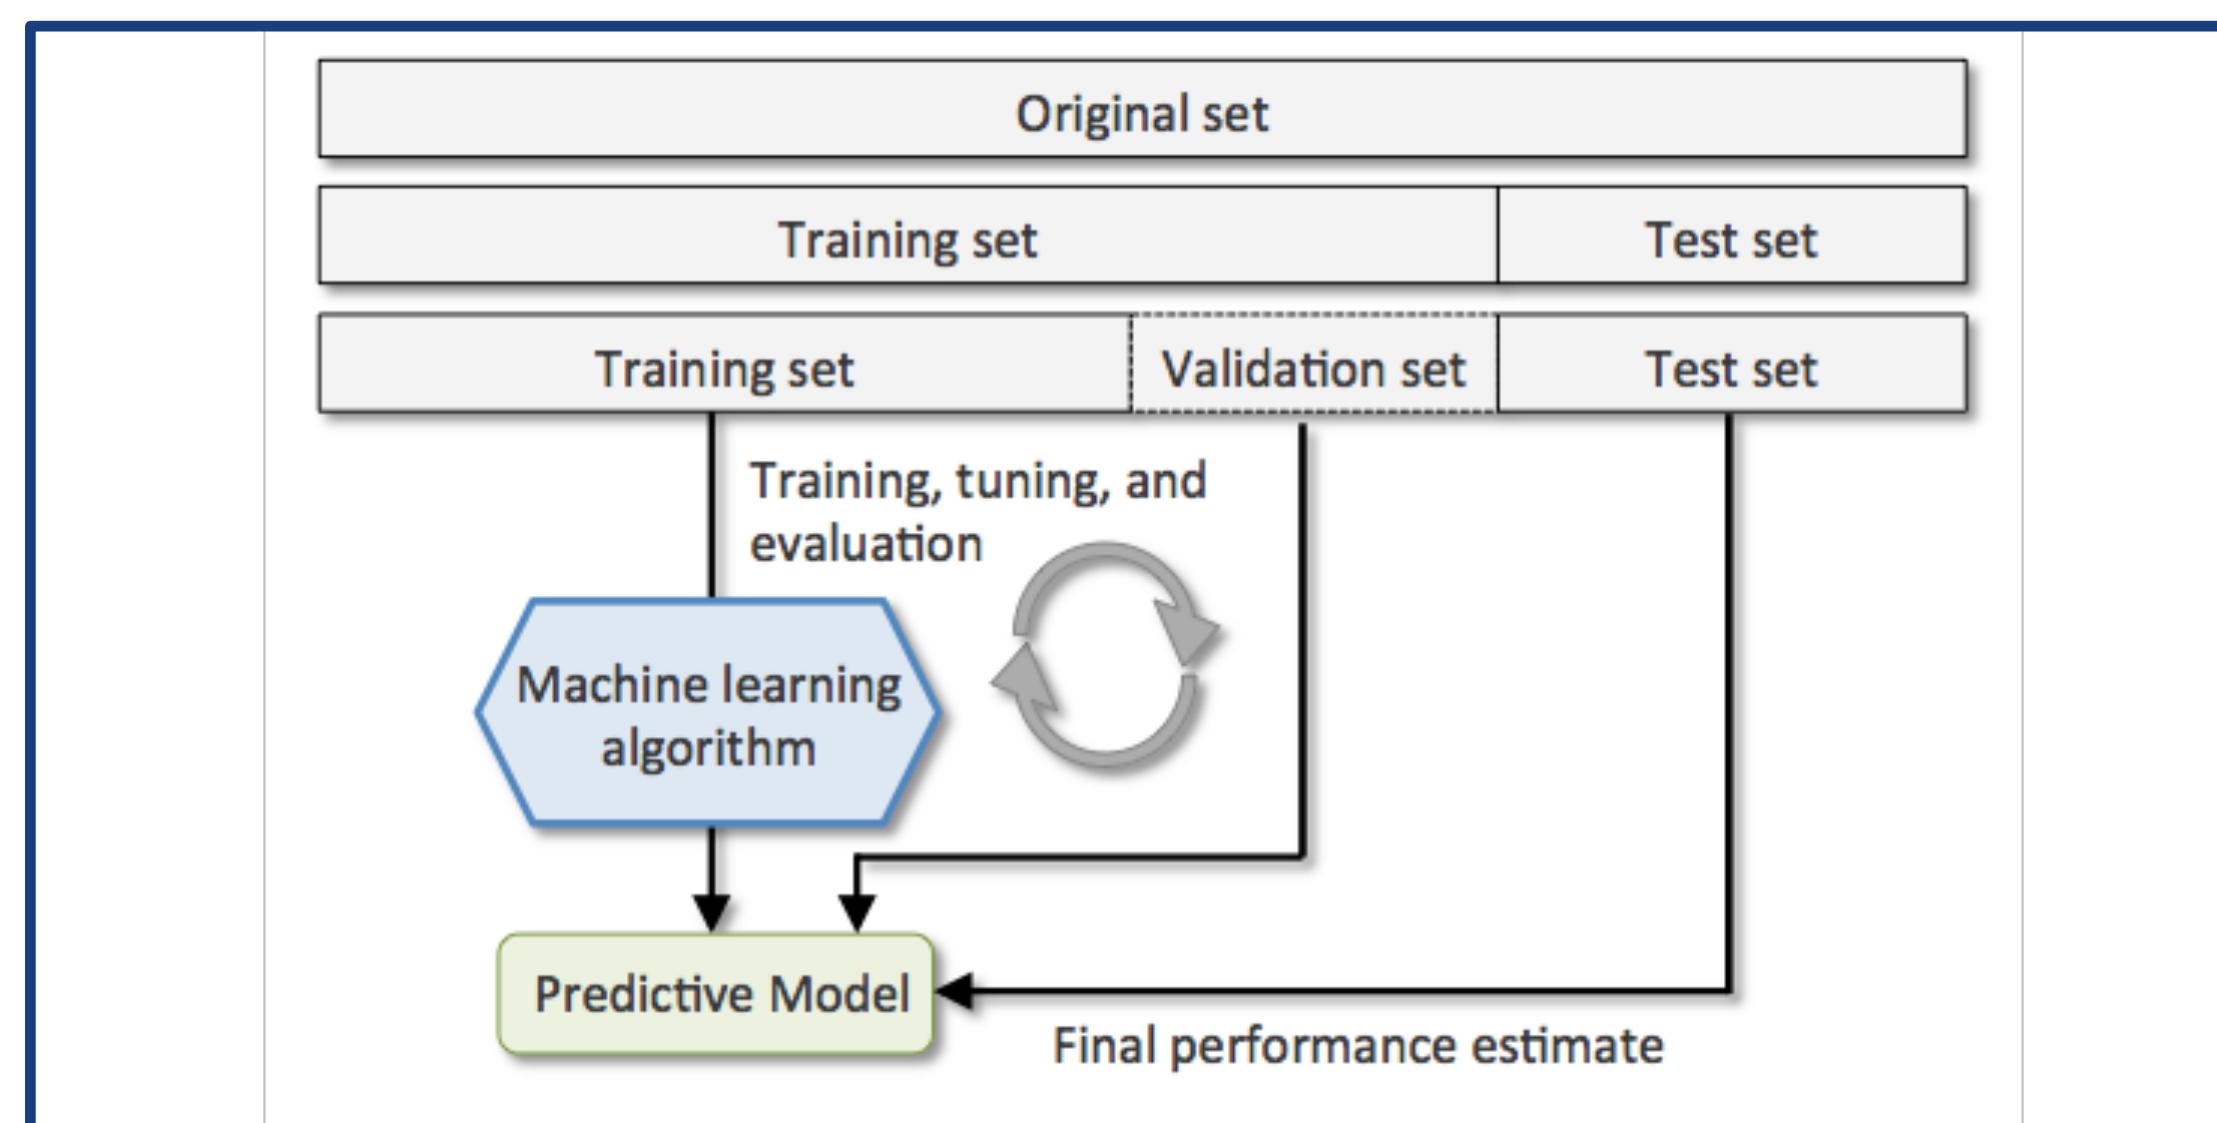
# Validation data

- Validation Data is used to assess how accurate our Neural Network is, based on our training so far
- Only forward pass (i.e. prediction); no backpropagation and updating of weights and biases (i.e. no learning takes place)
- Objective is to compare the accuracies of our Neural Networks that have different architectures



# Test data

- Test Data is to gauge the accuracy of a trained neural network
- Data points in the training data are usually excluded from the test data
- As a neural network has not seen the Test Data before, Test Data give us an unbiased estimate of its predictive accuracy



# Tensorflow and Keras

# TensorFlow

- TensorFlow, developed by Google, is a software library for machine learning
- Under the hood, we will leverage on TensorFlow to create, train and test our neural networks (or models)
- Keras is the high-level API of TensorFlow 2 and is a more approachable interface to work with TensorFlow

# Keras (Adding example)

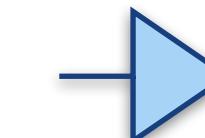
- Let's now create a neural network and train it to count 3 numbers
- So, we want to train a neural network by showing it examples of adding 3 numbers
  - E.g.  $2 + 3 + 4 = 9$
  - E.g.  $6 + 2 + 6 = 14$
- And when training is over, we will give him 3 numbers
  - E.g.  $7 + 10 + 3 = ?$
- Our trained neural network should give us a value that is very close to 20

# Keras (Adding example)

- First, create our training data by randomly generating tuples of 3 values

```
import numpy as np
import tensorflow as tf

# prepare training features
x_train = np.random.randint(-100, 100, size=(300, 3))
print("x_train.shape =", x_train.shape)
print("x_train =", x_train)
```

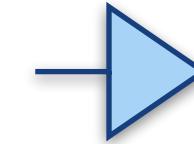


```
x_train.shape = (300, 3)
x_train = [[ 39 -23 -75]
 [ 58 -96 -43]
 [ 4  25 -20]
 [ 4  11  27]
 [-35 -23  34]
 [ 50  94 -90]
 [ 52  44  52]
 [ 63 -65  35]
 [ 83 -57 -90]
 [-61 -43 -59]
 [-9  54  37]]
```

# Keras (Adding example)

- Next, compute the sums of generated tuples
- These values are the “labels” to our samples

```
y_train = np.array([[sum(x)] for x in x_train])
print("y_train.shape =", y_train.shape)
print("y_train =", y_train)
```



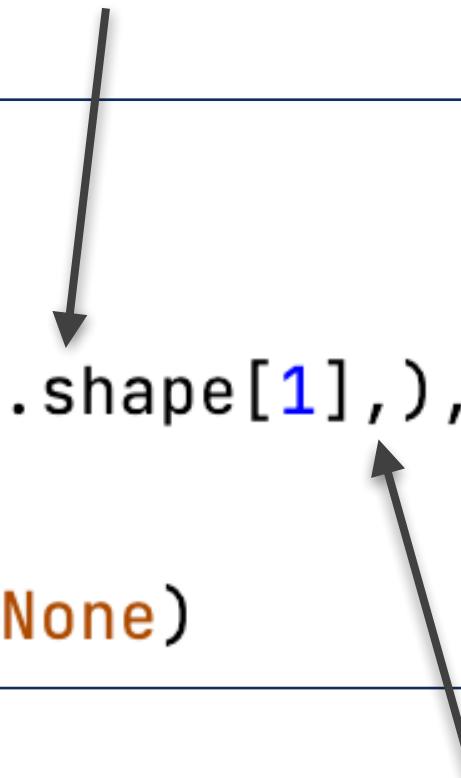
```
y_train.shape = (300, 1)
y_train = [[ -59]
[ -81]
[  9]
[ 42]
[ -24]
[ 54]
[ 148]
[ 33]
[ -64]
[-163]
[ 82]]
```

# Keras (Adding example)

- Create our model by using Sequential()
- Add hidden layers via Dense(...)
- The loss function is Mean Square Error

```
# construct neural network model
model = tf.keras.Sequential()
model.add(tf.keras.layers.Dense(500,
                               input_shape=(x_train.shape[1],),
                               activation='relu'))
model.add(tf.keras.layers.Dense(1))
model.compile(optimizer='adam', loss='mse', metrics=None)
```

no. of features in each training sample



must have the comma if  
there is only one dimension

# Keras (Adding example)

- Our model start training when `fit(...)` is called
- `x_train` contains training **features**
- `y_train` contains training **labels**
- Our Neural Network will be trained with our training **samples** 500 times (i.e. `epochs=500`)

```
# train the model
model.fit(x_train, y_train, epochs=500, verbose=1)
```

# Keras (Adding example)

- Once our training has completed, create some test data to test our trained neural network
- Use evaluate(...) to determine the loss (the difference between the ground truth and prediction)

```
# generate test data
x_test = np.random.randint(-100, 100, size=(20, 3))
y_test = np.array([[sum(x)] for x in x_test])

# perform auto-evaluation
loss = model.evaluate(x_test, y_test, verbose=1)
print('Loss = ', loss, '\n')
```

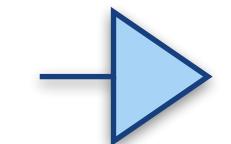


Loss = 0.026681125164031982

# Keras (Adding example)

- Examine our model's predicted results

```
# perform prediction
predictions = model.predict(x_test)
for i in np.arange(len(predictions)):
    print('Data: ', x_test[i],
          ', Actual: ', y_test[i],
          ', Predicted: ', predictions[i])
```



Data: [78 60 72] , Actual: [210] , Predicted: [209.98482]
Data: [-29 -33 74] , Actual: [12] , Predicted: [12.3939495]
Data: [72 16 98] , Actual: [186] , Predicted: [185.92139]
Data: [98 53 88] , Actual: [239] , Predicted: [239.03915]
Data: [-24 25 -41] , Actual: [-40] , Predicted: [-40.075893]
Data: [ 47 3 -17] , Actual: [33] , Predicted: [32.941364]
Data: [-77 -15 33] , Actual: [-59] , Predicted: [-58.995853]
Data: [-10 23 82] , Actual: [95] , Predicted: [94.7834]
Data: [-57 61 1] , Actual: [5] , Predicted: [4.686406]
Data: [ -1 -63 38] , Actual: [-26] , Predicted: [-26.01451]
Data: [ 68 3 -94] , Actual: [-23] , Predicted: [-23.03826]
Data: [-66 1 91] , Actual: [26] , Predicted: [26.239145]
Data: [-95 -45 59] , Actual: [-81] , Predicted: [-80.65134]
Data: [74 37 11] , Actual: [122] , Predicted: [121.929184]
Data: [-100 -24 68] , Actual: [-56] , Predicted: [-55.926502]
Data: [-11 -70 -69] , Actual: [-150] , Predicted: [-150.02951]
Data: [-11 -58 -52] , Actual: [-121] , Predicted: [-121.001015]
Data: [ -8 53 -57] , Actual: [-12] , Predicted: [-12.055797]
Data: [-63 -72 -70] , Actual: [-205] , Predicted: [-204.99951]
Data: [-37 86 -78] , Actual: [-29] , Predicted: [-29.145395]

# The End